
The Concept of GPU Compiler

Release 0.6

Chen Chung-Shu

Sep 28, 2026

CONTENTS

1	About	1
2	3D Modeling	3
3	Graphics HW and SW Stack	21
4	Basic geometry in computer graphics	31
5	OpenGL	53
6	GPU Architecture	115
7	Software Structure	155
8	Open Sources	179
9	Resources	181

- *Authors*
- *Build steps*
- *Revision history*
- *Licensing*
- *Motivation*
- *Preface*

1.1 Authors

Chen Chung-Shu 陳鍾樞

1.2 Build steps

<https://github.com/Jonathan2251/sa/blob/master/README.md>

1.3 Revision history

Version 0.7, not released yet.

Version 0.6, Released September 27, 2026.

opengl.rst: refine section '3D Rendering'. hw.rst: add section 'Framebuffer hierarchy'. hw.rst: move section 'System Features —Buffers' to opengl.rst 'OpenGL Buffers' and refine it. sw.rst: refine section 'Unified IR Conversion Flows'.

Version 0.5, Released September 20, 2026.

opengl.rst: add section 'Mesh Construction'. refine section '3D Rendering'. hw.rst: refine section 'Geometry Units'. Refine English.

Version 0.4, Released September 7, 2026.

opengl.rst: Add section 'Ray Tracing Pipeline'. opengl.rst: Refine the meshlet section. hw.rst: Add dgpu.gv and igpu.gv. hw-sw-stack.rst: Refine the VSync section. Correct typos.

Version 0.3, Released March 14, 2026.

geo-math.rst: refine ‘section Projection’ and ‘section Mapping data in GPU’. sw.rst: add ‘section Vector Processor’.

Version 0.2, Released February 28, 2026.

Reorganize sections.

Version 0.1, Released February 22, 2026.

Initial version.

1.4 Licensing

Chen Chung-Shu

1.5 Motivation

When I began my career as a graphics GPU compiler engineer, I could not find a book that clearly explained the material from a compiler engineer’s perspective. As a result, I decided to write this book and have been adding to it over time as I continue my studies.

1.6 Preface

Basically, a CPU is a SISD (Single Instruction Single Data) architecture in each core. The multimedia instructions in CPUs are smaller-scale forms of SIMD (Single Instruction Multiple Data), while GPUs are large-scale SIMD processors, capable of coloring millions of image pixels in just a few milliseconds.

Since 2D and 3D graphic processing offers great potential for parallel data processing, GPU hardware typically includes tens of thousands of functional units per chip, as seen in products by NVIDIA and other manufacturers.

This book provides an overview of how 3D animation is created and executed on a CPU—GPU system. It then introduces GPU compilers and hardware features relevant to graphics applications. Finally, it explains how GPUs have taken on computational tasks traditionally handled by CPUs through the GPGPU (general-purpose computing on graphics processing units) concept and related standards.

Website: Basic Theory of 3D Graphics with OpenGL¹.

¹ https://www3.ntu.edu.sg/home/ehchua/programming/opengl/CG_BasicsTheory.html

3D MODELING

- *Animation*
- *Node-Editor (shaders generator)*
 - *Node-Editor*
 - *Code Generation from Node-Editor*
- *3D Modeling Tools*

By creating 3D models with triangles or quads on a surface, the model is formed using a polygon mesh¹. This mesh consists of all the vertices shown in the first image as Fig. 2.1.




Fig. 2.1: Creating 3D model and texturing

¹ <https://www.quora.com/Which-one-is-better-for-3D-modeling-Quads-or-Tris>

After applying smooth shading^{Page 3, 1}, the vertices and edge lines are covered with color (or the edges are visually removed —**edges never actually have black outlines**). As a result, the model appears much smoother².

Furthermore, after texturing (texture mapping), the model looks even more realistic³.

2.1 Animation

★ Animation Layers: High → Low

This breakdown organizes animation systems from the **highest gameplay logic** down to the **lowest GPU skinning**, and clearly marks which parts are controlled by the **user** and which parts are handled by the **3D engine**.

1. Gameplay Animation Logic (High Level): set by user (game developer)

See video here⁴.

Examples

- Play “walk” when speed > 0.1
- Trigger “jump” on button press
- Switch to “attack” when enemy detected
- Blend run when velocity increases

Where it lives

Live in **gameplay scripts** (C#, Blueprints, GDScript, Python)

- Unity: C# scripts
- Unreal: Blueprints or C++
- Godot: GDScript
- ThinMatrix: Java (no scripting layer)

This layer decides *when* animations should play.

2. Animation State Machine / Animation Graph: set by user (game developer)

Examples

- Idle → Walk → Run transitions
- Blend trees
- Animation layers (upper body, lower body)
- Animation parameters (speed, grounded, direction)

Where it lives

- Unity: Animator Controller
- Unreal: Animation Blueprint
- Godot: AnimationTree
- ThinMatrix: *Does not have this layer*

This layer controls *which* animation clip is active and how transitions occur.

² <https://en.wikipedia.org/wiki/Shading>

³ https://en.wikipedia.org/wiki/Texture_mapping

⁴ https://www.youtube.com/watch?v=7QIcd6_TTys

3. Animation Clip Playback Layer: user chooses

An **animation clip** is a **sequence of keyframes** over a period of time that represents a **motion or action**.

Examples

- Play animation clip
- Loop animation
- Set animation speed
- Crossfade between clips
- Blend two clips together

Who sets it? User chooses which clip to play.

Who implements it? Engine handles blending, timing, and playback.

Where it lives

- Unity: Mecanim runtime
- Unreal: AnimInstance
- Godot: AnimationPlayer
- ThinMatrix: Java engine code (Animator.java)

This layer executes the user's choices.

4. Skeleton Animation System (Low Level): 3D engine implements it

Examples

- Bone hierarchy
- Keyframe interpolation
- Joint transforms
- Matrix palette generation
- Pose calculation

Who sets it? Engine

Who uses it? User indirectly, by playing animations.

Where it lives

- Unity: C++ engine core
- Unreal: C++ engine core
- Godot: C++ engine core
- ThinMatrix: Java engine code (he writes this manually)

This layer performs the mathematical work of animation.

5. GPU Skinning (Lowest Level): 3D engine implements it

Examples

- Vertex shader skinning
- Applying bone matrices
- Weighted vertex deformation

- Sending matrices to GPU

Who sets it? Engine

Who uses it? User never touches this layer directly (except in custom engines).

Where it lives

- Unity: C++ + HLSL
- Unreal: C++ + HLSL
- Godot: C++ + GLSL
- ThinMatrix: GLSL shader he writes manually

This is the final stage where the GPU deforms the mesh.

Full Hierarchy (Summary)

```
HIGH LEVEL (User)
-----
1. Gameplay Animation Logic
2. Animation State Machine / Animation Graph
3. Animation Clip Playback

LOW LEVEL (Engine)
-----
4. Skeleton Animation System
5. GPU Skinning
```

Developers only write the highest-level animation logic and designed transitions & blends as shown in Fig. 2.2. The engine automatically handles all lower-level animation work. Like the video^{Page 4, 4}, Jason's tutorials operate only in Level 1 and Level 2:

✓ Level 1 —Scripts

He writes code like:

```
animator.SetFloat("Speed", speed);
```

✓ Level 2 —State Machine

He configures transitions and parameters.

ThinMatrix's engine collapses the top three layers into Java because it has **no scripting layer** and **no animation graph**, so the user must modify the engine code directly.

According to the video on ThinMatrix's skeleton animation⁵, he is sampling the animation clip at different times. The animation clip already contains: keyframes, bone transforms, timestamps, interpolation curves. All of this comes from Blender's exported .dae file. Joints are positioned at different poses at specific times (keyframes), as illustrated in Fig. 2.3.

Although most of 3D game engines are written C++, **ThinMatrix's engine** is 100% Java. In this series of videos, you will see that he writes new Java engine modules, edits existing engine code, loads animation data from Blender, interpolates keyframes, updates bone matrices and sends them to the GPU. Because ThinMatrix's engine is **tiny and educational** for engine programmer or game developer, does not provide Scripting Layer (such as C#, Python, GDScript, Blueprints) most commercial 3D engines. Instead, he modifies ThinMatrix's Java engine directly, which differs from most other 3D engines operate.

Animation flow

⁵ <https://www.youtube.com/watch?v=f3Cr8Yx3GGA>




Fig. 2.2: Animation levels




Fig. 2.3: Get time points at keyframes


Every modern 3D animation tool comes with its own **built-in render engine**, and often more than one. In 3D game design, **game engines (Unity, Unreal, Godot)** use **real-time engines** for real-time animation.

Pipeline: Blender → Engine → OpenGL

```
+-----+
|  Blender  |
| (Modeling Tool) |
+-----+
|
| Exports assets:
| - Meshes (.obj, .fbx, .glTF)
| - Armatures / bones
| - Animations (keyframes)
| - Textures (PNG, EXR, TGA)
|
v
+-----+
|  Game Engine  |
| (Unity, Unreal, Godot, |
| LWJGL, JOGL, Custom) |
+-----+
|
| Engine loads assets:
| - Parses mesh data
| - Loads skeletons
| - Loads animation curves
| - Loads materials/shaders
|
| Engine code you write:
| - Java (JOGL/LWJGL)
```

(continues on next page)

(continued from previous page)



Note**3D modeling tools do store animation and movement data**

—but they do NOT store any rendering or API-specific code.

Game engines do store animation data

—but programmers still write the logic that plays, blends, and controls those animations.

Animation Data vs. Movement Speed in Games

List the animation types in a table for inclusion in this book.

Table 2.1: Animation Types

Animation Type	What Moves	Description	GPU Requirement
Transform Animation	Object transform	The entire mesh moves as a rigid body using position, rotation, and scale. No vertex-level deformation occurs.	Optional (fixed-function or shaders)
Skinning	Vertex positions	Vertices are blended by bone matrices to deform the mesh (arms bending, legs walking). Requires per-vertex matrix blending.	Requires shaders
Morph Target Animation	Vertex positions	Vertices blend between multiple stored shapes (facial expressions, muscle bulges). Uses morph weights to interpolate.	Requires shaders
Procedural Deformation	Vertex positions	Vertices are modified by mathematical functions (wind, waves, noise, squash-and-stretch). Driven by time or simulation parameters.	Requires shaders

Example: Walking Animation: Skinning + Transform Animation

When a character walks in a game, the animation is produced by two different systems working together:

1. **Skinning (Bone Animation)** Skinning is responsible for deforming the mesh. It drives the motion of limbs such as legs, arms, spine, and feet. Without skinning, the character would move as a rigid statue with no bending or articulation.
2. **Transform Animation (Rigid-Body Movement)** Transform animation moves the entire character through the world. This includes translation, rotation, and root motion. Without transform animation, the character would walk in place without actually moving forward.
 - Root bone: the bone that represents the entire object's transform —the top-most parent of the hierarchy. An example of person:

```

Root
├── Pelvis
│   ├── Spine
│   │   ├── Chest
│   │   │   ├── Neck
│   │   │   │   ├── Head
│   │   │   │   └── Shoulders
│   │   │       ├── Arm_L
│   │   │       └── Arm_R
│   └── ...
└── ...

```

(continues on next page)

(continued from previous page)

```
├ Leg_L
└ Leg_R
```

Both systems are required to create a complete walking animation:

- **Skinning** provides the internal limb motion.
- **Transform animation** provides the external world-space movement.

Together, they produce the final effect of a character walking naturally through the environment.

The following explains what animation data 3D modeling tools store, what game engines store, and what programmers must implement manually. It also clarifies the relationship between animation curves, movement speed, and gameplay logic.

1. What 3D Modeling Tools Actually Store

3D modeling and animation tools such as Blender, Maya, and 3ds Max store **animation data**, not gameplay logic.

They **do store**:

- Keyframes (frame 0, frame 10, frame 24, etc.)
- Bone transforms at each keyframe
- Interpolation curves (Bezier, linear, quaternion)
- Animation duration (e.g., 1.2 seconds)
- Frame rate (e.g., 24 fps)
- Skeleton hierarchy
- Skin weights (vertex-to-bone influences)
- Optional root bone motion (displacement over time)

Modeling tools produce **data**, not rendering code and not gameplay rules in OpenGL/DirectX code.

Example:

```
Bone "Arm" rotation at frame 0 = (0°, 0°, 0°)
Bone "Arm" rotation at frame 10 = (45°, 0°, 0°)
```

2. What Game Engines Actually Store

The **engine's built-in C++ renderer handles all OpenGL/Vulkan/Metal calls automatically**. Game engines such as Unity, Unreal Engine, Godot, or custom engines store and manage animation data, but still do not define gameplay movement speed.

They **do store**:

- Animation clips
- State machines (Idle → Walk → Run)
- Blend trees
- Transition rules
- Animation events
- Curves for rotation, scaling, and root motion

Again: data, not OpenGL/DirectX code.

Example:

```
If speed < 0.1 → Idle
If speed > 0.1 → Walk
If speed > 4.0 → Run
```

This is engine logic, not GPU code.

Game engines interpret animation data but rely on programmer logic to control how characters move in the world.

3. What Programmers Must Implement

Programmers write the **logic** that uses animation data to move objects.

Examples:

In Unity (C#)

```
animator.SetFloat("speed", playerVelocity);
...
float speed = 3.5f;
transform.position += direction * speed * Time.deltaTime;
```

In a custom engine (C++/OpenGL)

```
shader.setMatrix("boneMatrices[0]", boneMatrix);
...
float velocity = 3.5f;
position += velocity * deltaTime;
```

In JOGL/LWJGL (Java)

```
glUniformMatrix4fv(boneLocation, false, boneMatrixBuffer);
...
float velocity = 3.5f;
position += velocity * deltaTime;
```

Programmers write:

Programmers implement:

- Movement speed
 - E.g. Set the value for speed or velocity as the code above.
- Acceleration and deceleration
- Physics integration
- AI movement
- Player input
- Animation blending logic
- Uploading bone matrices to the GPU
- GLSL shader code for skinning

Animation data is *used* by code, not replaced by it.

4. Root Motion vs. Movement Speed

Some animations include **root motion**, where the root bone moves forward during a walk cycle. Modeling tools export this as bone displacement over time, but they still do **not** define speed.

Example:

If the root bone moves 1 meter in 0.5 seconds, the engine can compute:

```
speed = 1m / 0.5s = 2 m/s
```

However:

- Blender does not store “2 m/s”
- The engine derives speed from displacement
- Programmers decide whether to use root motion or in-place animation

5. Summary Table

Concept	Stored in Blender?	Stored in Engine?	Controlled by Programmer?
Keyframes	Yes	Yes	No
Bone transforms	Yes	Yes	No
Animation length	Yes	Yes	No
Movement speed	No	Yes (derived)	Yes
Physics movement	No	Yes	Yes
AI movement	No	Yes	Yes

6. Final Clarification

- **3D modeling tools store animation timing, not gameplay speed.**
- **Game engines store animation clips, not movement speed.**
- **Programmers control movement speed, physics, and gameplay behavior.**
- **No tool generates JOGL/OpenGL/Vulkan/DirectX code.**
- **All rendering API calls are written by engine developers or by you in a custom engine.**

Example for accelerating playing

Animation Speed vs Engine Rendering (5× Speed)

The following table shows how animation playback, movement speed, and GPU rendering interact when the gameplay speed is multiplied by five. The animation remains 24 fps internally, but its playback time advances five times faster. The GPU continues to render at 60 fps and samples the animation at the current time.

Property	Original Value	After 5× Speed
Animation FPS (baked)	24 fps	24 fps (unchanged)
Animation Playback Speed	1×	5×
Steps per Second	6 steps/sec	30 steps/sec
Movement Speed	6 m/sec	30 m/sec
GPU Rendering FPS	60 fps	60 fps
Engine Playing Frames (What GPU Displays)	Samples animation at 60 fps	Samples animation at 60 fps (skips/interpolates intermediate animation frames)

Summary:

- The animation does **not** become 120 fps; it is simply played 5× faster.

- The runner appears to take **30 steps per second** and move **30 meters per second**.
- The GPU still renders **60 frames per second**.
- The engine **samples** the animation at each rendered frame, so it effectively displays every fifth animation sample, using interpolation for smoothness. For this case, it may display **1 out of 2 animation frames** from 3D modeling.

2.2 Node-Editor (shaders generator)

- 3D animation tools (Blender, Maya, Houdini) use render engines and node editors for materials, lighting, and effects.
- Game engines (Unity, Unreal, Godot) use real-time engines and node editors for shaders, VFX, and sometimes logic.

A node editor defines the **entire material** that is applied to the surface of a 3D object. The shader generated from the node graph runs on **every pixel** (fragment) of the object's surface. In this sense, the node editor controls the **whole surface**, not only a specific region.

However, the node graph can include **masks**, **textures**, **vertex colors**, or **procedural patterns** that allow the artist to specify which *parts* of the surface receive a particular effect. These masks do not limit the shader to only part of the surface; instead, they instruct the shader how to behave differently across different regions.

2.2.1 Node-Editor

Example

Let's say you want:

- rust only on the edges
- metal everywhere else

In the node editor:

1. Load a rust texture
2. Load a metal texture
3. Use a mask (curvature or hand-painted)
4. Mix them using a Mix node

The shader still runs on the whole surface, but the mask tells it:

- "Use rust here"
- "Use metal here"

In summary:

- The node editor defines the **full material** for the **entire surface**.
- Artists can use masks or textures to **target specific areas** within that surface.
- The shader still executes globally, but its **output varies** based on the mask inputs.

Thus, a node editor controls the whole surface, while masks determine how different parts of that surface are shaded.

For 3D video game engines, the only case where mask data is inside the model file is vertex colors. Everything else lives in textures or material/shader assets.

Procedural Rust on Edges Using Shader Nodes

To demonstrate how to create a **rust-on-edges** material using Blender's shader node editor as shown in Fig. 2.4. The goal is to reproduce the effect commonly used on metal containers: clean metal on flat surfaces and rust accumulation along exposed edges.




Fig. 2.4: An example to Rust on Edges Using Shader Nodes⁶.

The technique relies on three core ideas:

1. Detecting edges using the *Bevel* or *Pointiness* attribute.
2. Creating a mask that isolates only the worn edges.
3. Blending a rust material with a metal material using that mask.

This workflow is fully procedural and does not require painting or external textures.


Procedural Edge Wear Node Graph (ASCII Diagram) to create Fig. 2.4 in video⁶ includes 1. edge detection, 2. mask breakup, 3. material creation, and 4. final blending as follows:



(continues on next page)

⁶ <https://odysee.com/@jsabbott:d/how-to-make-procedural-edge-wear-in:2>

(continued from previous page)



(continues on next page)

(continued from previous page)

```

Rust BSDF -----+
                  |
                  |
                  |
                Final Edge Wear Mask (Fac)

```

2.2.2 Code Generation from Node-Editor

Node-based shader editors are visual tools used in modern game engines and DCC (Digital Content Creation) software. They allow users to build shaders by connecting nodes instead of writing GLSL, HLSL, or Metal code manually. These editors **do generate shader code automatically**.

1. Node-Based Editors Do Generate Shader Code

Node-based shader editors such as:

- Unity Shader Graph
- Unreal Engine Material Editor
- Godot Visual Shader Editor
- Blender Shader Nodes (for Eevee/Cycles)

all **compile the visual node graph into real shader code**.

Depending on the engine, the generated code may be:

- GLSL (OpenGL / Vulkan)
- HLSL (DirectX)
- MSL (Metal)
- SPIR-V (Vulkan intermediate format)

The generated code is usually not shown to the user, but it is compiled and sent to the GPU at runtime.

2. How Users Generate Shader Code with Nodes

The workflow for generating shader code through a node editor typically looks like this:

1. The user opens the shader editor.
2. The user creates nodes representing:
 - math operations (add, multiply, dot product)
 - texture sampling
 - Determine FS.
 - lighting functions
 - color adjustments
 - UV transformations
3. The user connects nodes visually to define the shader logic.
 - Defines surface color, lighting, and texture sampling → **determine FS**.
 - For simple deformations (waves, wind, dissolve) → **determine the VS**.
 - Generate both **TCS and TES** code for **displacement and subdivision control**.

- **GS** code is typically **written manually** by graphics programmers since **primitives culling and clipping** are not related with model resolution and texture materials.

4. The engine converts the node graph into an internal shader representation.
5. The engine compiles this representation into platform-specific shader code (GLSL, HLSL, MSL, or SPIR-V).
6. The compiled shader is sent to the GPU and used for rendering.

The user never writes the shader code directly; the editor generates it automatically.

3. Who Is the “User” of Node-Based Editors?

The typical users of node-based shader editors are:

Graphics Designers / Technical Artists

- Primary users.
- They create visual effects, materials, and surface shaders.
- They usually do not write GLSL or HLSL manually.
- Node editors allow them to work visually without programming.

Software Programmers / Graphics Programmers

- Secondary users.
- They may create custom nodes or extend the shader system.
- They write low-level shader code when needed.
- They integrate the generated shaders into the rendering pipeline.

In most workflows:

- **Graphics designers** build the shader visually.
- **The engine** generates the shader code.
- **Programmers** handle advanced logic, optimization, or custom nodes.

4. Summary

- Node-based shader editors **do generate shader code** automatically.
- Users generate shaders by connecting visual nodes rather than writing GLSL/HLSL manually.
- The primary “user” is the **graphics designer or technical artist**.
- Programmers support the system by writing custom nodes or low-level shaders when needed.

The shaders introduction is illustrated in the next section *OpenGL*.

2.3 3D Modeling Tools

Every CAD software manufacturer, such as AutoDesk and Blender, has their own proprietary format. To solve interoperability problems, neutral or open source formats were created as intermediate formats to convert between proprietary formats.

Naturally, these neutral formats have become very popular. Two famous examples are STL (with a *.STL* extension) and COLLADA (with a *.DAE* extension). Below is a list showing 3D file formats along with their types.

Table 2.2: 3D file formats⁷

3D file format	Type
STL	Neutral
OBJ	ASCII variant is neutral, binary variant is proprietary
FBX	Proprietary
COLLADA	Neutral
3DS	Proprietary
IGES	Neutral
STEP	Neutral
VRML/X3D	Neutral

The four key features a 3D file can store include the model's geometry, the model's surface texture, scene details, and animation of the model⁷.

Specifically, they can store details about four key features of a 3D model, though it's worth bearing in mind that you may not always take advantage of all four features in all projects, and not all file formats support all four features!

3D printer applications do not to support animation. CAD and CAM such as designing airplane does not need feature of scene details.

DAE (Collada) appeared in the video animation above. Collada files belong to a neutral format used heavily in the video game and film industries. It's managed by the non-profit technology consortium, the Khronos Group.

The file extension for the Collada format is .dae. The Collada format stores data using the XML mark-up language.

The original intention behind the Collada format was to become a standard among 3D file formats. Indeed, in 2013, it was adopted by ISO as a publicly available specification, ISO/PAS 17506. As a result, many 3D modeling programs support the Collada format.

That said, the consensus is that the Collada format hasn't kept up with the times. It was once used heavily as an interchange format for Autodesk Max/Maya in film production, but the industry has now shifted more towards OBJ, FBX, and Alembic⁷.

- <https://registry.khronos.org/OpenGL-Refpages/>
- <https://www.mesa3d.org>
- <https://www.opengl.org/sdk/>, <https://www.opengl.org/sdk/libs/>

⁷ <https://all3dp.com/3d-file-format-3d-files-3d-printer-3d-cad-vrml-stl-obj/>

GRAPHICS HW AND SW STACK

- *HW Block Diagram*
- *SW Stack and Data Flow*
- *Displaying Pixels*
- *The Role and Purpose of Shaders*

This section provides a more detailed illustration of animation across the software and hardware stacks on both the CPU and GPU, and explains how data flows between the CPU, the GPU, and each layer of the software stack.

The previous section, *3D Modeling*, described the information stored in 3D models and how it is used for animation.

The following section, *SW Stack and Data Flow*, describes **how each frame is generated** to display **motion and skinning effects** using the small animation parameters stored in a 3D model and sent by the CPU.

The following section, *Role and Purpose of Shaders*, explains how different visual effects can be achieved by **switching shaders** and applying different materials across frames.

Reference:

- https://en.wikipedia.org/wiki/Free_and_open-source_graphics_device_driver

3.1 HW Block Diagram

The block diagram of the Graphic Processing Unit (GPU) is shown in Fig. 3.1.

The roles of the CPU and GPU in graphic animation are illustrated in Fig. 3.2.

- A GPU cannot directly read user input from a keyboard, mouse, or gamepad; play audio; or load files from a drive. Therefore, the GPU cannot handle all animation-related work on its own³.
- A graphics driver consists of an implementation of the OpenGL state machine and a compilation stack to compile the shaders into the GPU's machine language. This compilation, like most driver work, is executed on the CPU. The compiled shaders are then sent to and executed by the GPU. (SDL = Simple DirectMedia Layer)⁴.

The GPU driver writes commands and data from the CPU to GPU-visible memory over PCIe. On Arm GPUs, these commands form a Command Stream Frontend (CSF). A GPU chipset contains many SIMD processors (cores). To accelerate command processing, the CSF is represented in a compact form. As shown in Fig. 3.3, the chipset includes

¹ https://en.wikipedia.org/wiki/Graphics_processing_unit

² <https://en.wikipedia.org/wiki/Vulkan>

³ <https://stackoverflow.com/questions/47426655/cpu-and-gpu-in-3d-game-whos-doing-what>

⁴ [<https://en.wikipedia.org/wiki/Mesa_\(computer_graphics\)>](https://en.wikipedia.org/wiki/Mesa_(computer_graphics))

⁵ <https://developer.arm.com/documentation/102813/0107/GPU-activity>




Fig. 3.1: Components of a GPU: GPU has accelerated video decoding and encoding^{Page 21, 1}




Fig. 3.2: OpenGL and Vulkan are both rendering APIs. In both cases, the GPU executes shaders, while the CPU executes everything else^{Page 21, 2}.




Fig. 3.3: An MCU and specialized hardware circuits accelerate CSF (Command Stream Frontend) processing [Page 21, 5](#).

an MCU (microcontroller) and specialized hardware that translates the CSF into data structures for individual SIMD processors.

3.2 SW Stack and Data Flow

The driver runs on the CPU side as shown in Fig. 3.4. The OpenGL API eventually calls the driver's functions, and the driver executes these functions by issuing commands to the GPU hardware and/or sending data to the GPU.

Even so, the GPU's rendering work, which uses data such as 3D vertices and colors sent from the CPU and stored in GPU or shared memory, consumes more computing power than the CPU.




Fig. 3.4: Graphic SW Stack and data flow in initializing graphic model

✓ As described in the [Animation](#) section and Fig. 3.4, the **game engine's built-in C++ renderer handles OpenGL, Vulkan, or Metal calls automatically**. Users set values such as speed and velocity and customize the animation logic.

After the user creates a skeleton and textures for each model and sets keyframe times using a 3D modeling tool, users can write **gameplay scripts** (Java code, C#, Blueprints, GDScript, Python, etc.) to tell the engine to play animations⁶.

As described in the [Node-Editor \(shaders generator\)](#) section, graphics designers and technical artists create skin materials, while software and graphics programmers create additional materials with node-editor tools (shader generators). These tools generate the shaders used for rendering.

Shaders may call built-in functions written in Compute Shaders, SPIR-V, or LLVM-IR. LLVM *libclc* is a project for OpenCL built-in functions, which can also be used in OpenGL⁷. Like CPU built-ins, new GPU ISAs or architectures must implement their own built-ins or port them from open source projects like *libclc*.

The CPU evaluates the stored keyframes to produce motion and other animations, as illustrated in the [Animation](#) section.

The per-frame data is not the full set of vertices, but rather a small set of animation parameters called **uniform updates**, as shown in Fig. 3.5 and described later.

Note

Bone matrices determine the positions of triangles within a 3D model during animation. This bone transformation data is much **smaller** than the complete mesh of the 3D model. We will provide an example and explain this in more detail in the [Animation Example](#) section. Because this transformation data is small and constant across all

⁶ https://en.wikipedia.org/wiki/Java_OpenGL

⁷ <https://libclc.llvm.org>




Fig. 3.5: Graphic SW Stack and data flow in rendering

shader pipeline stages, it is stored in the GPU's global memory and can be cached in the **uniform/constant cache** for performance, as illustrated in Fig. 6.11 of *Processor Units and Memory Hierarchy in NVIDIA GPU 15* section.

The CPU updates only these small animation parameters and issues draw commands to the GPU.




Fig. 3.6: The input and output for GPU rendering

The 3D rendering pipeline is illustrated in Fig. 3.6.

Object-shape data is described by VAOs (Vertex Array Objects) in OpenGL. This is explained later in the *OpenGL section*. Additionally, OpenGL provides VBOs (Vertex Buffer Objects), which allow vertex array data to be stored in high-performance graphics memory on the GPU and enable efficient data transfer⁸⁹.

After the GPU receives the **uniform updates** from the CPU, it performs the computationally intensive per-vertex work within the rendering pipeline to generate the **final pixel values** for **each frame** displayed on the screen. These final pixel values are collectively referred to as the **Rendered Image**.

⁸ http://www.songho.ca/opengl/gl_vbo.html

⁹ If your models will be rigid, meaning you will not change each vertex individually, and you will render many frames with the same model, you will achieve the best performance not by storing the models in your class, but in vertex buffer objects (VBOs) <https://gamedev.stackexchange.com/questions/19560/what-is-the-best-way-to-store-meshes-or-3d-models-in-a-class>

✓ “**Rendered Image**” = the final per-frame output written into the framebuffer. The **Uniform Updates** will be described in detail later.

✓ The CPU updates only the small animation parameters called **uniform updates**, as shown in Fig. 3.5; the GPU performs the computationally intensive per-vertex work.

As mentioned in the previous section on *animation movement*, 3D modeling tools store keyframes, bone transforms at each keyframe, and related data, then use this information to perform animation.

The CPU updates only the **bone** transformation data ..., rather than updating the entire vertex or mesh data for each animation frame. These updates are very small—on the order of kilobytes rather than megabytes. For each rendered frame, the CPU sends these small updates to the GPU, and the **GPU takes over the animation work from the CPU**. This type of movement animation is called **skinning**, and is illustrated as follows:

Skinning

Skinning is a vertex deformation technique used to animate a mesh by attaching its vertices to a hierarchical skeleton (bones). Each vertex stores one or more bone indices and corresponding weights that describe how strongly each bone influences that vertex.

During animation, the application updates the bone transformation matrices. The vertex shader then computes the final vertex position by blending the transformed positions according to the stored weights. This allows the mesh to bend, twist, and deform smoothly as the skeleton moves.

Skinning does not create new geometry or smooth the surface topology. It only transforms the existing vertices of the mesh. Examples include bending an arm, flapping a wing, or deforming a flexible tube as its bones rotate.

The CPU updates only high-level animation state, such as:

- Current animation time
- **Bone matrices (small)**
- **Morph weights**
- Material parameters
- Particle emitter settings
- Global uniforms (camera, lights, etc.)

These are tiny updates —kilobytes, not megabytes.

$$finalPosition = \sum_{i=0}^{N-1} weight_i (boneMatrix_i \cdot originalPosition)$$

In practice, game engines normalize the weights so that their sum is 1.0: $\Rightarrow \sum_{i=0}^{N-1} weight_i = 1.0$

Example: Bending an Arm

Imagine a character’s arm mesh. Each vertex in the elbow area has weights like:

- 70% influenced by upper-arm bone
- 30% influenced by lower-arm bone

When the elbow bends:

- Upper-arm bone rotates
- Lower-arm bone rotates
- GPU blends the influence
- The elbow area deforms smoothly

This is skinning.

✓ After GPU processing, the color pixels are written to the framebuffer (video memory). The display device (monitor, LCD, OLED, etc.) fetches these pixels and displays them on the screen. The interface between the framebuffer and the display device is explained in the next section, *Displaying Pixels*.

3.3 Displaying Pixels

The interface between the framebuffer and the display device is shown in Fig. 3.7.




Fig. 3.7: PC with Frame Buffer to LCD Display

The GPU and display (monitor, LCD, OLED, etc.) use **VSync**, **NVIDIA G-SYNC**, or **AMD FreeSync** to prevent **screen tearing**, as described below:




Fig. 3.8: VSync¹⁰

- Double Buffering

While the display is reading from the frame buffer to display the current frame, we might be updating its contents for the next frame (not necessarily in a raster-scan manner). This would result in so-called tearing, in which the screen shows parts of the old frame and parts of the new frame. Double buffering resolves this issue. Instead of using a single framebuffer, a modern GPU uses a front buffer and a back buffer. The display reads from the front buffer while the GPU writes the next frame to the back buffer. When the frame is complete, the GPU swaps the front and back buffers (a buffer swap or page flip).

- VSync

A refresh period is roughly:

visible scan-out → blanking interval (including VSync) → next visible scan-out

- Visible scan-out: the display reads pixel rows from top to bottom and shows them.
- Blanking interval: a short time after the last row and before the next first row. Historically the electron beam returned to the top; modern LCD/LED displays retain a timing interval for coordination.
- **VSync: a timing signal within that blanking interval indicating the boundary between refreshes.**

vertical blanking interval = front porch + VSync pulse + back porch

- Front porch: brief delay after the last visible row finishes.
- VSync pulse: timing signal that marks synchronization between refreshes.
- Back porch: brief delay after the VSync pulse before the first visible row of the next frame starts.

In most display modes, the **visible scan-out becomes slower** because the **pixel clock is lower**. The blanking/VSync interval may also change somewhat, depending on the timing mode, but it is usually a small part of the total period.

Changing from a 5-cycle refresh period to a 10-cycle refresh period doubles the time for the entire refresh interval:

active scan-out + blanking/VSync interval = 10 cycles

- 5 cycles per refresh → VSync occurs every 5 cycles.
- 10 cycles per refresh → VSync occurs every 10 cycles.

If the refresh period is changed from 5 cycles to 10 cycles on an LCD or LED display, the display scans the screen more slowly: the active scan plus its blanking interval portion together occupy the 10 cycles.

¹⁰ https://www3.ntu.edu.sg/home/ehchua/programming/opengl/CG_BasicsTheory.html

Double buffering alone does not solve the entire problem, as the buffer swap might occur at an inappropriate time, for example, while the display is in the middle of displaying the old frame. This is resolved via the so-called vertical synchronization (or VSync) at the end of the raster-scan. When the application signals the GPU to perform a buffer swap, the GPU waits until the next VSync to perform the actual swap, after the entire current frame is displayed.

Tearing

To avoid tearing, the GPU runs at half the refresh rate of the display, as shown below.

```
GPU      | GPU is writing buffer A | GPU is writing buffer B |
Display  | VSync B |
Display  | VSync B and stay on B | VSync A and stay on A |
                                   ^
                                   |
                                   tearing
```

As the text diagram above illustrates, enabling VSync prevents the display from presenting frames faster than its refresh rate. If the GPU can render frames faster than the display refresh rate, the additional frames are not presented; VSync still prevents tearing.

If an application presents frames with VSync enabled at a fixed cadence, its displayed-frame cadence may be an integral fraction of the display refresh rate: 1/1, 1/2, 1/3, 1/4, and so on.

Examples on a 60 Hz display:

- GPU/app $\geq$ 60 FPS: 1/1 $\rightarrow$ 60 presented FPS (extra rendered frames may be dropped or replaced).
- GPU/app = 30 FPS: 1/2 $\rightarrow$ 30 presented FPS.
- GPU/app = 20 FPS: 1/3 $\rightarrow$ 20 presented FPS.
- GPU/app = 15 FPS: 1/4 $\rightarrow$ 15 presented FPS.

- NVIDIA G-SYNC and AMD FreeSync

If your monitor and graphics card both in your customer computer support NVIDIA G-SYNC, you're in luck. With this technology, a special chip in the display communicates with the graphics card. This lets the monitor vary the refresh rate to match the frame rate of the NVIDIA GTX graphics card, up to the maximum refresh rate of the display. This means that the frames are displayed as soon as they are rendered by the GPU, eliminating screen tearing and reducing stutter for when the frame rate is both higher and lower than the refresh rate of the display. This makes it perfect for situations where the frame rate varies, which happens a lot when gaming. Today, you can even find G-SYNC technology in gaming laptops!

AMD has a similar solution called FreeSync. However, this doesn't require a proprietary chip in the monitor. In FreeSync, the AMD Radeon driver, and the display firmware handle the communication. Generally, FreeSync monitors are less expensive than their G-SYNC counterparts, but gamers generally prefer G-SYNC over FreeSync as the latter may cause ghosting, where old images leave behind artifacts¹¹.

3.4 The Role and Purpose of Shaders

The flow for 3D/2D graphic processing is shown in Fig. 3.9.

Compiled shaders are sent to the GPU when the application calls `glLinkProgram()`. At that point, the driver prepares the shader binaries for GPU execution.

¹¹ <https://www.avadirect.com/blog/frame-rate-fps-vs-hz-refresh-rate/>




Fig. 3.9: OpenGL Flow

`glLinkProgram()` is called when an application finishes preparing a shader program —**not when creating a mesh, and not when issuing a draw command.**

When a game calls `glLinkProgram()` to relink a shader program, the shaders must be compiled and prepared for the GPU.

This usually happens during game startup, level loading, or the creation of a new shader variant (e.g., enabling shadows, fog, skinning).

Games switch shaders constantly —sometimes hundreds of times per frame—but they do not re-link them.

When a game runs, different materials, effects, and rendering passes use different shaders.

Examples of switching shaders:

- When the player enters a snowy biome, ice meshes use the ice shader.
- The axe blade uses a metal PBR shader. When the axe blade strikes stone, a particle shader renders the sparks.

BASIC GEOMETRY IN COMPUTER GRAPHICS

- *Color*
- *Transformation*
- *Cross Product*
- *Dot Product*
- *Projection*

This section introduces the fundamental geometry mathematics used in computer graphics. As discussed in the previous sections, 3D animation is primarily based on geometric representations such as meshes (vertices) and surface descriptions including textures, materials, shaders, and lighting models created in 3D content creation tools. Consequently, vertex transformations and lighting-based color computations form the mathematical foundation of modern computer graphics and animation.

The complete concept can be found in the book *Computer Graphics: Principles and Practice, 3rd Edition*, authored by John F. et al. However, the book contains over a thousand pages.

It is very comprehensive and may take considerable time to understand all the details.

4.1 Color

- Additive colors in light are shown in Fig. 4.1¹².
- In the case of paints, additive colors produce shades and become light gray due to the addition of darker pigments³.

Note

Additive colors

I know it doesn't match human intuition. However, additive RGB colors in light combine to produce white light, while additive RGB in paints result in light gray paint. This makes sense because light has no shade. This result stems from the way human eyes perceive color. Without light, no color can be sensed by the eyes.

Computer engineers should understand that exploring the underlying reasons falls into the realms of physics or the biology of the human eye structure.

¹ https://en.wikipedia.org/wiki/RGB_color_model

² https://www.youtube.com/watch?v=kEnz_3miiAc

³ <https://www.tiktok.com/@tonesterpaints/video/7059565281227853102>




Fig. 4.1: Additive colors in light

4.2 Transformation

Overview

Transformation matrices are taught in high school and college. However, these mathematical details are not always retained clearly in memory. The following section reviews the parts relevant to graphics rendering.

In both 2D and 3D graphics, every object transformation is performed by multiplying the object's vertex coordinates by one or more **transformation matrices**. Modern OpenGL uses **homogeneous coordinates** and **4×4 matrices** to unify translation, rotation, scaling, projection, and even animation (skinning) into a single mathematical framework.

A vertex in 3D is represented as:

$$\mathbf{v} = \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

A transformation is applied by matrix multiplication:

$$\mathbf{v}' = M\mathbf{v}$$

Multiple transformations are combined by multiplying matrices:

$$\mathbf{v}' = PVM\mathbf{v}$$


Where:

- M = Model matrix (object $\rightarrow$ world)
- V = View matrix (world $\rightarrow$ camera)
- P = Projection matrix (camera $\rightarrow$ clip space)

This is the core of the OpenGL rendering pipeline, as shown in Fig. 4.2.

- **Model space:** This is the vertex-position space described in *Root bone in Animation flow*. All vertex coordinates are calculated relative to the root bone.
- **Model Transform:** M = Model matrix (object $\rightarrow$ world). This represents the vertex positions mentioned under *Transform Animation in Animation flow*.

⁴ https://www3.ntu.edu.sg/home/ehchua/programming/opengl/CG_BasicsTheory.html

Fig. 4.2: Coordinate Transformation Pipeline^{Page 32, 4}

Details for Fig. 4.2 can be found on “4. Vertex Processing” of the website^{Page 32, 4}.

Transformation Matrices⁵

- Translation: Moves an object in 3D space.

$$T(x, y, z) = \begin{bmatrix} 1 & 0 & 0 & x \\ 0 & 1 & 0 & y \\ 0 & 0 & 1 & z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} = \begin{bmatrix} X + x \\ Y + y \\ Z + z \\ 1 \end{bmatrix}$$

- Scaling: Resizes an object.

$$S(s_x, s_y, s_z) = \begin{bmatrix} s_x & 0 & 0 & 0 \\ 0 & s_y & 0 & 0 \\ 0 & 0 & s_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} = \begin{bmatrix} s_x X \\ s_y Y \\ s_z Z \\ 1 \end{bmatrix}$$

- Rotation X: Rotates around the X-axis.

$$R_x(\theta) = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & \cos \theta & -\sin \theta & 0 \\ 0 & \sin \theta & \cos \theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} = \begin{bmatrix} X \\ Y \cos \theta - Z \sin \theta \\ Y \sin \theta + Z \cos \theta \\ 1 \end{bmatrix}$$

- Rotation Y: Rotates around the Y-axis.

$$R_y(\theta) = \begin{bmatrix} \cos \theta & 0 & \sin \theta & 0 \\ 0 & 1 & 0 & 0 \\ -\sin \theta & 0 & \cos \theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} = \begin{bmatrix} X \cos \theta + Z \sin \theta \\ Y \\ -X \sin \theta + Z \cos \theta \\ 1 \end{bmatrix}$$

- Rotation Z: Rotates around the Z-axis.

$$R_z(\theta) = \begin{bmatrix} \cos \theta & -\sin \theta & 0 & 0 \\ \sin \theta & \cos \theta & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} = \begin{bmatrix} X \cos \theta - Y \sin \theta \\ X \sin \theta + Y \cos \theta \\ Z \\ 1 \end{bmatrix}$$

- Shear in X: - Skews geometry along axis X.

$$\text{Shear}_X(a, b) = \begin{bmatrix} 1 & a & b & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} = \begin{bmatrix} X + aY + bZ \\ Y \\ Z \\ 1 \end{bmatrix}$$

- Shear in Y: Skews geometry along axis Y.

⁵ https://en.wikipedia.org/wiki/Transformation_matrix

$$\text{Shear}_Y(c, d) = \begin{bmatrix} 1 & 0 & 0 & 0 \\ c & 1 & d & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} = \begin{bmatrix} X \\ cX + Y + dZ \\ Z \\ 1 \end{bmatrix}$$

- Shear in Z: Skews geometry along axis Z.

$$\text{Shear}_Z(e, f) = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ e & f & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} = \begin{bmatrix} X \\ Y \\ eX + fY + Z \\ 1 \end{bmatrix}$$

- Reflection: Mirrors across a plane.

Reflect_{XY} , $\text{Reflect}(\mathbf{n})$

The “4.2 Model Transform (or Local Transform, or World Transform)” of on the website^{Page 32, 4} provides conceptual coverage of transformations. List the websites that provide proofs of the non-obvious transformation formulas below.

Rotation

The mathematical proof is given below.

1. https://en.wikipedia.org/wiki/Rotation_matrix

- Prove the 2D formula and then intuitively extend it to 3D along the X, Y, and Z axes⁶.

2. Proof in greater details:

https://austinmorlan.com/posts/rotation_matrices/

Shear (Skew)

Shear is a skewing transformation as shown in Fig. 4.3.




Fig. 4.3: 3D shear

Shear in X: plane $x = 0$ (the YZ-plane), slides points parallel to the X-axis.

The mathematical proof is given below.

https://en.wikipedia.org/wiki/Shear_mapping

Reflection

Reflection is nothing but a mirror image of an object.

⁶ https://en.wikipedia.org/wiki/Rotation_matrix

Reflection across the XY-plane:

$$\text{Reflect}_{XY} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & -1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Reflection across an arbitrary plane with unit normal $\mathbf{n}$:

$$R = I - 2\mathbf{nn}^T$$

The mathematical proof is given below.

<https://www.geeksforgeeks.org/computer-graphics/computer-graphics-reflection-transformation-in-3d/>

The following Quaternion Product (Hamilton product) is from the wiki⁷ since it is not covered in the book.

$$\mathbf{i}j = -ji = k, jk = -kj = i, ki = -ik = j.$$

4.3 Cross Product

Todo:

- Computing the direction of the line of intersection between two planes (via $n_1 \times n_2$)

end-of-Todo

Both triangles and quads are polygons. So, objects can be formed with polygons in both 2D and 3D. The transformation in 2D or 3D is well covered in almost every computer graphics book. This section introduces the most important concept and method for **determining inner and outer planes**. Then, a **point or object can be checked for visibility** during 2D or 3D rendering.

Any **area** of a polygon can be calculated by dividing it into triangles or quads. The area of a triangle or quad can be calculated using the cross product in 3D.

✓ The role of cross product:

In 2D geometry mathematics, v_0, v_1 and v_2 can form the area of a parallelogram as shown in Fig. 4.4. The fourth vertex, v_3 , can then be determined to complete the parallelogram.

The area of the parallelogram is given by:

$$\mathbf{a} = v_1 - v_0, \mathbf{b} = v_2 - v_0 \\ \|\mathbf{a} \times \mathbf{b}\| = \|\mathbf{a}\| \|\mathbf{b}\| \sin(\Theta)$$

The area of a parallelogram is same in both 2D and 3D. To extend the definition of the cross product to 3D, all we must additionally consider the orientation of the plane, since a plane has two possible faces.

$$\mathbf{a} \times \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \sin(\Theta) \mathbf{n}$$

- $\mathbf{n}$ is a unit vector perpendicular to the plane. $\Rightarrow$ direction.

As shown in Fig. 4.5, the plane determined by v_0, v_1, v_2 with CCW ordering defines a unique orientation.

The area of the parallelogram remains unchanged after rotation as shown in Fig. 4.6, which means the area and plane face determined by extending the definition of cross product from 2D to 3D correctly.

⁷ <https://en.wikipedia.org/wiki/Quaternion>




Fig. 4.4: The area determined by v_0, v_1, v_2 in 2D




Fig. 4.5: The area and plane face determined by v_0, v_1, v_2 with CCW ordering before rotation z axis.




Fig. 4.6: The area and plane face determined by v_0, v_1, v_2 with CCW ordering after rotation z axis.

The area of the triangle is obtained by dividing the parallelogram by 2:

$$\frac{1}{2} \|\mathbf{a} \times \mathbf{b}\| \quad \dots \text{ (triangle area)}$$

✓ Matrix Notation for Cross Product:

The cross product in **2D** is defined by a formula and can be represented with matrix notation, as proven here¹⁰¹¹.

The cross product in **2D** is defined by a formula and can be represented with matrix notation, as proven here¹⁰¹¹.

$$\mathbf{a} \times \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \sin(\Theta)$$

$$\mathbf{a} \times \mathbf{b} = \begin{vmatrix} \mathbf{i} & \mathbf{j} & \mathbf{k} \\ a_1 & a_2 & 0 \\ b_1 & b_2 & 0 \end{vmatrix} = \begin{bmatrix} a_1 & a_2 \\ b_1 & b_2 \end{bmatrix}$$

¹⁰ <https://www.xarg.org/book/linear-algebra/2d-perp-product/>

¹¹ <https://www.nagwa.com/en/explainers/175169159270/>

After the above matrix form is proven, the antisymmetry property may be demonstrated as follows:

$$a \times b = \begin{bmatrix} a \\ b \end{bmatrix} = \begin{bmatrix} a_1 & a_2 \\ b_1 & b_2 \end{bmatrix} = a_1 b_2 - a_2 b_1 =$$

$$-b_1 a_2 - (-b_2 a_1) = \begin{bmatrix} -b_1 & -b_2 \\ a_1 & a_2 \end{bmatrix} = \begin{bmatrix} -b \\ a \end{bmatrix} = -b \times a$$

✓ Determine the area in a plane:

As described earlier in this section, three vertices form a parallelogram or triangle. Its area can be determined because the angle between $v_1 - v_0$ and $v_2 - v_1$ satisfies $0 < \Theta < 180^\circ$ under CCW orientation. In fact **one single vector** $v_1 - v_0$ **is sufficient** to determine the area. This is described below.

In 2D, any two points from P_i to P_{i+1} can form a vector and determine the inner or outer side.

For example, as shown in Fig. 4.7, Θ is the angle from $P_i P_{i+1}$ to $P_i P'_{i+1} = 180^\circ$.

Using the right-hand rule and counter-clockwise order, any vector $P_i Q$ between $P_i P_{i+1}$ and $P_i P'_{i+1}$, with angle θ such that $0^\circ < \theta < 180^\circ$, indicates the inward direction.



Fig. 4.7: Inward edge normals




Fig. 4.8: Inward and outward in 2D for a vector.

Based on this observation, the rule for inward and outward vectors is shown in Fig. 4.7. Facing the same direction as a specific vector, the left side is inward and the right side is outward, as shown in Fig. 4.8.

For each edge $P_i - P_{i+1}$, the inward edge normal is the vector $\times v_i$; the outward edge normal is $-\times v_i$, where $\times v_i$ is the cross-product of v_i , as shown in Fig. 4.7.

A polygon can be created from a set of vertices. Suppose $(P_0, P_1, \dots, P_n)$ defines a polygon. The line segments $P_0 P_1, P_1 P_2$, etc., are the polygon's edges. The vectors $v_0 = P_1 - P_0, v_1 = P_2 - P_1, \dots, v_n = P_0 - P_n$ represent those edges.

Using counter-clockwise ordering, the left side is considered inward. Thus, the inward region of a polygon can be determined, as shown in Fig. 4.9 and Fig. 4.10.




Fig. 4.9: Triangle with CCW




Fig. 4.10: Hexagon with CCW




Fig. 4.11: Triangle with CW

For a convex polygon with vertices listed in counter-clockwise order, the inward edge normals point toward the interior of the polygon, and the outward edge normals point toward the unbounded exterior. This matches our usual intuition.

However, if the polygon vertices are listed in clockwise (CW) order, the interior and exterior definitions are reversed. Fig. 4.11 shows an example where P_0, P_1, P_2 are arranged in CW order.

This cross product has an important property: going from v to $\times v$ involves a 90° rotation in the same direction as the rotation from the positive x-axis to the positive y-axis.




Fig. 4.12: Draw a polygon with vertices counter clockwise

As shown in Fig. 4.12, when drawing a polygon with vectors (lines) in counter-clockwise order, the polygon will be formed, and the two sides of each vector (line) can be identified¹².

Furthermore, whether a point is inside or outside the polygon can be determined.

One simple method to test whether a point lies inside or outside a simple polygon is to cast a ray from the point in any fixed direction and count how many times it intersects the edges of the polygon.

If the point is outside the polygon, the ray will intersect its edges an even number of times. If the point is inside the polygon, it will intersect the edges an odd number of times¹³.

In the same way, by following the counter-clockwise direction to create a 2D polygon step by step, a 3D polygon can be constructed.

¹² Figure 7.19 of Book: Computer graphics principles and practice 3rd edition

¹³ https://en.wikipedia.org/wiki/Point_in_polygon





Fig. 4.13: Cross product definition in 3D

As shown in Fig. 4.13 from the wiki⁹, the inward direction is determined by $a \times b < 0$, and the outward direction is determined by $a \times b > 0$ in OpenGL.

Replacing a and b with x and y , as shown in Fig. 4.14, the positive Z-axis ($z+$) represents the outer surface, while the negative Z-axis ($z-$) represents the inner surface¹⁴.

Fig. 4.14: OpenGL pointing outwards, indicating the outer surface (z axis is $+$)

Position each triangle in front of the camera and construct it with counter-clockwise (CCW) ordering, as shown in Fig. 4.9. Building every triangle with CCW ordering defines a consistent outer surface (front face). Fig. 4.15 shows a 3D polygon created from 2D triangles. The direction of a triangle's plane is given by the line perpendicular to it.

Cast a ray from the 3D point along the X-axis and count how many intersections with the outer object occur. Depending on the number of intersections along each axis (even or odd), you can determine whether **the point (or camera) is inside or outside**¹⁵.

An odd number means inside, and an even number means outside. As shown in Fig. 4.16, points on the line passing through the object satisfy this rule.

⁹ https://en.wikipedia.org/wiki/Cross_product

¹⁴ Normals are used to differentiate the front- and back-face, and for other processing such as lighting. Right-hand rule (or counter-clockwise) is used in OpenGL. The normal is pointing outwards, indicating the outer surface (or front-face). https://www3.ntu.edu.sg/home/ehchua/programming/opengl/CG_BasicsTheory.html

¹⁵ <https://stackoverflow.com/questions/63557043/how-to-determine-whether-a-point-is-inside-or-outside-a-3d-model-computationally>




Fig. 4.15: 3D polygon with directions on each plane




Fig. 4.16: Point is inside or outside of 3D object

✓ Summary:

Based on these description of this section, this means:

- ✓ Each mesh (triangle or primitive) has a fixed “outer” and “inner” side, determined by CCW ordering in object space.
- ✓ By reading these CCW-ordered vertices in a triangle sequentially, the shape and surface orientation of the 3D model can be constructed.
- ✓ There is no need to wait for the entire mesh to be received; once three CCW-ordered vertices in a triangle are available, each triangle can be processed correctly as shown in Fig. 4.17 from the camera position p_0 .
- ✓ When the camera moves to the p_1 inside an object: CCW $\leftrightarrow$ CW flips as shown in Fig. 4.17.
- ✓ As shown in *Triangle Area Calculation*, when $0 < \Theta < 180^\circ$ under CCW orientation, the area of a triangle area is given by:

$$\frac{1}{2} \|(v_1 - v_0) \times (v_2 - v_0)\| = \frac{1}{2} \|(v_1 - v_0)\| \|(v_2 - v_0)\| \sin(\Theta)$$

- ✓ Though each triangle can be correctly identified and processed using its CCW ordering. As mentioned in Fig. 4.2 of section *Transformation*, the Coordinate Transformation Pipeline maps geometry from camera space to Clipping Space (Clipping Volume). This transformation significantly simplifies the calculations required to discard and clip triangles, as described in the section *Projection* later.

How does OpenGL render (draw) the inner face of a triangle?

OpenGL does NOT determine front/back in world space.

When the camera moves to the inner space of a object:

- The projection changes
- The triangle’s screen-space orientation changes




Fig. 4.17: A triangle can be constructed as soon as three vertices are received

- CCW $\leftrightarrow$ CW flips
- So the GPU flips front/back classification

OpenGL uses counter clockwise and pointing outwards as default

```
// unit cube
// A cube has 6 sides and each side has 4 vertices, therefore, the total number
// of vertices is 24 (6 sides * 4 verts), and 72 floats in the vertex array
// since each vertex has 3 components (x,y,z) (= 24 * 3)
//      v6----- v5
//      /|          /|
//      v1-----v0|
//      | |          | |
//      | v7-----|-v4
//      | /          | /
//      v2-----v3

// vertex position array
GLfloat vertices[] = {
    .5f, .5f, .5f, -.5f, .5f, .5f, -.5f, -.5f, .5f, .5f, -.5f, .5f, // v0,v1,v2,v3
    ↪ (front)
    .5f, .5f, .5f, .5f, -.5f, .5f, .5f, -.5f, -.5f, .5f, .5f, -.5f, // v0,v3,v4,v5
    ↪ (right)
    .5f, .5f, .5f, .5f, .5f, -.5f, -.5f, .5f, -.5f, -.5f, .5f, .5f, // v0,v5,v6,v1
    ↪ (top)
    -.5f, .5f, .5f, -.5f, .5f, -.5f, -.5f, -.5f, -.5f, -.5f, .5f, .5f, // v1,v6,v7,v2
    ↪ (left)
    -.5f, -.5f, -.5f, .5f, -.5f, -.5f, .5f, -.5f, .5f, -.5f, -.5f, .5f, // v7,v4,v3,v2
    ↪ (bottom)
    .5f, -.5f, -.5f, -.5f, -.5f, -.5f, -.5f, .5f, -.5f, .5f, .5f, -.5f, // v4,v7,v6,v5
    ↪ (back)
};
```

From the code above, we can see that OpenGL uses counter-clockwise¹⁶ and pointing outwards as the default. However, OpenGL provides `glFrontFace(GL_CW)` for clockwise winding¹⁷.

For a group of objects, a scene graph provides better animation support and saves memory¹⁸.

4.4 Dot Product

Dot Product

- Ray—plane (line—plane) intersection
- Determining angles between vectors
- Lighting (Lambertian shading)
- Solving for a point on the intersection line of two planes (because plane equations use dot products)

Described in wiki here:

https://en.wikipedia.org/wiki/Dot_product

¹⁶ http://www.songho.ca/opengl/gl_vbo.html

¹⁷ <https://registry.khronos.org/OpenGL-Refpages/gl4/html/glFrontFace.xhtml>

¹⁸ https://en.wikipedia.org/wiki/Scene_graph

✓ As described in the previous section *Cross Product*, the cross-product is:

$$\mathbf{a} = v_1 - v_0, \mathbf{b} = v_2 - v_0$$

$$\mathbf{a} \times \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \sin(\Theta) \mathbf{n}$$

- $\mathbf{n}$ is a unit vector perpendicular to the plane $\Rightarrow$ direction.

The dot product definition is:

$$\mathbf{a} \cdot \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \cos(\Theta)$$

✓ Since $\mathbf{n}$ is the outward normal vector for a CCW-ordered triangle, we have:

- $(\mathbf{p} - v_0) \cdot \mathbf{n} > 0 \Rightarrow \mathbf{p}$ lies on the front (outer) side of the plane.
- $(\mathbf{p} - v_0) \cdot \mathbf{n} = 0 \Rightarrow \mathbf{p}$ lies on the plane.
- $(\mathbf{p} - v_0) \cdot \mathbf{n} < 0 \Rightarrow \mathbf{p}$ lies on the back (inner) side of the plane.

✓ A plane is represented by:

$$\mathbf{n} \cdot (\mathbf{x}_1 - \mathbf{x}_0) = 0$$

where:

- $\mathbf{n}$ is the plane's normal vector
- x_0, x_1 are any points on the plane

$$\mathbf{n} \cdot (\mathbf{x}_1 - \mathbf{x}_0) = 0$$

$$\Rightarrow \mathbf{n} \cdot \mathbf{x}_1 - \mathbf{n} \cdot \mathbf{x}_0 = 0$$

$$\Rightarrow \mathbf{n} \cdot \mathbf{x}_1 = \mathbf{n} \cdot \mathbf{x}_0$$

Let's define the scalar constant d by:

$$d = -\mathbf{n} \cdot \mathbf{x}_0$$

Thus, the set of all points $\mathbf{p}$ satisfying

$$\mathbf{n} \cdot \mathbf{p} + d = 0$$

✓ Ray-plane (line-plane) intersection

For an edge between vertices $\mathbf{p}_0$ and $\mathbf{p}_1$, parameterized as:

$$\mathbf{p}(t) = \mathbf{p}_0 + t(\mathbf{p}_1 - \mathbf{p}_0)$$

the intersection with a clipping plane is found by solving:

$$\mathbf{n} \cdot \mathbf{p}(t) + d = 0$$

This yields:

$$t = \frac{-(\mathbf{n} \cdot \mathbf{p}_0 + d)}{\mathbf{n} \cdot (\mathbf{p}_1 - \mathbf{p}_0)}$$

4.5 Projection

✓ Reason:

As described in the previous section *Cross Product*, each mesh (triangle or primitive) has a fixed “outer” and “inner” side, determined by CCW ordering in object space. By reading these CCW-ordered vertices sequentially, the shape and surface orientation of the 3D model can be constructed, and hidden primitives can be clipped or discarded.

However primitive clipping and discarding can be performed much more efficiently by mapping the view frustum to **clip space**, where the GPU can **easily clip or discard primitives**, as shown Fig. 4.18 from the earlier section *Transformation* again for clarity. Performing clipping and discarding in **world space** would be significantly more **difficult**.




Fig. 4.18: Coordinate Transformation Pipeline^{Page 32, 4}

✓ Projection Area:




Fig. 4.19: Clipping-Volume Cuboid

Only objects within the cone between near and far planes are projected to 2D in perspective projection.

Perspective and orthographic projections (used in CAD tools) from 3D to 2D can be represented by transformation matrices as described in wiki here⁸.

The “4.4 Projection Transform - Perspective Projection” of on the website^{Page 32, 4} provides conceptual coverage of projections.

Camera Space Setup

Assume a right-handed camera coordinate system as shown in Fig. 4.19:

- The camera is located at the origin.
- The camera looks down the negative z axis.
- The near plane is located at $z = -n$.
- The far plane is located at $z = -f$.
- The view frustum bounds on the near plane are:
 - left: l
 - right: r

⁸ https://en.wikipedia.org/wiki/3D_projection

- bottom: b
- top: t

✓ Matrix P_{persp} converts 3D coordinates into clip space:

Perspective projection P_{persp} (general form): Converts 3D $\rightarrow$ clip space with depth

$$P = \begin{bmatrix} \frac{2n}{r-l} & 0 & \frac{r+l}{r-l} & 0 \\ 0 & \frac{2n}{t-b} & \frac{t+b}{t-b} & 0 \\ 0 & 0 & -\frac{f+n}{f-n} & -\frac{2fn}{f-n} \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

A point p in camera space is represented as:

$$\mathbf{p} = \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Converting from camera space to clipping space produces a homogeneous coordinate of the form $[x_c, y_c, z_c, w_c]$:

$$P\mathbf{p} = \begin{bmatrix} \frac{2n}{r-l}x + \frac{r+l}{r-l}z \\ \frac{2n}{t-b}y + \frac{t+b}{t-b}z \\ -\frac{f+n}{f-n}z - \frac{2fn}{f-n} \\ -z \end{bmatrix} = \begin{bmatrix} x_c \\ y_c \\ z_c \\ w_c \end{bmatrix}$$

After transforming to **clip space**, each vertex coordinate is expressed in **homogeneous** form, and the **view frustum boundaries are encoded in the coordinate values**. A vertex lies inside the view frustum if the following conditions are satisfied:

$$-w_c \leq x_c \leq w_c \quad -w_c \leq y_c \leq w_c \quad -w_c \leq z_c \leq w_c$$

✓ Matrix P_{persp} Derivation:

★ **Idea:**

The perspective projection matrix P_{persp} is derived by choosing its coefficients such that, after perspective division, the frustum boundaries l, r, b, t, n, f are mapped to the **canonical cube** $[-1, 1]^3$.

More explicitly, we impose the following constraints:

$$\begin{aligned} x = l &\rightarrow x_{ndc} = -1 \\ x = r &\rightarrow x_{ndc} = 1 \\ y = b &\rightarrow y_{ndc} = -1 \\ y = t &\rightarrow y_{ndc} = 1 \\ z = -n &\rightarrow z_{ndc} = -1 \\ z = -f &\rightarrow z_{ndc} = 1 \end{aligned}$$

These conditions determine the coefficients of the matrix.

X Coordinate Mapping

Since the near plane is located at $z = -n$, by similar triangles, the projected x-coordinate on the near plane is proportional to the ratio between $-n$ and z as follows:

$$x_n = \frac{-n}{z} x$$

The near-plane bounds map to NDC such that $x = l \Rightarrow x_{ndc} = -1$ and $x = r \Rightarrow x_{ndc} = +1$. Since the midpoint $\frac{l+r}{2}$ is generally not equal to 0, the mapping is not centered at the origin. Therefore, a linear mapping requires an offset term B_x as follows:

$$x_{ndc} = A_x x_n + B_x$$

Applying the near constraints: substituting $x_n = l$ and $x_{ndc} = -1$:

$$-1 = A_x l + B_x \quad \dots(1)$$

Applying the far constraints: substituting $x_n = r$ and $x_{ndc} = 1$:

$$1 = A_x r + B_x \quad \dots(2)$$

Solving equations (1) and (2) to get A_x :

$$2 = A_x(r - l) \Rightarrow A_x = \frac{2}{r - l} \quad \dots(3)$$

Substituting equation (3) to (2):

$$1 = A_x r + B_x \Rightarrow 1 = \frac{2}{r - l} r + B_x \Rightarrow B_x = \frac{r - l - 2r}{r - l} = -\frac{r + l}{r - l} \quad \dots(4)$$

From (3) and (4): Solving for A_x and B_x yields:

$$A_x = \frac{2}{r - l}, \quad B_x = -\frac{r + l}{r - l}$$

Substituting $x_n = \frac{-n}{z} x$ gives the resulting mapping is:

$$x_{ndc} = A_x x_n + B_x = \frac{-2n}{r - l} \frac{x}{z} - \frac{r + l}{r - l}$$

Since $x_{ndc} = \frac{x_c}{w_c}$ and $w_c = -z$, therefore:

$$x_c = x_{ndc}(-z) = \frac{2n}{r - l} x + \frac{r + l}{r - l} z.$$

The physical meaning of $x_c = x_{ndc}(-z)$ is:

Multiplying $x_c = x_{ndc}$ by $(-z)$ converts the normalized coordinate back into the **real horizontal position at that depth**.

Y Coordinate Mapping:

Using the same derivation for the y-axis:

$$y_n = \frac{-n}{z} y$$

The resulting mapping is:

$$A_y = \frac{2}{t - b}, \quad B_y = -\frac{t + b}{t - b}$$

y_{ndc} :

$$y_{ndc} = A_y y_n + B_y = \frac{-2n}{t - b} \frac{y}{z} - \frac{t + b}{t - b}$$

y_c :

$$y_c = y_{ndc}(-z) = \frac{2n}{t - b} y + \frac{t + b}{t - b} z.$$

Z Coordinate Mapping

Depth is mapped linearly such that:

$$z = -n \Rightarrow z_{ndc} = -1 \quad z = -f \Rightarrow z_{ndc} = +1$$

Assume:

$$z_c = A_z z + B_z$$

Then:

$$z_{ndc} = \frac{A_z z + B_z}{-z}$$

Applying the near constraints: substituting $z = -n$ and $z_{ndc} = -1$:

$$-1 = \frac{A_z(-n) + B_z}{-(-n)} \Rightarrow -n = A_z(-n) + B_z \quad \dots(1)$$

Applying the far constraints: substituting $z = -f$ and $z_{ndc} = 1$:

$$1 = \frac{A_z(-f) + B_z}{-(-f)} \Rightarrow f = A_z(-f) + B_z \quad \dots(2)$$

Solving equations (1) and (2) to get A_z :

$$-n - f = A_z(-n + f) \Rightarrow A_z = \frac{-n - f}{-n + f} = -\frac{f + n}{f - n} \quad \dots(3)$$

Substituting equation (3) to (2):

$$\begin{aligned} f &= A_z(-f) + B_z \Rightarrow f = -\frac{f + n}{f - n}(-f) + B_z \\ \Rightarrow B_z &= f + \frac{f + n}{f - n}(-f) = \frac{(f^2 - fn) + (-f^2 - fn)}{f - n} = -\frac{2fn}{f - n} \quad \dots(4) \end{aligned}$$

From (3) and (4):

$$A_z = -\frac{f + n}{f - n}, \quad B_z = -\frac{2fn}{f - n}$$

z_{ndc} :

$$z_{ndc} = \frac{A_z z + B_z}{-z} = \frac{-\frac{f+n}{f-n}z - \frac{2fn}{f-n}}{-z}$$

z_c :

$$z_c = z_{ndc}(-z) = -\frac{f + n}{f - n}z - \frac{2fn}{f - n}$$

Perspective Projection Matrix

Combining all components, the perspective projection matrix is:

$$P = \begin{bmatrix} \frac{2n}{r-l} & 0 & \frac{r+l}{r-l} & 0 \\ 0 & \frac{2n}{t-b} & \frac{t+b}{t-b} & 0 \\ 0 & 0 & -\frac{f+n}{f-n} & -\frac{2fn}{f-n} \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

A point p in camera space is represented as:

$$\mathbf{p} = \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Converting from camera space to clipping space produces a homogeneous coordinate of the form $[x_c, y_c, z_c, w_c]$:

$$P\mathbf{p} = \begin{bmatrix} x_c \\ y_c \\ z_c \\ w_c \end{bmatrix}$$

As mentioned the **physical meaning** of $x_c = x_{ndc}(-z)$, $y_c = y_{ndc}(-z)$, $z_c = z_{ndc}(-z)$ is:

Converts the normalized coordinate back into the **real horizontal position at that depth**.

If $x_c = x_{ndc}(-z_i)$, $y_c = y_{ndc}(-z_i)$, $z_c = z_{ndc}(-z_i)$, where z_i is the depth of plane P_i (a plane located between the near and far planes), then this operation projects NDC corrdinates onto **the plane P_i** .

For all vertices that survived clipping, the resulting coordinates satisfy:

$$-1 \leq x_{ndc} \leq 1 \quad -1 \leq y_{ndc} \leq 1 \quad -1 \leq z_{ndc} \leq 1$$

After transforming to **clip space**, each vertex corrodinate is expressed in **homogeneous** form, and the **view frustum boundaries are encoded in the coordinate values**. A vertex lies inside the view frustum if the following conditions are satisfied:

$$-w_c \leq x_c \leq w_c \quad -w_c \leq y_c \leq w_c \quad -w_c \leq z_c \leq w_c$$

Based on the clip-space representation, when w_c is $-z$ the coordinates $[x, y, z]$ **can be clipped according to their depth values**.

✓ Map points to NDC:

After applying matrix P_{persp} , any vertex in view space is mapped to clip space:

$$(x, y, z, 1) \rightarrow (x_c, y_c, z_c, w_c)$$

For the standard perspective matrix:

$$w_c = -z$$

The normalized device coordinates (NDC) are obtained by perspective division:

$$x_{ndc} = \frac{x_c}{w_c}, \quad y_{ndc} = \frac{y_c}{w_c}, \quad z_{ndc} = \frac{z_c}{w_c}$$

Substituting the matrix coefficients:

$$\begin{aligned} w_c &= -z \\ x_{ndc} &= \frac{x_c}{w_c} = \frac{\frac{2n}{r-l}x + \frac{r+l}{r-l}z}{-z} = \frac{2n}{r-l} \frac{x}{-z} + \frac{r+l}{r-l}, \\ y_{ndc} &= \frac{y_c}{w_c} = \frac{\frac{2n}{t-b}y + \frac{t+b}{t-b}z}{-z} = \frac{2n}{t-b} \frac{y}{-z} + \frac{t+b}{t-b}, \\ z_{ndc} &= \frac{z_c}{w_c} = \frac{-\frac{f+n}{f-n}z - \frac{2fn}{f-n}}{-z} = \frac{f+n}{f-n} + \frac{2fn}{(f-n)z} \end{aligned}$$

This matrix maps the view frustum in camera space to the normalized cube in NDC after homogeneous division.

For all vertices that survived clipping, the resulting coordinates satisfy:

$$-1 \leq x_{ndc} \leq 1 \quad -1 \leq y_{ndc} \leq 1 \quad -1 \leq z_{ndc} \leq 1$$

Therefore, the visible region lies inside the **canonical cube**:

$$[-1, 1]^3$$

✓ Comparison of clipping and discarding in world space and clip space

When a triangle intersects the view frustum, it must be clipped so that only the portion inside the frustum is rasterized. Although the clipping procedure is conceptually similar in world space and clip space, the mathematical complexity differs significantly. Clipping and discarding in clip space will save **85%** in instructions.

1A. Discarding in world space:

As described in the section *Dot Product*:

The definition of Dot Product is:

$$\mathbf{a} = v_1 - v_0, \mathbf{b} = v_2 - v_0$$

$$\mathbf{a} \cdot \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \cos(\Theta)$$

When $(\mathbf{p} - v_0) \cdot \mathbf{n} < 0 \Rightarrow \mathbf{p}$ lies on the back (inner) side of the plane.

For the ray-plane (line-plane) intersection, the d_i can be obtained by choosing any point $\mathbf{p}_0$ on the plane with normal vector $\mathbf{n}_i$.

$$d_i = -\mathbf{n}_i \cdot \mathbf{p}_0$$

In world (or view) space, the view frustum is bounded by six arbitrary planes, each defined by a normal vector $\mathbf{n}$ and distance d .

For each vertex $\mathbf{p}$, discarding requires testing against all planes:

$$(\mathbf{p} - \mathbf{p}_0) \cdot \mathbf{n}_i < 0$$

$$\Rightarrow \mathbf{n}_i \cdot \mathbf{p} + d_i < 0 \quad \text{for any } i \in [1, 6]$$

Cost per vertex

- 6 dot products (each ≈ 3 multiplications + 2 additions)
- 6 additions with plane constants
- 6 comparisons

Approximate arithmetic cost:

- 18 multiplications
- 18 additions
- 6 comparisons

1B. Discarding in clip space:

In clip space, vertices are represented in homogeneous coordinates (x_c, y_c, z_c, w_c) . The view frustum becomes an axis-aligned volume defined by:

$$-w_c \leq x_c \leq w_c \quad -w_c \leq y_c \leq w_c \quad -w_c \leq z_c \leq w_c$$

Approximate arithmetic cost:

- 6 comparisons

Overall arithmetic instruction reduction **85% ~ 95%**.

2A. Clipping in World Space:

Edge-plane intersection

As described in the section *Dot Product*, the ray-plane (line-plane) intersection can be derived as follows:

$$t = \frac{-(\mathbf{n} \cdot \mathbf{p}_0 + d)}{\mathbf{n} \cdot (\mathbf{p}_1 - \mathbf{p}_0)}$$

Each frustum plane requires a separate equation and dot-product evaluation.

Triangle reconstruction

After computing all intersection points:

- New vertices are inserted along intersecting edges
- The original triangle is split into one or more triangles
- Perspective projection is applied afterward

Care must be taken to preserve perspective correctness during interpolation.

2B. Clipping in Clip Space:

Edge-plane intersection

Edges are interpolated linearly in homogeneous space:

$$\mathbf{v}(t) = \mathbf{v}_0 + t(\mathbf{v}_1 - \mathbf{v}_0)$$

Intersection with a clipping boundary is found by solving equations such as:

$$x(t) = \pm w(t), \quad y(t) = \pm w(t), \quad z(t) = \pm w(t)$$

Each case reduces to a single scalar equation for t .

$$\begin{aligned} x(t) &= x_0 + t(x_1 - x_0) & w(t) &= w_0 + t(w_1 - w_0) \\ x(t) &= w(t) \rightarrow \\ x_0 + t(x_1 - x_0) &= w_0 + t(w_1 - w_0) \\ x_0 - w_0 &= t[(w_1 - w_0) - (x_1 - x_0)] \\ t &= \frac{x_0 - w_0}{(x_0 - w_0) - (x_1 - w_1)} \end{aligned}$$

Compare $t = \frac{x_0 - w_0}{(x_0 - w_0) - (x_1 - w_1)}$ and the equation from world space $t = \frac{-(\mathbf{n} \cdot \mathbf{p}_0 + d)}{\mathbf{n} \cdot (\mathbf{p}_1 - \mathbf{p}_0)}$, it saves **85%** for reducing two dot products and additional operations.

Triangle reconstruction

After clipping:

- New vertices remain in homogeneous coordinates
- Perspective division is deferred
- Linear interpolation remains perspective-correct

The final step applies the perspective divide:

$$(x_c, y_c, z_c, w_c) \rightarrow \left(\frac{x_c}{w_c}, \frac{y_c}{w_c}, \frac{z_c}{w_c} \right)$$

4.3 Comparison and Practical Implications

- World-space clipping and discarding uses general plane equations and complex geometry.
- Clip-space clipping and discarding uses axis-aligned bounds and simple interpolation.
- Perspective correctness is naturally preserved in clip space.
- GPU hardware can implement clip-space clipping and discarding efficiently.

For these reasons, modern graphics pipelines perform triangle clipping and discarding in clip space, not in world space.

Reference:

1. Every computer graphics book covers the topic of transformation of objects and their positions in space. Chapter 4 of the *Blue Book: OpenGL SuperBible, 7th Edition* provides a concise yet useful 40-page overview of transformation concepts and is good material for gaining a deeper understanding of transformations. description of transformation.
2. Chapter 7 of the Red Book covers transformations and projections.
3. https://en.wikipedia.org/wiki/3D_projection

- *Example of OpenGL program*
- *3D Rendering*
 - *Animation Parameters*
 - *3D Rendering Pipeline*
 - *Mesh Construction*
 - *Tessellation Example*
 - *Mobile GPU 3D Rendering*
 - * *TBDR —Tile-Based Deferred Rendering*
 - * *TBDR Rendering*
 - *Mesh-Shader Pipeline*
 - *Animation Example*
 - *Ray Tracing Pipeline*
- *GLSL (GL Shader Language)*
 - *Background*
 - *Examples*
 - *Goals*
 - *GLSL vs. C: Feature Overview*
 - *GLSL Qualifiers by Shader Stage*
- *OpenGL Buffers*
- *OpenGL Shader Compiler*

5.1 Example of OpenGL program

The following example is from the OpenGL Red Book and its example code¹².

¹ <http://www.opengl-redbook.com>

² <https://github.com/openglredbook/examples>

References/triangles.vert

```
#version 400 core

layout( location = 0 ) in vec4 vPosition;

void
main()
{
    gl_Position = vPosition;
}
```

References/triangles.frag

```
#version 450 core

out vec4 fColor;

void main()
{
    fColor = vec4(0.5, 0.4, 0.8, 1.0);
}
```

References/01-triangles.cpp

```
1  //////////////////////////////////////
2  //
3  //  Triangles.cpp
4  //
5  //////////////////////////////////////
6
7  #include "vgl.h"
8  #include "LoadShaders.h"
9
10 enum VAO_IDs { Triangles, NumVAOs };
11 enum Buffer_IDs { ArrayBuffer, NumBuffers };
12 enum Attrib_IDs { vPosition = 0 };
13
14 GLuint  VAOs[NumVAOs];
15 GLuint  Buffers[NumBuffers];
16
17 const GLuint  NumVertices = 6;
18
19 //-----
20 //
21 //  init
22 //
23
24 void
25 init( void )
26 {
```

(continues on next page)

(continued from previous page)

```

27     glGenVertexArrays( NumVAOs, VAOs ); // Same with glCreateVertexArray( NumVAOs,
↪VAOs );
28     // https://stackoverflow.com/questions/24441430/glgen-vs-glcreate-naming-
↪convention
29     // Make the new VAO:VAOs[Triangles] active, creating it if necessary.
30     glBindVertexArray( VAOs[Triangles] );
31     // opengl->current_array_buffer = VAOs[Triangles]
32
33     GLfloat vertices[NumVertices][2] = {
34         { -0.90f, -0.90f }, { 0.85f, -0.90f }, { -0.90f, 0.85f }, // Triangle 1
35         { 0.90f, -0.85f }, { 0.90f, 0.90f }, { -0.85f, 0.90f } // Triangle 2
36     };
37
38     glCreateBuffers( NumBuffers, Buffers );
39
40     // Make the buffer the active array buffer.
41     glBindBuffer( GL_ARRAY_BUFFER, Buffers[ArrayBuffer] );
42     // Attach the active VBO:Buffers[ArrayBuffer] to VAOs[Triangles]
43     // as an array of vectors with 4 floats each.
44     // Kind of like:
45     // opengl->current_vertex_array->attributes[attr] = {
46     //     type = GL_FLOAT,
47     //     size = 4,
48     //     data = opengl->current_array_buffer
49     // }
50     // Can be replaced with glVertexArrayVertexBuffer(VAOs[Triangles], Triangles,
51     // buffer[ArrayBuffer], ArrayBuffer, sizeof(vmath::vec2));,
↪glVertexArrayAttribFormat(), ...
52     // in OpenGL 4.5.
53
54     glBufferStorage( GL_ARRAY_BUFFER, sizeof(vertices), vertices, 0);
55
56     ShaderInfo shaders[] =
57     {
58         { GL_VERTEX_SHADER, "media/shaders/triangles/triangles.vert" },
59         { GL_FRAGMENT_SHADER, "media/shaders/triangles/triangles.frag" },
60         { GL_NONE, NULL }
61     };
62
63     GLuint program = LoadShaders( shaders );
64     glUseProgram( program );
65
66     glVertexAttribPointer( vPosition, 2, GL_FLOAT,
67                           GL_FALSE, 0, BUFFER_OFFSET(0) );
68     glEnableVertexAttribArray( vPosition );
69     // Above two functions specify vPosition to vertex shader at layout (location = 0)
70 }
71
72 //-----
73 //
74 // display
75 //

```

(continues on next page)

(continued from previous page)

```

76
77 void
78 display( void )
79 {
80     static const float black[] = { 0.0f, 0.0f, 0.0f, 0.0f };
81
82     glClearBufferfv(GL_COLOR, 0, black);
83
84     glBindVertexArray( VAOs[Triangles] );
85     glDrawArrays( GL_TRIANGLES, 0, NumVertices );
86 }
87
88 //-----
89 //
90 // main
91 //
92
93 #ifndef _WIN32
94 int CALLBACK WinMain(
95     _In_ HINSTANCE hInstance,
96     _In_ HINSTANCE hPrevInstance,
97     _In_ LPSTR     lpCmdLine,
98     _In_ int       nCmdShow
99 )
100 #else
101 int
102 main( int argc, char** argv )
103 #endif
104 {
105     glfwInit();
106
107     GLFWwindow* window = glfwCreateWindow(800, 600, "Triangles", NULL, NULL);
108
109     glfwMakeContextCurrent(window);
110     gl3wInit();
111
112     init();
113
114     while (!glfwWindowShouldClose(window))
115     {
116         display();
117         glfwSwapBuffers(window);
118         glfwPollEvents();
119     }
120
121     glfwDestroyWindow(window);
122
123     glfwTerminate();
124 }

```

Init():

- Generate Vertex Array VAOs and bind VAOs[0].

```
(glGenVertexArrays(NumVAOs, VAOs); glBindVertexArray(VAOs[Triangles]); glCreateBuffers(NumBuffers, Buffers);)
```

A vertex-array object holds various data related to a collection of vertices. Those data are stored in buffer objects and managed by the currently bound vertex-array object.

– `glBindBuffer(GL_ARRAY_BUFFER, Buffers[ArrayBuffer]);`

Because there are many different places where buffer objects can be in OpenGL, when we bind a buffer, we need to specify what we'd like to use it for. In our example, because we're storing vertex data into the buffer, we use `GL_ARRAY_BUFFER`. The place where the buffer is bound is known as the binding target.

- According to the counter-clockwise rule in the previous section, triangle primitives are defined in variable *vertices*. After binding OpenGL VBO Buffers[0] to *vertices*, vertex data will be sent to the memory of the server (GPU).

Think of the “active” buffer as just a global variable, and there are a bunch of functions that use the active buffer instead of taking using a parameter. These global state variables are the ugly side of OpenGL⁶ and can be replaced with `glVertexArrayVertexBuffer()`, `glVertexArrayAttribFormat()`, etc. Then call `glBindVertexArray(vao)` before drawing in OpenGL 4.5⁷⁸.

- `glVertexAttribPointer(vPosition, 2, GL_FLOAT, GL_FALSE, 0, BUFFER_OFFSET(0));`

During GPU rendering, each vertex position will be held in *vPosition* and passed to the “triangles.vert” shader through the *LoadShaders(shaders)* function.

`glfwSwapBuffers(window);`

- You've already used double buffering for animation. Double buffering is done by making the main color buffer have two parts: a front buffer that's displayed in your window; and a back buffer, which is where you render the new image. When you swap the buffers (by calling `glfwSwapBuffers()`, for example), the front and back buffers are exchanged⁹.

`display();`

- Bind VAOs[0], set render mode to `GL_TRIANGLES` and send vertex data to Buffer (gpu memory, OpenGL pipeline). Next, GPU will do rendering pipeline described in next section.

The *triangles.vert* has input *vPosition* and no output variable, so using `gl_Position` default variable without declaration. The *triangles.frag* has not defined input variable and has defined output variable *fColor* instead of using `gl_FragColor`.

The “in” and “out” in shaders above are “type qualifier”. A type qualifier is used in the OpenGL Shading Language (GLSL) to modify the storage or behavior of global and locally defined variables. These qualifiers change particular aspects of the variable, such as where they get their data from and so forth¹⁰.

Though attribute and varying are removed from later version 1.4 of OpenGL, many materials in website using them¹¹¹². It's better to use “in” and “out” to replace them as the following code. OpenGL has a few ways to binding API's variable with shader's variable. `glVertexAttrib*` as the following code and `glBindAttribLocation()`¹³, ...

replace attribute and varying with in and out

```
uniform float scale;
layout (location = 0) attribute vec2 position;
// layout (location = 0) in vec2 position;
```

(continues on next page)

⁶ <https://stackoverflow.com/questions/21652546/what-is-the-role-of-glbindvertexarrays-vs-glbindbuffer-and-what-is-their-relatio>

⁷ <https://registry.khronos.org/OpenGL-Refpages/gl4/html/glBindVertexBuffer.xhtml>

⁸ Page 152 of Blue book: OpenGL SuperBible 7th Edition.

⁹ Section of Color Buffer, page 222-223 of book “OpenGL Programming Guide 9th Edition”^{Page 53, 1}.

¹⁰ [<https://www.khronos.org/opengl/wiki/Type_Qualifier_\(GLSL\)>](https://www.khronos.org/opengl/wiki/Type_Qualifier_(GLSL))

¹¹ [<https://www.khronos.org/opengl/wiki/Type_Qualifier_\(GLSL\)#Removed_qualifiers>](https://www.khronos.org/opengl/wiki/Type_Qualifier_(GLSL)#Removed_qualifiers)

¹² <https://github.com/vispy/vispy/issues/242>

¹³ [<https://www.khronos.org/opengl/wiki/Layout_Qualifier_\(GLSL\)>](https://www.khronos.org/opengl/wiki/Layout_Qualifier_(GLSL))

(continued from previous page)

```
layout (location = 1) attribute vec4 color;
// layout (location = 1) in vec4 color;
varying vec4 v_color;
// out v_color

void main()
{
    gl_Position = vec4(position*scale, 0.0, 1.0);
    v_color = color;
}
```

```
// OpenGL API
GLfloat attrib[] = { x * 0.5f, x * 0.6f, x* 0.4f, 0.0f };
// Update the value of input attribute 1 : layout (location = 1) in vec4 color
glVertexAttrib4fv(1, attrib);
```

```
varying vec4 v_color;
// in vec4 v_color;

void main()
{
    gl_FragColor = v_color;
}
```

An OpenGL program is made of two shaders¹⁴¹⁵:

- The vertex shader is (commonly) executed once for every vertex we want to draw. It receives some attributes as input, computes the position of this vertex in space and returns it in a variable called `gl_Position`. It also defines some varyings.
- The fragment shader is executed once for each pixel to be rendered. It receives some varyings as input, computes the color of this pixel and returns it in a variable called `fColor`.

Since we have 6 vertices in our buffer, this shader will be executed 6 times by the GPU (once per vertex)! We can also expect all 6 instances of the shader to be executed in parallel, since a GPU has many cores.

5.2 3D Rendering

3D animation is the process of creating moving images by manipulating digital objects within a three-dimensional space. 3D rendering is the process of converting 3D models into 2D images on a computer²³.

Based on the previous section of 3D modeling, the 3D modeling tool produces a 3D vertex model. Then, the 3D gaming designers or programmers will manually write the scripting language that instructs the 3D engine to issue corresponding OpenGL calls, as shown in Fig. 5.1.

In section *SW Stack and Data Flow*, we mentioned the GPU will generate the rendering image for each frame according the 3D Information and Uniform Updates sent from CPU, and write each of the final frame of data in the form of color pixels to framebuffer (video memory) as Fig. 5.2.

However, the 3D engine only executes the animation within the frustum-projection space used for on-screen rendering, as illustrated in Fig. 5.3.

¹⁴ <https://engineering.monstar-lab.com/en/post/2022/03/01/Introduction-To-GPUs-With-OpenGL/>

¹⁵ <https://glumpy.github.io/modern-gl.html>

²³ https://en.wikipedia.org/wiki/3D_rendering




Fig. 5.1: Graphic SW Stack and data flow in rendering




Fig. 5.2: The input and output for GPU rendering




Fig. 5.3: The needed animation space

5.2.1 Animation Parameters

✓ CPU only updates small animation parameters named **Uniform Updates** as appeared in Fig. 5.1; GPU computes the heavy per-vertex work.

The 3D animation will trigger the 3D rendering process for each 2D image drawing according to the **Uniform Updates**.

The “small animation parameters” updated by the CPU are formally called:

- ✓ Uniform updates
- ✓ Constant buffer updates
- ✓ Per-frame / per-draw constants
- ✓ Bone matrix palette updates (for skinning)
- ✓ Morph weight updates (for morphing)

These are the correct technical terms used in modern graphics pipelines.

⚓ The Proper Term: “Uniform Updates”

The most accurate and universal name is:

- ✓ Uniform updates

or

- ✓ Updating uniform buffers

Because the CPU is updating uniform data that the GPU reads during shading.

Examples of uniform data:

- bone matrices
- morph weights
- animation time
- material parameters

- camera matrices
- light parameters

These are small, constant-for-the-draw values.

⚙ More Specific Terms Used in Game Engines

1. Animation Parameters

Used in animation systems:

- “animation parameters”
- “skinning parameters”
- “bone palette”
- “morph weights”

2. Per-Frame Constants

Used in engine architecture:

- “frame constants”
- “per-frame constant buffer”
- “global shader constants”

3. Per-Draw Constants

Used in render pipelines:

- “per-draw uniform block”
- “per-object constant buffer”
- “material constant buffer”

⚙ In Modern APIs (GL, Vulkan, DirectX)

OpenGL

- Uniforms
- Uniform Buffer Objects (UBOs)
- Shader Storage Buffer Objects (SSBOs)

DirectX

- Constant Buffers (CBuffers)

Vulkan

- Descriptor sets
- Uniform buffers


All refer to the same concept: small CPU-updated data that the GPU reads during shading.

5.2.2 3D Rendering Pipeline

The steps are shown in Fig. 5.4.

- A fragment can be treated as a pixel in 3D spaces, which is aligned with the pixel grid, with attributes such as position, color, normal and texture.

³ https://www3.ntu.edu.sg/home/ehchua/programming/opengl/CG_BasicsTheory.html



3D Graphics Rendering Pipeline: Output of one stage is fed as input of the next stage. A vertex has attributes such as (x, y, z) position, color (RGB or RGBA), vertex-normal (n_x, n_y, n_z) , and texture. A primitive is made up of one or more vertices. The rasterizer raster-scans each primitive to produce a set of grid-aligned fragments, by interpolating the vertices.

Fig. 5.4: 3D Graphics Rendering Pipeline^{Page 61, 3}

From the previous Fig. 3.5 and Fig. 3.6 in section *SW Stack and Data Flow*, we introduce the 3D animation data are classified as follows:

- **Vertex Data = 3D model information** (the mesh (geometry), such as VBO/VAO)
- **Animation Parameters = per-frame uniform updates** (transforms, bone matrices, camera, materials, ...)

The complete steps of 3D Rendering pipeline, **excluding animation** are shown in the Fig. 5.5 from the OpenGL website²⁴ and in the Fig. 5.6. The website also provides a description for each stage. To clarify the modern GPU pipeline, Fig. 5.7 shows the use of Primitive Assembly (fixed-function) and Primitive Setup (fixed-function).

²⁴ https://www.khronos.org/opengl/wiki/Rendering_Pipeline_Overview




Fig. 5.5: Diagram of the Rendering Pipeline. The blue boxes are programmable shader stages. Shaders with dashed outlines indicate optional shader stages.




Figure 1.2 OpenGL pipeline

Fig. 5.6: OpenGL pipeline in blue book

In Fig. 5.5, the “Primitive Assembly” is the “Output Primitive Assembly”. A more accurate ordering is:

```

Vertex Shader
  ↓
Patch assembly / Input Primitive Assembly
  ↓
Tessellation
  ├── Tessellation Control Shader
  ├── Primitive Generation
  └── Tessellation Evaluation Shader
  ↓
Geometry Shader
  ↓
Output Primitive Assembly
  ↓
Rasterization
  
```

As shown in Fig. 5.7:

- Vertex Shader and Tessellation: process and transform **vertex** data.
- Primitive Processing: processes and transforms **primitive** data.




Fig. 5.7: Modern GPU Pipeline

- Rasterizer: **Primitives → Fragments.**
- Fragment Shader: **Fragments → Colored Fragments.**

As illustrated in the *Cross Product* section,

- ✓ Each mesh (triangle or primitive) has a fixed “outer” and “inner” side, determined by CCW ordering in object space.
- ✓ By reading these CCW-ordered vertices in a triangle sequentially, the shape and surface orientation of the 3D model can be constructed.
- ✓ There is no need to wait for the entire mesh to be received; once three CCW-ordered vertices in a triangle are available, each triangle can be processed correctly.
- ✓ When the camera moves to the inside an object: CCW ↔ CW flips.

This means:

- ✓ Vertex Shader and Tessellation: **may process each vertex independently** as long as the vertex order in a triangle is preserved.
- ✓ Once **three CCW-ordered vertices are available**, Primitive Assembly can convert them into a triangle and pass it to the next pipeline stage.

After vertices are assembled into primitives (such as triangles), the front-facing and back-facing surfaces can be determined, and the hidden primitives can be removed.

The Red Book and Blue Book show only **Vertex Specification** and **Vertex Data** in the rendering flow because they **never** show **Animation Parameters** as part of the rendering flow. The animation flow from CPU to GPU is shown in Fig. 5.8, based on Fig. 5.4.

Each draw call may correspond to:




Fig. 5.8: CPU and GPU Pipeline For Shaders

- one mesh
- one submesh
- one meshlet (in mesh-shader pipelines)
- or many meshes batched together

Although the Rendering Pipeline shown in Fig. 5.5 and Fig. 5.6 do not explicitly include per-frame animation flow—because the inputs are labeled **Vertex Specification** and **Vertex Data** and they do not show Animation Parameters as part of the rendering process—the pipeline is still applicable.

However the following table from OpenGL rendering pipeline Figure 1.2 and its stages from the book *OpenGL Programming Guide, 9th Edition* ^{Page 53, 1} is broad enough to cover animation.

Table 5.1: OpenGL rendering pipeline from page 10 and page 36 of book
 “OpenGL Programming Guide 9th Edition”^{Page 53, 1} and ^{Page 62, 24}.

Stage.	Description
Vertex Shading	Vertex → Vertex and other data such as color for later passes. For each vertex issued by a drawing command, a vertex shader processes the data associated with that vertex. Vertex Shader: provides the Vertex → Vertex transformation effects controlled by the users.
Tessellation Shading	Create more detail on demand when zoomed in. After the vertex shader processes each vertex, the tessellation shader stage (if active) continues processing. The tessellation stage is actually divided into two shaders known as the tessellation control shader and the tessellation evaluation shader . A single patch from Tessellation Control Shader (TCS) and Tessellation Evaluation Shader (TVS) can generate millions of micro-triangles . See reference below.
Primitive Assembly Geometry Shader	This is a fixed-function hardware stage: forms triangles/lines/points. Primitive Transformation: output zero primitives (cull), output one primitive (pass-through), output many primitives (amplify) and output different topology (e.g., point → quad) Allows additional processing of geometric primitives. This stage may create new primitives before rasterization. The Geometry shading stage is another optional stage that can modify entire geometric primitives within the OpenGL pipeline. This stage operates on individual geometric primitives allowing each to be modified. In this stage, you might generate more geometry from the input primitive, change the type of geometric primitive (e.g., converting triangles to lines), or discard the geometry altogether.
Culling Clipping	Remove entire primitives that are hidden or outside the viewport. Clip the hidden portions of the primitive, separating it into visible and hidden parts and discarding the hidden portions.
Primitive Setup (rasterization preparation)	This stage: takes the final primitive (after GS), computes edge equations, computes barycentric interpolation coefficients, determine rasterization rules and prepare for triangle traversal.
Rasterization	Geometric Primitives → Fragment. The job of the rasterizer is to determine which screen locations are covered by a particular piece of geometry (point, line, or triangle). Knowing those locations, along with the input vertex data, the rasterizer linearly interpolates the data values for each varying variable in the fragment shader and sends those values as inputs into your fragment shader. A fragment can be treated as a pixel in 3D spaces, which is aligned with the pixel grid, with attributes such as position, color, normal and texture. Early Depth and Stencil Tests (Early-Z): reject hidden fragments before shading.
Fragment Shading	Fragment → Colored Fragment. Determine color for each pixel. In this stage, a fragment’s color and depth values are computed and then sent for further processing in the fragment-testing and blending parts of the pipeline. The final stage where you have programmable control over the color of a screen location is fragment shading. In this shader stage, you use a shader to determine the fragment’s final color (although the next stage, per-fragment operations, can modify the color one last time) and potentially its depth value. Fragment shaders are very powerful, as they often employ texture mapping to augment the colors provided by the vertex processing stages. A fragment shader may also terminate processing a fragment if it determines the fragment shouldn’t be drawn; this process is called fragment discard. A helpful way of thinking about the difference between shaders that deal with vertices and fragment shaders is this: vertex shading (including tessellation and geometry shading) determines where on the screen a primitive is, while fragment shading uses that information to determine what color that fragment will be.

Table 5.2: Continue OpenGL rendering pipeline from page 10 and page 36 of book “OpenGL Programming Guide 9th Edition”^{Page 53, 1} and^{Page 62, 24}.

Stage.	Description
Per-Fragment Operations	During this stage, a fragment’s visibility is determined using depth testing (also commonly known as z-buffering) and stencil testing. If a fragment successfully makes it through all of the enabled tests, it may be written directly to the framebuffer, updating the color (and possibly depth value) of its pixel, or if blending is enabled, the fragment’s color will be combined with the pixel’s current color to generate a new color that is written into the framebuffer.
Compute shading stage	Compute shader: may be applied in any stage. This is not part of the graphical pipeline like the stages above, but stands on its own as the only stage in a program. A compute shader processes generic work items, driven by an application-chosen range, rather than by graphical inputs like vertices and fragments. Compute shaders can process buffers created and consumed by other shader programs in your application. This includes framebuffer post-processing effects or really anything you want. Compute shaders are described in Chapter 12 of Red Book, “Compute Shaders” ²⁵ .

Tessellation

- Tessellation Shading: The core problem that Tessellation deals with is the static nature of 3D models in terms of their detail and polygon count. The thing is that when we look at a complex model such as a human face up close we prefer to use a highly detailed model that will bring out the tiny details (e.g. skin bumps, etc). A highly detailed model automatically translates to more triangles and more compute power required for processing. ...One possible way to solve this problem using the existing features of OpenGL is to generate the same model at multiple levels of detail (LOD). For example, highly detailed, average and low. We can then select the version to use based on the distance from the camera. This, however, will require more artist resources and often will not be flexible enough. ...Let’s take a look at how Tessellation has been implemented in the graphics pipeline. The core components that are responsible for Tessellation are two new shader stages and in between them a **fixed function** stage that can be configured to some degree but does not run a shader. The first shader stage is called **Tessellation Control Shader (TCS)**, the **fixed function** stage is called the **Primitive Generator (PG)**, and the second shader stage is called **Tessellation Evaluation Shader (TES)**. Some GPUs do not implement this fixed-function stage in hardware and do not provide TCS, TES, or geometry shaders. A user can write **compute shaders** instead to generate detail on the fly. This surface is usually defined by some **polynomial formula** and the idea is that moving a **CP** has an effect on the entire surface. ...The group of CPs is usually called a **Patch**²⁶. The data flow in the Tessellation Stage between the TCS, Fixed-Function Tessellator, and TES is illustrated in Fig. 5.9. Chapter 9 of Red Book^{Page 53, 1} has details. The next section *Tessellation Example* describes the details for the tessellation with an example.
- Tessellation **cannot** decrease the resolution of vertices from the VS. The Geometry Shader can **reduce geometry** (by discarding primitives), but it **cannot** reduce the number of input vertices coming from VS/TES. The rasterizer can **reduce fragments**, but it cannot reduce vertices.

Data Flow

The following diagrams summarize the OpenGL rendering pipeline: Fig. 5.9 and Fig. 5.10.

²⁵ Page 36 of book “OpenGL Programming Guide 9th Edition”^{Page 53, 1}.

²⁶ <https://ogldev.org/www/tutorial30/tutorial30.html>




Fig. 5.9: The part 1 of GPU Rendering Pipeline Stages



Fig. 5.10: The part 2 of GPU Rendering Pipeline Stages

The data flow through the OpenGL shader pipeline and the details of the TCS, Fixed-Function Tessellator, and TES are described below.

Table 5.3: Data Flow Through the OpenGL Shader Pipeline

Shader Stage	Input Data (from CPU or previous stage)	Output Data (to next stage)	How GPU Hardware Uses These Data (with Stage Name)
Vertex Shader	- Per-vertex attributes: - Positions (vec3/vec4) - Normals, tangents - Texture coordinates - Vertex colors - Skinning weights/indices - Uniforms and UBOs - Textures / samplers	- gl_Position (clip-space) - Varyings - Optional point size	Vertex Processing Stage: - ALUs transform vertices - Writes positions into Primitive Assembly - Stores varyings in interpolation registers
Tessellation Control Shader (TCS)	- Patch control points - Uniforms - Per-patch attributes	Modified control points: gl_out Tessellation levels: gl_TessLevelInner, gl_TessLevelOuter	Tessellation Control Stage: - Writes tessellation levels to fixed-function tessellator - Stores control points in patch memory
Fixed-Function Tessellator (TS)	- Modified control points: gl_out - Tessellation levels: gl_TessLevelInner, gl_TessLevelOuter - Per-patch attribute (triangles, quads, isolines) - Partitioning mode (integer, fractional_even, fractional_odd) - Winding order	Tessellated coordinates (u,v,w): gl_TessCoord - Bypass modified Control Points	Fixed-Function Tessellator (TS), also known as Primitive Generator (PG): - Generates tessellated domain coordinates (u,v,w) to TES
Tessellation Evaluation Shader (TES)	- Tessellated coordinates (u,v,w): gl_TessCoord - modified Control Points - Uniforms - gl_PrimitiveID	Tessellated Vertices: gl_Position - Any per-vertex varyings for GS or rasterizer - Optional custom attributes	Tessellation Evaluation Stage: - ALUs compute final vertex positions - Outputs to Primitive Assembly - Sends varyings to interpolation hardware
Geometry Shader	- Assembled primitives - All varyings	- Zero or more primitives - New varyings - New gl_Position	Geometry Processing Stage: - Allocates per-primitive scratch memory - Emits new primitives - Expands or reduces geometry

Table 5.4: Data Flow Through the OpenGL Shader Pipeline Continue

Shader Stage	Input Data (from CPU or previous stage)	Output Data (to next stage)	How GPU Hardware Uses These Data (with Stage Name)
Rasterizer (Fixed Function)	- Primitives (triangles/lines/points) - Per-vertex varyings	- Fragments - Interpolated varyings - gl_FragCoord	Rasterization Stage: - Barycentric units interpolate varyings - Generates fragments - Sends fragments to fragment shader cores
Fragment Shader	- gl_FragCoord - Interpolated varyings - Textures / samplers - Uniforms	- gl_FragColor or user-defined outputs - Depth override (optional)	Fragment Processing Stage: - ALUs compute pixel color - Texture units fetch texels - Outputs color/depth to ROP
Output Merger / ROP (Fixed Function)	- Fragment shader outputs - Depth/stencil values - Blending state	- Final framebuffer color - Updated depth/stencil buffers	Output Merger Stage: - Performs depth/stencil tests - Applies blending - Writes final pixels to framebuffer memory - Handles MSAA resolve

Varying

A varying is a piece of data that:

- Comes out of the vertex shader
- Gets interpolated by the rasterizer
- Arrives as input to the fragment shader

It is called **varying** because its value **varies across the surface of a triangle**.

Table 5.5: Examples of Common Varyings

Varying Name	Meaning	Why It Varies Across the Primitive
vNormal	Surface normal at each vertex	Lighting requires a smoothly changing normal across the triangle so per-pixel shading can compute correct diffuse and specular terms
vUV	Texture coordinates	Each pixel needs its own UV to sample the correct texel from the texture
vColor	Vertex color (per-vertex material tint)	Enables smooth color gradients or per-vertex painting effects
vWorldPos	World-space position of the vertex	Used for per-pixel lighting, reflections, shadows, and screen-space effects; must be interpolated so each fragment knows its own world position

In 2D animation, a model is a single flat view and can be viewed only from that side. To display different views, multiple 2D models must be created and switched from frame to frame²⁷.

²⁷ <https://tw.video.search.yahoo.com/search/video?fr=yfp-search-sb&p=2d+animation#id=12&vid=46be09edf57b960ae79e9cd077eea1ea&action=view>

5.2.3 Mesh Construction

As described in section *3D Rendering Pipeline*:

✓ By reading these CCW-ordered vertices in a triangle sequentially, the shape and surface orientation of the 3D model can be constructed.

Because a mesh contains many vertices shared across triangles, this section explains how OpenGL and GPUs avoid reprocessing those redundant vertices.”

OpenGL Triangles

OpenGL has:

- GL_TRIANGLES —every 3 vertices form an independent triangle
- GL_TRIANGLE_STRIP —consecutive vertices share edges
- GL_TRIANGLE_FAN —all triangles share the first vertex

For example:

```
glDrawArrays(GL_TRIANGLE_STRIP, 0, 6);
```

With vertices:

```

v0  ——— v2  ——— v4
|  /    |  /    |
| /     | /     |
| /     | /     |
v1  ——— v3  ——— v5
```

vertices:

v0, v1, v2, v3, v4, v5

Triangle list for GL_TRIANGLES:

indices:

```
(0,1,2)
(2,1,3)
(2,3,4)
(4,3,5)
```

Triangle strip for GL_TRIANGLE_STRIP:

indices:

```
0,1,2,3,4,5
```

The triangles are generated as:

```

Triangle 0: v0, v1, v2
Triangle 1: v2, v1, v3
Triangle 2: v2, v3, v4
Triangle 3: v4, v3, v5
```

So after the first 3 vertices, each additional vertex creates one new triangle.

For a GPU compiler, triangle strip is important because the primitive assembly stage has to keep track of the previous vertices and **alternates triangle winding order** for successive triangles.

For the first triangle:

$v_0 \rightarrow v_1 \rightarrow v_2$

For the second triangle, the shared edge is v_1 - v_2 . To maintain the same winding direction, OpenGL uses:

$v_2 \rightarrow v_1 \rightarrow v_3$

rather than:

$v_1 \rightarrow v_2 \rightarrow v_3$

That causes the second triangle's winding to be opposite.

3D models may create meshes with individual triangles or triangle strips. However, real 3D meshes are usually not one continuous strip. A model has:

- branches
- holes
- disconnected surfaces
- different materials
- hard edges
- UV seams
- normal seams

These make a single efficient triangle strip difficult. Modern GPUs also have **post-transform vertex caches**, so well-organized triangle lists can reuse vertex shader results efficiently.

The mesh itself is a collection of triangles. `GL_TRIANGLES` and `GL_TRIANGLE_STRIP` are primarily different primitive assembly methods for interpreting the index/vertex stream. Also, modern APIs such as Vulkan and Direct3D still support triangle strips, but triangle lists are the dominant representation for general-purpose 3D meshes.

Post-Transform Vertex Cache

A modern GPU usually has a small **post-transform vertex cache** that stores the results of vertex shader execution.

This cache allows the GPU to avoid executing the vertex shader multiple times for the same vertex when that vertex is referenced by several triangles.

For example, consider the following indexed triangle list:

```
Triangle 0: 0, 1, 2
Triangle 1: 2, 1, 3
Triangle 2: 2, 3, 4
```

The vertices 1 and 2 are shared by the first two triangles. The GPU can process the first triangle as follows:

Index	Cache	Action
0	miss	Run vertex shader
1	miss	Run vertex shader
2	miss	Run vertex shader

When processing the second triangle:

Index	Cache	Action
2	hit	Reuse vertex shader result
1	hit	Reuse vertex shader result
3	miss	Run vertex shader

Therefore, although vertex 1 and vertex 2 appear multiple times in the index buffer, their vertex shaders do not need to be executed again as long as their transformed results are still in the cache.

The cache is called a *post-transform* cache because it stores the results *after* vertex transformation, rather than merely caching the original vertex attributes.

A simplified GPU pipeline is:

```

Index Buffer
  |
  v
Vertex Fetch
  |
  v
Vertex Shader
  |
  v
Post-Transform Vertex Cache
  |
  v
Primitive Assembly
  |
  v
Rasterization

```

Triangle Order Matters

The cache has limited capacity. If a vertex is referenced again after it has been evicted from the cache, the vertex shader must execute again.

For example, consider a small cache:

```

T0: 0, 1, 2
T1: 3, 4, 5
T2: 0, 1, 2

```

After processing T0, vertices 0, 1, and 2 may be in the cache. Processing T1 replaces them with 3, 4, and 5. When T2 is processed, vertices 0, 1, and 2 are no longer in the cache, so their vertex shaders must execute again.

A different triangle order can keep shared vertices in the cache for longer:

```

T0: 0, 1, 2
T1: 2, 1, 3
T2: 2, 3, 4
T3: 3, 4, 5

```

This ordering has good locality because neighboring triangles share recently processed vertices.

Mesh Optimization

A **mesh optimizer** can reorder the triangles, and therefore the indices, to improve vertex-cache locality.

The mesh optimizer does not remove duplicated vertex processing by itself. Instead, it arranges the triangle order so that the GPU's post-transform vertex cache can reuse more previously computed vertex shader results.

The relationship is:

```
Mesh Optimizer
|
| Reorders triangles / indices
v
Better Vertex Locality
|
v
Post-Transform Vertex Cache
|
| More cache hits
v
Fewer Vertex Shader Executions
```

This is one reason that a well-organized **triangle list** can achieve substantial vertex reuse on modern GPUs, even without using a triangle strip.

Triangle strips provide vertex reuse through their primitive-assembly rules, whereas an optimized triangle list relies on the post-transform cache to reuse previously computed vertex shader results.

The two mechanisms are therefore different:

```
Triangle Strip
|
+-- Implicit vertex reuse in primitive assembly

Optimized Triangle List
|
+-- Explicit indices
|
+-- Good triangle ordering
|
+-- Post-transform cache reuse
```

5.2.4 Tessellation Example

In Chapter 9 (Tessellation), the Red Book^{Page 53, 1} focuses on:

- `gl_TessLevelOuter[]`
- `gl_TessLevelInner[]`

It never mentioned to generate modified CPs in TCS. The following example gives the output for (TCS → TS → TES) in patching a single rectangle.

An example for Inflated 4×4 Bézier Patch (TCS → TS → TES)

The following diagram illustrates the complete OpenGL tessellation pipeline for a **4×4** bicubic Bézier patch on **1 single rectangle**. Only the four interior control points (5, 6, 9, 10) are lifted off the plane, producing a smooth inflated surface.

Tessellation Control Shader (TCS): output:

- **modified Control Points (CPs, Patch):** `gl_out`
- **Tessellation level:** `gl_TessLevelInner`, `gl_TessLevelOuter`

Another name for **CPs** is **Patch**.

The TCS outputs 16 CPs arranged in a 4×4 grid. Only CPs 5, 6, 9, and 10 are elevated to create curvature.

```
#version 450 core
layout(vertices = 16) out;

void main()
{
    // Copy all CPs
    gl_out[gl_InvocationID].gl_Position =
        gl_in[gl_InvocationID].gl_Position;

    // Inflate interior CPs
    if (gl_InvocationID == 5 ||
        gl_InvocationID == 6 ||
        gl_InvocationID == 9 ||
        gl_InvocationID == 10)
    {
        gl_out[gl_InvocationID].gl_Position +=
            vec4(0.0, 0.0, 1.0, 0.0);
    }

    // Tessellation levels
    if (gl_InvocationID == 0) {
        gl_TessLevelOuter[0] = 4.0;
        gl_TessLevelOuter[1] = 4.0;
        gl_TessLevelOuter[2] = 4.0;
        gl_TessLevelOuter[3] = 4.0;

        gl_TessLevelInner[0] = 4.0;
        gl_TessLevelInner[1] = 4.0;
    }
}
```

Fixed-Function Tessellator (TS), also known as the **Primitive Generator (PG)**: output:

- **Tessellated coordinates (u,v,w)**: gl_TessCoord

The PG uses the tessellation levels to generate a **set of points** inside the triangle. Each point has its own barycentric coordinates. This set is called the **tessellated coordinates**.

The grid size depends on tessellation levels:

- If gl_TessLevelOuter[0..3] = 4.0 and gl_TessLevelInner[0..1] = 4.0 → you get a 5×5 grid, **Tessellated coordinates (u,v,w)**
- If you set 8.0 → you get a 9×9 grid
- If you set 2.0 → you get a 3×3 grid

The fixed-function tessellator generates a 5×5 evaluation grid for tessellation level 4.0. No shading language code is written for this stage.

Tessellation Evaluation Shader (TES): output:

- **Tessellated Vertices**: gl_Position

For each (u, v) , the TES computes the surface point $P(u, v)$ as:

$$P(u, v) = \sum_{i=0}^3 \sum_{j=0}^3 B_i(u) B_j(v) P_{ij}$$

where the Bernstein basis functions are:

$$B_0(t) = (1-t)^3, \quad B_1(t) = 3t(1-t)^2, \quad B_2(t) = 3t^2(1-t), \quad B_3(t) = t^3.$$

The TES evaluates the Bézier surface at each tessellated (u, v) coordinate using the 16 CPs.

```
#version 450 core
layout(quads, equal_spacing, cw) in;

float B(int i, float t) {
    if (i == 0) return (1 - t) * (1 - t) * (1 - t);
    if (i == 1) return 3 * t * (1 - t) * (1 - t);
    if (i == 2) return 3 * t * t * (1 - t);
    return t * t * t;
}

void main()
{
    float u = gl_TessCoord.x;
    float v = gl_TessCoord.y;

    vec4 p = vec4(0.0);
    int idx = 0;

    for (int i = 0; i < 4; ++i) {
        float bu = B(i, u);
        for (int j = 0; j < 4; ++j) {
            float bv = B(j, v);
            p += gl_in[idx].gl_Position * (bu * bv);
            idx++;
        }
    }

    gl_Position = p;
}
```

Result

The output for (TCS $\rightarrow$ TS $\rightarrow$ TES) in patching a single rectangle as the following table.

Inflated Bézier Patch: Control Points and Evaluated Surface (vec4)

All control points use homogeneous coordinates $(x, y, z, w = 1.0)$. Evaluated surface points $P(u,v)$ are also vec4.

CP Index	Grid Position (i, j)	Control Point (x, y, z, w)	Evaluated P(u,v) = vec4
0	(0, 0)	(0, 0, 0, 1)	(0.0, 0.0, 0.0, 1)
1	(1, 0)	(1, 0, 0, 1)	(1.0, 0.0, 0.0, 1)
2	(2, 0)	(2, 0, 0, 1)	(2.0, 0.0, 0.0, 1)
3	(3, 0)	(3, 0, 0, 1)	(3.0, 0.0, 0.0, 1)
4	(0, 1)	(0, 1, 0, 1)	(0.0, 1.0, 0.0, 1)
5	(1, 1)	(1, 1, 1, 1)	(1.0, 1.0, 0.5625, 1)
6	(2, 1)	(2, 1, 1, 1)	(2.0, 1.0, 0.5625, 1)
7	(3, 1)	(3, 1, 0, 1)	(3.0, 1.0, 0.0, 1)
8	(0, 2)	(0, 2, 0, 1)	(0.0, 2.0, 0.0, 1)
9	(1, 2)	(1, 2, 1, 1)	(1.0, 2.0, 0.5625, 1)
10	(2, 2)	(2, 2, 1, 1)	(2.0, 2.0, 0.5625, 1)
11	(3, 2)	(3, 2, 0, 1)	(3.0, 2.0, 0.0, 1)
12	(0, 3)	(0, 3, 0, 1)	(0.0, 3.0, 0.0, 1)
13	(1, 3)	(1, 3, 0, 1)	(1.0, 3.0, 0.0, 1)
14	(2, 3)	(2, 3, 0, 1)	(2.0, 3.0, 0.0, 1)
15	(3, 3)	(3, 3, 0, 1)	(3.0, 3.0, 0.0, 1)




Fig. 5.11: The final rendering result for 5x5 tessellated mesh.




Fig. 5.12: Geometry Shader (GS) can expand a 5x5 tessellated grid into a 6x6 mesh

The following TCS GLSL code from the Red Book adjusts the resolution of control points at runtime according to the distance of the square's vertices.

Specifying Tessellation Level Factors Using Perimeter Edge Centers.

```
#version 450 core

// Each patch has four precomputed edge centers:
//   edgeCenter[0] = left edge center
```

(continues on next page)

(continued from previous page)

```

//   edgeCenter[1] = bottom edge center
//   edgeCenter[2] = right edge center
//   edgeCenter[3] = top edge center
struct EdgeCenters {
    vec4 edgeCenter[4];
};

// Array of edge-center data, one entry per patch
uniform EdgeCenters patch[];

// Camera position in world space
uniform vec3 EyePosition;

layout(vertices = 16) out;

void main()
{
    // Pass through control points unchanged
    gl_out[gl_InvocationID].gl_Position =
        gl_in[gl_InvocationID].gl_Position;

    // Synchronize all invocations
    barrier();

    // Only invocation 0 computes tessellation levels
    if (gl_InvocationID == 0)
    {
        // Loop over the four perimeter edges
        for (int i = 0; i < 4; ++i)
        {
            // Distance from eye to this edge center
            float d = distance(
                patch[gl_PrimitiveID].edgeCenter[i],
                vec4(EyePosition, 1.0)
            );

            // Scale factor controlling how quickly tessellation increases
            const float lodScale = 2.5;

            // Convert distance to tessellation level
            float tessLOD = mix(
                0.0,
                gl_MaxTessGenLevel,
                d * lodScale
            );

            gl_TessLevelOuter[i] = tessLOD;
        }
    }
    #if 1
        // Compute the inner tessellation as the average of opposing outer
        // edges: differently from Red Book.
        // It's what most engines (Unreal, Unity HDRP, Vulkan samples) do.

```

(continues on next page)

(continued from previous page)

```

// Inner tessellation is the average of outer levels
float inner = 0.5 *
    (gl_TessLevelOuter[0] + gl_TessLevelOuter[2]);

inner = clamp(inner, 0.0, gl_MaxTessGenLevel);
gl_TessLevelInner[0] = inner;
gl_TessLevelInner[1] = inner;
#else
// The Red Book computes outer tessellation levels first, then
// derives the inner levels from the last computed tessLOD.
tessLOD = clamp(0.5 * tessLOD, 0.0, gl_MaxTessGenLevel);
gl_TessLevelInner[0] = tessLOD;
gl_TessLevelInner[1] = tessLOD;
#endif
}

```

The texture function with the argument `DisplacementMap` in the Red Book, as shown in the following code, does not return color data as in the Fragment Shader. It returns the vertex position data for displacement, such as a roughness map or anything related to surface appearance.

```
p += texture(DisplacementMap, gl_TessCoord.xy);
```

5.2.5 Mobile GPU 3D Rendering

The traditional desktop GPUs is **IMR —Immediate-Mode Rendering**: Cache misses dominate bandwidth.

TBDR —Tile-Based Deferred Rendering: Cache misses are nearly eliminated.

Note

Idea:

1. TBDR divides the whole frame into small tiles that fit entirely into on-chip **SRAM**.
2. Remove stages of Tessellation Control Shader (TCS), Tessellation Evaluation Shader (TES) and Geometry Shader (GS) since they are optional stages are shown in Fig. 5.6. Instead, developers use **compute shaders** before the graphics pipeline to generate meshlets, perform LOD selection, or add extra geometric detail for close-up **room-in** effects is shown as Fig. 5.15.

★ TBDR reduces **cache-miss** rate by roughly **10×—50×** compared to IMR, because all intermediate color/depth/stencil traffic stays in on-chip tile memory instead of going to L2/DRAM.

★ Desktop GPUs adopt IMR partly because GS/Tess/Mesh Shaders cannot run efficiently on TBDR. In addition, desktop GPUs adopt IMR because they have the **power**, **bandwidth**, and architectural freedom to **support unpredictable geometry pipelines** and massive workloads that would break TBDR's tile-based constraints.

TBDR —Tile-Based Deferred Rendering

△ For **low power** mobile device, mobile GPUs use **tile-based** rendering to **reduce the traffic to DRAM** as described below:

The traditional desktop GPUs is IMR —Immediate-Mode Rendering:

1. IMR: “Draw call arrives → render immediately”

CPU issues DrawCall #1

→ GPU transforms vertices

→ GPU rasterizes fragments

→ GPU writes to DRAM

It never waits to see the rest of the frame.

2. TBDR: “Draw call arrives → store geometry, don’t render yet”. TBDR processes it into two phases as follows:

Phase 1 —Binning (Full-Frame Geometry Processing) is shown as Fig. 5.13.




Fig. 5.13: TBDR Pipeline

When a draw call arrives on a TBDR GPU:

→ It runs the vertex shader

→ It transforms all triangles

→ It determines which tiles each triangle touches

→ It stores triangle references in **per-tile lists** is shown as follows:

Example of per-tile lists

```
Tile 0 → triangles: [T1, T7, T8, T20]
Tile 1 → triangles: [T2, T3, T7]
Tile 2 → triangles: [T4, T5, T6, T9, T10]
...
```

Phase 2 —Tile Rendering (Deferred Shading)

For each tile:

→ Load tile’s triangle list

→ Rasterize only those triangles

→ Shade only visible fragments

→ Keep all intermediate buffers in on-chip **SRAM**

→ Write final tile to DRAM once

★ A tile is a small portion of the rendered frame. In Phase 2—tile rendering—the GPU renders each tile and keeps its intermediate result in **SRAM**.

TBDR Rendering

✓ Rendering flow:

- Vertex Shader → Primitive Setup → **Tile-Based Culling and Clipping** → Rasterization → Fragment Shader

TBDR architectures depend on:

- predictable geometry counts,

- small on-chip tile memory,
- minimal external memory traffic.

△ As described in section *3D Rendering Pipeline*, Geometry Shader (GS) can generate both more vertices and more primitives than it receives. A single patch from Tessellation Control Shader (TCS) and Tessellation Evaluation Shader (TVS) can generate millions of micro-triangles. GS and Tessellation introduce **unbounded geometry amplification**, which breaks these assumptions and forces expensive DRAM spills for TBDR as shown in Fig. 5.14,




Fig. 5.14: CPU and GPU Pipeline For Shaders in Mobile Device

This is why Mali, PowerVR, Apple, and Adreno mobile GPUs all omit these stages^{28,29}.

Developers manually invoke **compute shaders** to generate meshlets or additional geometry, adding extra geometric detail for close-up **zoom-in** effects. Both Mali and PowerVR GPUs then run the standard vertex shader on the generated results.

✓ Step 1 —Developer dispatches a compute shader

This compute shader can do things like:

- break a **big mesh** into **meshlets** as Fig. 5.15. The mesh and meshlets are described in the *Mesh-Shader Pipeline* next section.
- generate **more vertices** for detail (subdivision, displacement)
- perform **LOD** selection
- cull invisible meshlets
- generate new index/vertex buffers

This is developer-controlled, not automatic.

The compute shader writes results into:

- SSBOs
- vertex buffers
- index buffers

These buffers now contain the final geometry you want to render.

✓ Step 2 —Developer issues a normal draw call

The Mali and PowerVR's rendering flow is illustrated as Fig. 5.16.

- Modern mobile engines instead use **compute shaders** for culling, LOD, meshlet prep, and procedural geometry.

²⁸ ARM Developer: *Why mobile GPUs avoid geometry shaders* <https://developer.arm.com/documentation/102476/latest/>

²⁹ Imagination Technologies: *Why Geometry Shaders Are Not Supported* <https://www.imgtec.com/blog/why-powervr-does-not-support-geometry-shaders/>




Fig. 5.15: CPU and GPU Pipeline For Shader's in Mobile Device

✓ Geometry Shaders are notoriously inefficient even on desktop GPUs. GPU vendors (NVIDIA + AMD + Intel) designed the mesh-shader pipeline described in the section *Mesh-Shader Pipeline*.

5.2.6 Mesh-Shader Pipeline

A single 3D object can contain 1 mesh, multiple meshes or hundreds of meshes (complex characters, vehicles, weapons).

Reasons

The purpose of converting a **mesh into small clusters (meshlets)** is to give the GPU **small, coherent**, cullable, cache-friendly work units that dramatically improve parallelism, memory locality, and LOD efficiency.

Raw meshes can have anywhere from thousands to millions of vertices/triangles, while meshlets intentionally restrict clusters to ~32–128 vertices and ~32–256 triangles to maximize GPU efficiency.

Motivation

NVIDIA, AMD, and Intel all needed:

- a compute-like geometry pipeline
- meshlet-based processing
- better culling
- GPU-driven rendering
- a **replacement** for VS → TCS → TES → GS
 - TCS → Fixed-Function Tessellator → TES: geometry amplification.
 - * Fixed-Function Tessellator: subdivides the patch, generates new domain coordinates and creates the tessellated grid.
 - * *Mesh-Shader replaces the fixed-function tessellator** with compute-like geometry pipeline.
 - * Mesh Shading provides a programmable alternative to the traditional tessellation and geometry-shader stages. A Mesh Shader can implement tessellation-like subdivision and geometry amplification while also integrating operations such as culling, LOD selection, and procedural geometry generation.

So the vendors co-designed the hardware pipeline.

✓ Microsoft and Khronos (Vulkan) each standardized it in their own APIs

✓ **Solution:** As shown in Fig. 5.17.

Rendering Pipeline

✓ GPU vendors (NVIDIA + AMD + Intel) designed the mesh-shader pipeline:




Fig. 5.16: CPU and GPU Pipeline For Shaders in Mobile Device




Fig. 5.17: Meshlet Offline To Render

3D Modeling Tool Output (big mesh)

→ CPU Meshlet Generator Tool (offline)

- Converting **big mesh** into **small clusters (meshlets)** to maximize GPU efficiency.

→ Precomputed meshlets (static clusters)

→ Task Shader (optional)

→ Mesh Shader

3D modeling tools do NOT generate meshlets. Meshlets are always generated later, using specialized meshlet-generation

→ tools, most commonly:

- NVIDIA meshlet generator (NV_mesh_shader ecosystem)
- meshoptimizer (Khronos-recommended, open source)
- Engine-specific meshlet builders (Unreal, Frostbite, etc.)

So the meshlet conversion happens after the model is exported —not inside Blender, Maya, 3ds Max, etc.

The animation flow from CPU to GPU for **Traditional**, **Compute Shader** based and **Mesh Shader** are shown in Fig. 5.18, Fig. 5.19 and Fig. 5.20.

Mesh shading (Vulkan VK_EXT_mesh_shader, similarly in NV mesh shader) replaces the fixed vertex-input + VS + optional tess/GS stages with a compute-like geometry pipeline:






Fig. 5.18: CPU and GPU Traditional Pipeline For Shaders

As in Fig. 5.20, NVIDIA/AMD desktop provide mesh-shader to do the following pipeline.

Task Shader Responsibilities

The **Task Shader** acts as a coarse-grained work distributor.

Key responsibilities:

Fig. 5.19: CPU and GPU **Mobile** Pipeline For ShadersFig. 5.20: CPU and GPU **Mesh Shader** Pipeline For Shaders

- Perform coarse culling at the meshlet or instance level.
- Select appropriate **LODs** for distant geometry.
 - But it does not create new detail like the Tessellation Shaders.
- Build a compact list of meshlets to be processed.
- Determine how many mesh shader workgroups to launch.
- Pass a payload (task data) to mesh shader workgroups.

The task shader does *not* emit vertices or primitives.

Mesh Shader Responsibilities

The **mesh shader** replaces the vertex shader, tessellation, and often the geometry shader. It operates on meshlets inside workgroups.

Key responsibilities:

- Load meshlet vertices and indices from GPU memory.
- Apply transforms, skinning, morphing, and procedural deformation.
- Mesh Shader outputs exactly what the meshlet contains usually, unless you code it manually.
 - Custom procedural code inside a Mesh Shader can generate more vertices, subdivide triangles, procedurally generate detail, amplify geometry. Mesh Shaders replace Vertex Shader, Geometry Shader and Tessellation (optional). But they do not perform automatic tessellation. They simply take a meshlet, run a workgroup and output the triangles inside that meshlet.
- Perform fine-grained culling: - frustum culling - backface culling - small triangle culling - cluster-level culling
- Generate the final set of vertices and primitives.
- Emit primitives directly to the rasterizer.

Because mesh shaders run in workgroups, they can use shared memory and synchronize threads, enabling efficient reuse of vertex data.

Why Meshlets Fit GPU Architecture Well

Meshlets align naturally with GPU hardware for several reasons:

- **Workgroup-Friendly:** Each meshlet maps cleanly to a single workgroup, keeping memory usage predictable and minimizing divergence.
- **Cache Efficiency:** Meshlets maximize vertex reuse and reduce memory bandwidth by grouping spatially local geometry.
- **Hierarchical Culling:** - Task shader: coarse culling of entire meshlets. - Mesh shader: fine culling of individual primitives.
- **Reduced CPU Overhead:** The GPU can perform culling, LOD selection, and primitive generation without CPU intervention, enabling GPU-driven rendering.
- **Scalable Parallelism:** Each meshlet is processed independently, allowing thousands of workgroups to run in parallel across GPU SMs.

Both mobile GPUs and mesh-shader GPUs convert a source mesh into smaller meshlets and render them efficiently using GPU SIMT execution and the memory hierarchy. The comparison is shown in the following table.

Comparison: Mobile GPU (Compute-Shader Based) vs. Desktop Mesh-Shader GPU

The **Mesh Shader** is similar to the previous section of **Mobile Compute Shader** based Meshlets as the following table:

Table 5.6: Mobile GPU vs Desktop Mesh-Shader GPU —Concept Comparison

Concept	Mobile GPU (Compute-Shader Based)	Desktop Mesh-Shader GPU
Meshlet generation	Compute Shader generates meshlets at runtime	CPU Meshlet Generator Tool (offline)
Tile-based	Yes	No
Work distribution	Compute Shader dispatch groups handle distribution	Task Shader distributes meshlet workloads
Meshlet expansion	Vertex Shader processes vertices after compute pre-processing	Mesh Shader expands meshlets and emits triangles
Culling & LOD	Compute Shader performs culling and LOD before raster	Task + Mesh Shader perform culling and LOD selection
Draw submission	Compute Shader writes indirect draw commands	Mesh Shader emits primitives directly to rasterizer
Pipeline family	Traditional Pipeline (VS → Raster → FS)	Mesh-Shader Pipeline (Task → Mesh → Raster → FS)

Summary

Meshlets and the mesh-shader pipeline transform geometry processing into a compute-like workflow. By organizing geometry into small, cache-friendly clusters and distributing work across task and mesh shaders, modern GPUs achieve higher throughput, better culling efficiency, and reduced CPU overhead compared to the traditional vertex-processing pipeline.

5.2.7 Animation Example

The skinning formula is described in *SW Stack and Data Flow* section as follows:

$$finalPosition = \sum_{i=0}^{N-1} weight_i (boneMatrix_i \cdot originalPosition)$$

The following code implements the formula shown above.

GLSL Vertex Shader

Listing 5.1: Example GPU skinning

```
layout(location = 0) in vec3 position;
layout(location = 1) in uvec4 boneIndex;
layout(location = 2) in vec4 boneWeights;

// Simple Uniforms (non-UBO)
uniform mat4 boneMatrices[100];
uniform mat4 model;
uniform mat4 view;
uniform mat4 projection;

vec4 skinnedPos = vec4(0.0);
for (int i = 0; i < 4; ++i) {
    skinnedPos += boneMatrices[boneIndex[i]] * vec4(position, 1.0) * boneWeight[i];
}
gl_Position = projection * view * model * skinnedPos;
```

Here:

- **position, boneIndex, boneWeight = vertex attributes**
- **boneMatrices, model, view, projection = uniforms**
- The OpenGL code used to pass these variables to GLSL will be shown in *OpenGL API Commands That Trigger GPU Skinning* later. The OpenGL API sets position, boneIndex and boneWeight to locations 0, 1 and 2, respectively, using **glVertexAttribPointer**.
 - void glVertexAttribPointer(**GLuint index**, **GLint size**, **GLenum type**, GLboolean normalized, GLsizei stride, const GLvoid * pointer);

Examples:

- * glBindBuffer(GL_ARRAY_BUFFER, vboPositions);
glVertexAttribPointer(**0, 3, GL_FLOAT, GL_FALSE**, stride, offset); → layout(**location = 0**) in **vec3** position;
- * glBindBuffer(GL_ARRAY_BUFFER, vboBoneIndex);
glVertexAttribPointer(**1, 4, GL_UNSIGNED_BYTE**, stride, offset); → layout(**location = 1**) in **uvec4** boneIndex;
 - Bone indices are small integers (0–255), so storing them as GL_UNSIGNED_BYTE: OpenGL will automatically zero-extend 8-bit unsigned integers into 32-bit unsigned integers inside the shader.

✓ Why boneIndex[] and boneWeight[] are 3D Model Information

These two arrays describe how the mesh is bound to the skeleton.

They are part of the static mesh data, created during rigging in Blender/Maya/etc.

boneIndex[] → tells which bone

- For each vertex: which bones influence it
- Example: { 3, 7, 12, 0 }

boneWeight[] → tells how much

- For each vertex: how much each bone influences it
- Example: { 0.5, 0.3, 0.2, 0.0 }

These values never change during animation. They are baked into the mesh and stored in the VBO as vertex attributes.

✓ Why boneMatrices[] is Animation Parameters

boneMatrices[] → tells where the bone is this frame

- Example: boneMatrices[3] (upper arm bone this frame)

[0.87 -0.49 0.00 0.12] [0.49 0.87 0.00 0.03] [0.00 0.00 1.00 0.00] [0.00 0.00 0.00 1.00]

This matrix might represent:

- a 30° rotation of the upper arm
- plus a small translation (0.12, 0.03, 0.0)

Animation Parameters are dynamic per-frame data, such as:

- bone matrices
- animation time
- morph weights
- blend factors
- procedural animation inputs

These change every frame.

✓ OpenGL API Commands That Trigger GPU Skinning

Overview

In OpenGL, animation is not built into the API. Instead, animation occurs because the application updates *Animation Parameters* (such as bone matrices) and the *vertex shader* interprets them. The GPU performs the animation math during the draw call.

The following sections describe the exact OpenGL commands involved in triggering GPU-based vertex animation.

1. Updating Animation Parameters (Uniforms or UBOs)

Animation Parameters such as boneMatrices[] are updated every frame. They are supplied to the vertex shader as uniforms or through a uniform buffer object (UBO).

Uniform array example:

```
// Matrix Uniforms
GLint locModel = glGetUniformLocation(program, "model");
GLint locView  = glGetUniformLocation(program, "view");
GLint locProj  = glGetUniformLocation(program, "proj");

glUniformMatrix4fv(locModel, 1, GL_FALSE, glm::value_ptr(modelMatrix));
glUniformMatrix4fv(locView,  1, GL_FALSE, glm::value_ptr(viewMatrix));
glUniformMatrix4fv(locProj,  1, GL_FALSE, glm::value_ptr(projMatrix));

// Bone Matrix Array
GLint loc = glGetUniformLocation(program, "boneMatrices");
glUniformMatrix4fv(loc, boneCount, GL_FALSE, boneMatrixData);
```

Uniform Buffer Object example:

```
glBindBuffer(GL_UNIFORM_BUFFER, boneUBO);
glBufferSubData(GL_UNIFORM_BUFFER, 0, size, boneMatrixData);
glBindBufferBase(GL_UNIFORM_BUFFER, bindingPoint, boneUBO);
```

These commands send the per-frame animation data to the GPU.

2. Binding Vertex Data (Mesh Information)

Static mesh data such as positions, normals, `boneIndex[]` and `boneWeight[]` is stored in vertex buffer objects (VBOs) and attached to a vertex array object (VAO).

```
glBindVertexArray(vao);

glBindBuffer(GL_ARRAY_BUFFER, vboPositions);
glVertexAttribPointer(0, 3, GL_FLOAT, GL_FALSE, stride, offset);

// Activate attribute location 1, then the shader's layout(location = 1)
// input receives real data.
glEnableVertexAttribArray(0);

glBindBuffer(GL_ARRAY_BUFFER, vboBoneIndex);
glVertexAttribPointer(1, 4, GL_UNSIGNED_BYTE, stride, offset);
glEnableVertexAttribArray(1);

glBindBuffer(GL_ARRAY_BUFFER, vboBoneWeight);
glVertexAttribPointer(2, 4, GL_FLOAT, GL_FALSE, stride, offset);
glEnableVertexAttribArray(2);
```

These commands provide the static 3D model information to the vertex shader.

3. Activating the Shader Program

The vertex shader containing the skinning logic must be activated before drawing.

```
glUseProgram(program);
```

This step ensures that the GPU will execute the correct vertex shader when the draw call is issued.

4. Issuing the Draw Call (Animation Trigger)

The draw call is the moment when the GPU executes the vertex shader for each vertex. This is where animation actually happens.

```
glDrawElements(GL_TRIANGLES, indexCount, GL_UNSIGNED_INT, 0);
```

or:

```
glDrawArrays(GL_TRIANGLES, 0, vertexCount);
```

The vertex shader runs once per vertex, combining:

- vertex attributes (`position`, `boneIndex[]`, `boneWeight[]`)
- animation parameters (`boneMatrices[]`)

to compute the animated vertex position.

Summary Table

Purpose	Data Type	OpenGL API	Static or Dynamic
Mesh data (positions, bone indices, bone weights)	Vertex Attributes	glVertexAttribPointer glEnableVertexAttribArray	Static (stored in VBO)
Animation Parameters (bone matrices)	Uniforms / UBO	glUniformMatrix4fv glBufferSubData	Dynamic (updated every frame)
Activate shader program	Shader Program	glUseProgram	Per draw
Trigger animation	Draw Call	glDrawElements / glDrawArrays	Per frame

Conclusion

OpenGL does not provide a built-in animation system. Instead, animation occurs because the application updates Animation Parameters each frame and the vertex shader applies animation math during the draw call. The GPU performs the animation only when the draw command is issued.

5.2.8 Ray Tracing Pipeline




Fig. 5.21: Relationships between the texturing concept³⁶.

The key conceptual difference between rasterization and ray tracing is the direction of processing.

Rasterization starts from **geometry** and determines which pixels are covered by each triangle.

Ray tracing starts from **pixels** and determines which object is visible through each pixel.

Typically, each ray must be tested for intersection with some subset of all the objects in the scene. Once the nearest object has been identified, the algorithm will estimate the incoming light at the point of intersection, examine the material properties of the object, and combine this information to calculate the final color of the pixel³⁶.

Ray Tracing (Pixel → Object)

Ray tracing reverses this process.

```
Camera
|
▼
For every pixel
|
```

(continues on next page)

³⁶ <[<https://en.wikipedia.org/wiki/Ray_tracing_\(graphics\)>](https://en.wikipedia.org/wiki/Ray_tracing_(graphics))>

(continued from previous page)

```

▼
Generate Viewing Ray
|
▼
Traverse BVH
|
▼
Find Closest Intersection
|
▼
Evaluate Material
|
▼
Pixel Color

```

Instead of asking

Which pixels belong to this triangle?

ray tracing asks

What object is visible through this pixel?

Primary Ray Generation

For an image of size

1920 × 1080

the renderer generates approximately

2,073,600

primary rays (one per pixel).

For pixel (x, y):

```

origin = camera_position

direction =
    normalize(pixel_position - camera_position)

```

Each ray represents the viewing direction through one screen pixel.

Finding the Visible Object

Suppose the scene contains

```

Triangle A
Triangle B
Sphere
Floor

```

The generated ray is tested against the acceleration structure (typically a Bounding Volume Hierarchy, or BVH).

```
Ray
|
| Triangle A ? no
|
| Triangle B ? yes
|           distance = 8.2
|
| Sphere ?    yes
|           distance = 15.4
|
| Floor ?     yes
|           distance = 32
```

The closest intersection determines the visible object.

```
Visible Object = Triangle B
```

Material Evaluation

Once the closest object is found, its material is evaluated.

For example,

```
Diffuse Texture
Normal Map
Roughness
Metallic
Emissive
```

Lighting calculations are then performed similarly to a fragment shader, except that the visible surface was determined by ray intersection instead of rasterization.

Texture Coordinate Interpolation

Each triangle still stores per-vertex attributes.

```
Vertex 0
  position
  normal
  uv

Vertex 1
  position
  normal
  uv

Vertex 2
  position
  normal
  uv
```

After the ray intersects the triangle, barycentric coordinates are computed.

For example,

```
u = 0.2
v = 0.3
w = 0.5
```

The texture coordinates at the hit point are interpolated as

```
uv =
  u * uv0 +
  v * uv1 +
  w * uv2
```

The interpolated texture coordinates are then used to sample textures exactly as in rasterization.

Comparison

Rasterization	Ray Tracing
Process one triangle	Process one pixel
Triangle → Pixels	Pixel → Ray
Rasterizer finds covered pixels	BVH traversal finds intersected object
Interpolate vertex attributes	Interpolate attributes at the hit point
Sample textures	Sample textures
Compute lighting	Compute lighting
Write pixel	Write pixel

Summary

Ray tracing uses the same geometric and material information as rasterization, including:

- Triangle meshes
- Vertex normals
- Texture coordinates
- Materials
- Textures

The primary difference lies in how the visible surface is determined.

- Rasterization determines visibility by projecting triangles onto the screen and finding the pixels they cover.
- Ray tracing determines visibility by casting a viewing ray through each pixel and finding the closest surface intersected by that ray.

After the visible surface has been identified, the remaining shading process (interpolating attributes, sampling textures, and evaluating materials) is very similar to traditional OpenGL fragment shading.

5.3 GLSL (GL Shader Language)

OpenGL is a standard specification for designing 2D and 3D graphics and animation in computer graphics. To support advanced animation and rendering, OpenGL provides a large set of APIs (functions) for graphics processing. Popular 3D modeling and animation tools—such as Maya, Blender, and others—can utilize these APIs to handle 3D-to-2D projection and rendering directly on the computer.

The hardware-specific implementation of these APIs is provided by GPU manufacturers, ensuring that rendering is optimized for the underlying hardware.

5.3.1 Background

In the previous section *SW Stack and Data Flow* described **how each frame is generated** to display the **movement animation or skinning effects** using the small animation parameters stored in 3D model and sent from CPU.

Based on description of section *SW Stack and Data Flow*, we know the animation can be implemented using **fixed-function skinning**. The following are the animation examples for **shader-less era**.

- ✓ Some consoles and mobile GPUs did have fixed-function skinning.
- ✓ In those systems, you could upload bone matrices and let hardware animate vertices.

✗ **But you could not change the formulas —only use the built-in ones.**

The following console GPUs did have fixed-function skinning:

PlayStation 2 (PS2) —VU0/VU1 Microcode

PS2 had fixed hardware instructions for:

- skinning
- morphing
- matrix blending

Developers could upload bone matrices and let the hardware do the blending. No shaders existed yet.

Nintendo GameCube / Wii —XF Unit

The GameCube GPU had a fixed-function transform unit that supported:

- matrix palette skinning (up to 10 matrices)
- per-vertex weighted blending

Again, no shaders —but hardware skinning existed.

The previous *section Role and Purpose of Shaders* also explained how different visual effects can be achieved by **switching shaders** and applying different materials across frames.

✗ However **the fixed-function pipeline (OpenGL 1.x / early 2.x without shaders)** has:

- no per-vertex programmable math
- no access to bone matrices
- no ability to blend multiple positions
- no ability to apply time-based deformation
- no ability to read custom vertex attributes
- no ability to modify vertex positions except via the model-view matrix

✗ As a result, the shader-less (fixed-function) pipeline in early OpenGL did not support GPU-based skinning. Skinning had to be implemented on the CPU, which imposed limitations on both **animation capability and performance**, as described below:

Major Disadvantages of a Shader-less (Fixed-Function) Pipeline

- **No GPU-side animation**
 - Cannot perform skinning, morphing, or procedural deformation on the GPU.
 - All animation must be computed on the CPU, causing performance bottlenecks.
- **Limited lighting and materials**
 - Only fixed-function lighting is available.

- No custom BRDFs, PBR workflows, toon shading, or stylized effects.
 - **No procedural or time-based effects**
 - Cannot implement UV animation, distortion, dissolve, holograms, or particle effects.
 - No access to noise functions or time-driven logic in the pipeline.
 - **No post-processing**
 - Motion blur, bloom, depth of field, color grading, and screen-space effects are impossible.
 - **Rigid data flow**
 - Cannot define custom vertex attributes, varyings, or uniform buffers.
 - Material and animation systems cannot be data-driven.
 - **Poor scalability and performance**
 - CPU must update all animated geometry every frame.
 - GPU parallelism is unused, limiting scene complexity.
 - **Deprecated and non-portable**
 - Fixed-function pipeline is removed in modern OpenGL core profiles.
 - Not compatible with contemporary engines or hardware.
- ✓ **All modern consoles** (PS5, PS5 Pro, PS6-class hardware of Sony, Switch 2 of Nintendo) use **programmable shader architectures** rather than fixed-function animation hardware. Fixed-function animation is now **obsolete**.
- Sony's current and upcoming GPUs are based on AMD RDNA architectures, which are fully programmable shader GPUs.
 - Nintendo's upcoming Switch 2 uses a custom Nvidia Ampere GPU.

5.3.2 Examples

An OpenGL program typically follows a structure like the example below:

Vertex shader

```
#version 330 core
layout (location = 0) in vec3 aPos; // the position variable has attribute position 0

out vec4 vertexColor; // specify a color output to the fragment shader

void main()
{
    gl_Position = vec4(aPos, 1.0); // see how we directly give a vec3 to vec4's_
    ↪ constructor
    vertexColor = vec4(0.5, 0.0, 0.0, 1.0); // set the output variable to a dark-red_
    ↪ color
}
```

Fragment shader

```
#version 330 core
out vec4 FragColor;

in vec4 vertexColor; // the input variable from the vertex shader (same name and same
↳type)

void main()
{
    FragColor = computeColorOfThisPixel(...);
}
```

OpenGL user program

```
int main(int argc, char ** argv)
{
    // init window, detect user input and do corresponding animation by calling opengl
↳api
    ...
}
```

The last *main()* function in an OpenGL application is written by the user, as expected. Now, let's explain the purpose of the first two main components of the OpenGL pipeline.

As discussed in the *Concepts of Computer Graphics* textbook, OpenGL provides a rich set of APIs that allow programmers to render 3D objects onto a 2D computer screen. The general rendering process follows these steps:

1. The user sets up lighting, textures, and object materials.
2. The system calculates the position of each vertex in 3D space.
3. The GPU and rendering pipeline automatically determine the color of each pixel based on lighting, textures, and interpolation.
4. The final image is displayed on the screen by writing pixel colors to the framebuffer.

To give programmers the flexibility to add custom effects or visual enhancements—such as modifying vertex positions for animation or applying unique coloring—OpenGL provides two programmable stages in the graphics pipeline:

- **Vertex Shader:** Allows the user to customize how vertex coordinates are transformed and processed.
- **Fragment Shader:** Allows the user to define how each pixel (fragment) is shaded and colored, enabling effects like lighting, textures, and transparency.

These shaders are written by the user and compiled at runtime, providing powerful control over the rendering process.

OpenGL uses fragment shader instead of pixel is : “Fragment shaders are a more accurate name for the same functionality as Pixel shaders. They aren't pixels yet, since the output still has to pass several tests (depth, alpha, stencil) as well as the fact that one may be using antialiasing, which renders one-fragment-to-one-pixel non-true³⁰. Programmer is allowed to add their converting functions that compiler translate them into GPU instructions running on GPU processor. With these two shaders, new features have been added to allow for increased flexibility in the rendering pipeline at the vertex and fragment level³¹. Unlike the shaders example here³², some converting functions for coordinate in vertex shader or for color in fragment shade are more complicated according the scenes of animation. Here is an example³³. In wiki

³⁰ <https://community.khronos.org/t/pixel-vs-fragment-shader/52838>

³¹ https://en.m.wikipedia.org/wiki/OpenGL_Shading_Language

³² <https://learnopengl.com/Getting-started/Shaders>

³³ <https://www.youtube.com/watch?v=LyoSSoYyfVU> at 5:25 from beginning: combine different textures.

shading page⁴, Gouraud and Phong shading methods make the surface of object more smooth by glsl. Example glsl code of Gouraud and Phong shading on OpenGL api are here³⁴. Since the hardware of graphic card and software graphic driver can be replaced, the compiler is run on-line meaning driver will compile the shaders program when it is run at first time and kept in cache after compilation³⁵.

The shaders program is C-like syntax and can be compiled in few mini-seconds, add up this few mini-seconds of on-line compilation time in running OpenGL program is a good choice for dealing the cases of driver software or gpu hardware replacement³⁷.

5.3.3 Goals

Goals of GLSL Shader Language:

GLSL was designed for real-time graphics using programmable GPUs.

1. Programmable Pipeline:
 - Custom control over vertex, fragment, and other pipeline stages
 - Enables dynamic effects, lighting, animation, and transformations
2. GPU Acceleration
 - Executes on GPU cores for massive parallel performance
 - Optimized for matrix and vector operations common in graphics
3. Cross-Platform Compatibility:
 - Runs consistently across OSes and hardware via OpenGL
 - Avoids vendor lock-in for portable shader code
4. C-Like Syntax
 - Familiar syntax for developers used to C-style languages
 - Supports functions, loops, conditionals, and custom types
5. Fine-Grained Rendering Control
 - Direct access to geometry, color, texture, lighting parameters
 - Enables advanced effects like shadows, fog, reflections
6. Real-Time Interactivity
 - Responds to user input, time, and animations at runtime
 - Suitable for games, simulations, and creative tools
7. Minimal Host Dependency
 - Executes within the graphics driver context
 - No need for external libraries, file I/O, or system calls

⁴ <https://en.wikipedia.org/wiki/Shading>

³⁴ <https://github.com/ruange/Gouraud-Shading-and-Phong-Shading>

³⁵ Compiler and interpreter: (<https://www.guru99.com/difference-compiler-vs-interpreter.html>). AOT compiler: compiles before running; JIT compiler: compiles while running; interpreter: runs (reference <https://softwareengineering.stackexchange.com/questions/246094/understanding-the-differences-traditional-interpreter-jit-compiler-jit-interp>). Both online and offline compiler are AOT compiler. User call OpenGL api to run their program and the driver call call online compiler to compile user's shaders without user compiling their shader before running their program. When user run a CPU program of C language, he must compile C program before running the program. This is offline compiler.

³⁷ <https://community.khronos.org/t/offline-glsl-compilation/61784>

5.3.4 GLSL vs. C: Feature Overview

GLSL expands upon C for GPU-based graphics programming.

Additions to C:

1. Specialized Data Types
 - `vec2`, `vec3`, `vec4`: float vectors
 - `mat2`, `mat3`, `mat4`: float matrices
 - `bvec`, `ivec`, `uvec`, `dvec`: boolean and integer vectors
 - `sampler2D`, `samplerCube`: texture samplers
2. Pipeline Qualifiers
 - `attribute`, `varying` (legacy)
 - `in`, `out`, `inout`: stage and parameter I/O
 - **uniform**: uniform variables are set externally by the host application (e.g., OpenGL) and remain constant across all shader invocations for a draw call.
 - **layout(location = x)**: set GPU variable locations. See *Animation Example* section.
 - precision qualifiers: `lowp`, `mediump`, `highp`
3. Built-in Functions
 - `texture()`, `reflect()`, `refract()`, `normalize()`
 - `mix()`, `smoothstep()`: interpolation and blending
 - `dot()`, `cross()`, `transpose()`, `inverse()`: math ops
 - `dFdx()`, `dFdy()`, `fwidth()`: pixel derivatives
4. Swizzling
 - `.xyzw`, `.rgba`, `.stpq` access vector components
 - e.g., `vec4 pos = vec3(1, 2, 3).xyzx`
5. Shader-Specific Keywords
 - `discard`: drop fragments early
 - `gl_Position`, `gl_FragColor`, `gl_VertexID`: built-ins
 - `subroutine`, `patch`, `sample`: advanced pipeline control

Removals and Restrictions:

1. No Pointers or Memory Access
 - No `*` or `&` operators
 - No `malloc`, `free`
2. No File I/O or Standard C Libs
 - No `stdio.h`, `printf()`, `fopen()`
3. No Recursion
 - Recursive functions not allowed
4. No `#include` Support

- Files can't be included via preprocessor

5. Limited Control Flow

- goto not allowed
- Loops must be statically determinable in many cases for compiler optimization as follows:

Example for loops must be statically determinable in many cases

```
const int MAX_LIGHTS = 10;
for (int i = 0; i < MAX_LIGHTS; ++i) {
    // Safe: MAX_LIGHTS is a compile-time constant
}
```

6. Restricted C Keywords

- typedef, union, enum, class, namespace, inline, etc.
- Reserved or disallowed

Notes:

- Changes help GPU execute safely in parallel
- Designed for real-time, interactive graphics

5.3.5 GLSL Qualifiers by Shader Stage

The CPU and GPU Pipeline For Shaders is introduced in section *3D Rendering Pipeline*. The Fig. 5.22 is the summary of GLSL Qualifiers below.




Fig. 5.22: Shaders input and output

Vertex Shader:

- in: Receives per-vertex attributes from buffer objects, it is **3D Model Information** described in Fig. 5.22.
- out: Passes data to next stage (e.g., fragment shader)
- **uniform**: Global parameters like matrices or lighting, it is **Animation Parameters**, also referred to as **Uniform Updates** described in Fig. 5.22. The *Animation Example* section provides an example in Vertex Shader (VS) that demonstrates animation using both the 3D model information and Uniform Updates.
- **layout(location = x)**: Binds input/output to attribute index. See *Animation Example* section.
- **const**: Compile-time constants
- Cannot use interpolation qualifiers on inputs

Fragment Shader:

- in: Receives **interpolated data** from previous stage as shown in Fig. 5.22.
- out: Writes **Final Fragment Color** to FrameBuffer
- uniform: Global parameters like **textures** or **lighting** as shown in Fig. 5.22. **Uniform data remains unchanged** across all pipeline stages and is shared by all shaders in the pipeline. This means that uniform data represents global parameters for 3D GPU rendering.
- flat: Disables interpolation; uses provoking vertex
- **smooth**: Enables perspective-correct interpolation (default); **noperspective**: Use Linear interpolation in screen space instead of “smooth” option.
- centroid: Samples within primitive area (for multisampling)
- sample: Per-sample interpolation (GLSL 4.0+)
- discard: Terminates fragment processing early

Compute Shader:

- layout(local_size_x = x): Defines workgroup size
- uniform: Input parameters from host
- buffer: Shader storage buffer access
- shared: Shared memory among invocations in a workgroup
- image2D, image3D: Direct image access
- coherent, volatile, restrict: Memory access control
- readonly, writeonly: Access mode for image/buffer
- Compute shader: **may be applied in any stage** as described in section *3D Rendering Pipeline*.

Common Across Stages:

- const: Immutable values
- uniform: Host-set global parameters
- layout(binding = x): Bind uniform/buffer/image to index
- precise: Ensures consistent computation
- invariant: Prevents variation across shader executions

Notes:

- attribute and varying are deprecated (use in/out instead)
- Interpolation qualifiers only affect fragment shader inputs
- Uniforms are shared across all stages and remain constant

Examples of GLSL Qualifiers by Shader Stage

```
// =====  
// Vertex Shader: Qualifier Summary (GLSL)  
// =====  
  
// Vertex inputs  
layout(location = 0) in vec3 aPosition;    // in: per-vertex attribute
```

(continues on next page)

(continued from previous page)

```

layout(location = 1) in vec3 aNormal;

// Outputs to fragment shader
out vec3 vNormal;                                // out: passes to next stage

// Uniforms
uniform mat4 uModelMatrix;                        // uniform: global parameter
uniform mat4 uViewProjectionMatrix;

// Constants
const float PI = 3.14159265;                      // const: compile-time constant

void main() {
    vNormal = aNormal;
    gl_Position = uViewProjectionMatrix * uModelMatrix * vec4(aPosition, 1.0);
}

// =====
// Fragment Shader: Qualifier Summary (GLSL)
// =====

// Inputs from vertex shader
in vec3 vNormal;                                // in: interpolated input

// Output to framebuffer
out vec4 fragColor;                              // out: final pixel color

// Uniforms
uniform vec3 uLightDirection;                    // uniform: shared global input
uniform vec3 uBaseColor;

// Interpolation control
// flat in vec3 vFlatColor;                        // flat: no interpolation
// smooth in vec3 vSmoothColor;                    // smooth: default interpolation
// noperspective in vec3 vLinearColor;             // noperspective: screen-space linear

void main() {
    float brightness = max(dot(normalize(vNormal), uLightDirection), 0.0);
    fragColor = vec4(uBaseColor * brightness, 1.0);
}

// =====
// Compute Shader: Qualifier Summary (GLSL)
// =====

#version 430

// Workgroup size
layout(local_size_x = 16, local_size_y = 16) in;

// Shared memory
shared float tileData[256];                      // shared: intra-group memory

```

(continues on next page)

(continued from previous page)

```
// Uniforms
uniform float uTime;                // uniform: global input

// Buffer access
layout(std430, binding = 0) buffer DataBuffer {
    float values[];
};

// Image access
layout(binding = 1, rgba32f) uniform image2D uImage;

// Memory qualifiers
// coherent, volatile, restrict, readonly, writeonly

void main() {
    uint idx = gl_GlobalInvocationID.x;
    values[idx] += sin(uTime);        // buffer write
    imageStore(uImage, ivec2(idx, 0), vec4(values[idx])); // image write
}
```

5.4 OpenGL Buffers

CPU and GPU provides different Buffers to speedup OpenGL pipeline rendering¹⁶.

¹⁶ Page 155 - 185 of book "OpenGL Programming Guide 9th Edition" Page 53, 1.

Table 5.7: Graphics Buffers

Buffer Type	Access	Location	API/Usage	Function	Description
Vertex Buffer (VBO)	Read	GPU	OpenGL, Vulkan	Store vertex attributes	Holds data like position, normal, and texture coords for drawing geometry.
Index Buffer (IBO/EBO)	Read	GPU	OpenGL, Vulkan	Reuse vertex data	Stores indices into the vertex buffer to avoid duplication.
Uniform Buffer (UBO)	Read	GPU or Shared	OpenGL, Vulkan	Constant input data	Shares transformation matrices, lighting, or material data across shaders.
Shader Storage Buffer (SSBO)	Read/Write	GPU or Shared	OpenGL, Vulkan	General data exchange	Flexible, large buffers accessible for structured shader I/O.
Constant Buffer	Read	GPU or Shared	DirectX, Vulkan	Fast uniform access	Optimized for fast access to frequently read small data.
Image / Texture Buffer	Read/Write	GPU	OpenGL, Vulkan	Sample/store pixels	Stores image data for sampling or read/write image operations in shaders.
Color Buffer	Write	GPU	OpenGL, Vulkan	Store final pixel color	Stores output of fragment shaders; used for display or post-processing.
Depth Buffer (Z-Buffer)	Write/Read	GPU	OpenGL, Vulkan	Visibility testing	Stores per-pixel depth values for hidden surface removal.
Frame Buffer	Write	GPU	OpenGL, Vulkan	Store render output	Holds final color, depth, or other rendered output.
Stencil Buffer	Read/Write	GPU	OpenGL, Vulkan	Pixel masking	Used to conditionally discard or preserve pixels in the pipeline.

✓ Uniforms, UBOs, and SSBOs

This section explains the three major data-passing mechanisms in OpenGL: `uniform` variables, `Uniform Buffer Objects` (UBOs), and `Shader Storage Buffer Objects` (SSBOs). They differ in storage location, capacity, access rules, and intended usage.

1. Uniform (Default Uniform Block)

A *uniform* is a small piece of read-only data provided to shaders. Examples:

```
uniform mat4 view;
uniform vec3 lightPos;
uniform float exposure;
```

Characteristics:

- Stored inside the *Default Uniform Block* of the program object.
- Managed entirely by the OpenGL driver.
- Updated using `glUniform*` functions.
- Not a GPU buffer object.
- **Not mappable and not shareable across programs.**
- Uses opaque, driver-defined layout rules.
- Best suited for small parameters such as scalars, vectors, and matrices.

Typical size limit: approximately 16 KB (implementation dependent).

2. Uniform Buffer Object (UBO)

A *Uniform Buffer Object* stores uniforms declared inside a named uniform block:

```
layout(std140) uniform Camera {  
    mat4 view;  
    mat4 proj;  
    vec3 eyePos;  
};
```

UBOs are real GPU buffer objects allocated by the application:

```
glGenBuffers(1, &ubo);  
glBindBuffer(GL_UNIFORM_BUFFER, ubo);  
glBufferData(GL_UNIFORM_BUFFER, size, data, GL_STATIC_DRAW);
```

Characteristics:

- Lives in GPU VRAM as a buffer object.
- Read-only in shaders.
- Uses `std140` or `std430` structured layouts.
- Bindable to buffer binding points.
- **Shareable across multiple shader programs.**
- Mappable and updatable through buffer APIs.
- Typically limited to about 64 KB per block due to hardware constant-cache constraints.

Best suited for medium-sized structured data such as camera matrices, lighting parameters, and per-frame or per-pass data.

3. Shader Storage Buffer Object (SSBO)

A *Shader Storage Buffer Object* is a general-purpose GPU buffer accessible for both reading and writing from shaders:

```
layout(std430, binding = 0) buffer Storage {  
    vec4 positions[];  
};
```

SSBOs are allocated similarly to other buffer objects:

```
glBindBuffer(GL_SHADER_STORAGE_BUFFER, ssbo);  
glBufferData(GL_SHADER_STORAGE_BUFFER, size, nullptr, GL_DYNAMIC_COPY);
```

Characteristics:

- Lives in GPU VRAM as a buffer object.
- Readable and writable from shaders.
- Supports very large sizes (hundreds of megabytes).
- Uses `std430` layout for compact packing.
- Suitable for compute workloads, GPU-driven rendering, and large dynamic datasets.

Best suited for large arrays, particle systems, visibility buffers, and any data too large or too dynamic for UBOs.

Comparison Table

The following table summarizes the differences:

Feature	Uniform	UBO	SSBO
Storage	Program object	GPU buffer	GPU buffer
Access	Read-only	Read-only	Read/write
Update	<code>glUniform*</code>	Buffer update	Buffer update
Layout	Opaque	<code>std140/std430</code>	<code>std430</code>
Max size	~16 KB	~64 KB	Very large (MBs)
Share across programs	No	Yes	Yes
Shader write	No	No	Yes
Best for	Small parameters	Medium structured data	Large dynamic data

✓ `std140` and `std430` Layout Rules

OpenGL defines two major memory layout standards for uniform and storage buffer blocks: `std140` and `std430`. These rules determine how GLSL types are aligned and padded inside `uniform` blocks (UBOs) and `buffer` blocks (SSBOs). Understanding these layouts is essential for correctly matching GPU-side GLSL structures with CPU-side data structures. **`Std140` is more performance-efficient and `std430` is more memory-efficient.**

`std140` Layout

`std140` is the default layout for `uniform` blocks and is designed for compatibility with older hardware. It enforces strict alignment and padding rules to ensure predictable access through the GPU's constant cache.

Key properties:

- Scalars (`float`, `int`) have 4-byte alignment.
- `vec2` has 8-byte alignment.
- `vec3` and `vec4` both have 16-byte alignment.
- Matrices are stored as arrays of column vectors, each column aligned to 16 bytes.
- Arrays have a stride of 16 bytes per element, regardless of element type.
- Structs are padded so that each member follows the above rules, and the struct itself is aligned to 16 bytes.

Example:

```
layout(std140) uniform Camera {
    mat4 view;      // 64 bytes
    mat4 proj;      // 64 bytes
    vec3 eyePos;    // 16 bytes (vec3 padded to vec4)
    float exposure; // 4 bytes, but padded to 16 bytes
};
```

`std140` is predictable but often wastes memory due to padding.

`std430` Layout

`std430` is the default layout for `buffer` blocks (SSBOs) and is designed for modern GPUs. It relaxes many of the padding rules found in `std140` and allows tighter packing of data.

Key properties:

- Scalars have 4-byte alignment.
- `vec2` has 8-byte alignment.
- `vec3` and `vec4` have 16-byte alignment (same as `std140`).
- Arrays use the natural alignment of their element type (no forced 16-byte stride).
- Structs follow natural alignment rules without extra padding beyond what each member requires.
- Matrices are still stored as arrays of column vectors, but column alignment follows natural rules.

Example:

```
layout(std430, binding = 0) buffer Storage {
    vec4 positions[]; // tightly packed array of vec4
};
```

`std430` is more memory-efficient and is preferred for large datasets such as particle systems, mesh data, and compute workloads.

Comparison Table

Feature	<code>std140</code>	<code>std430</code>
Used in	UBOs (uniform blocks)	SSBOs (buffer blocks)
Padding	Heavy	Minimal
Array stride	Always 16 bytes	Natural alignment
Struct alignment	16 bytes	Natural alignment
Best for	Small uniform data	Large dynamic data

Summary

- `std140` provides strict, predictable alignment for uniform blocks, but often wastes memory.
- `std430` provides compact, efficient layouts for storage buffers, ideal for large arrays and compute workloads.
- Both layouts ensure that CPU-side data structures can be matched reliably with GPU-side GLSL declarations.
- Color buffer

They contain the RGB or sRGB color data and may also contain alpha values for each pixel in the framebuffer. There may be multiple color buffers in a framebuffer. You've already used double buffering for animation. Double buffering is done by making the main color buffer have two parts: a front buffer that's displayed in your window; and a back buffer, which is where you render the new image¹⁷.

- Depth buffer (Z buffer)

Depth is measured in terms of distance to the eye, so pixels with larger depth-buffer values are overwritten by pixels with smaller values¹⁸¹⁹²⁰.

¹⁷ Page 155 of book "OpenGL Programming Guide 9th Edition"^{Page 53, 1}.

¹⁸ Page 156 of book "OpenGL Programming Guide 9th Edition"^{Page 53, 1}.

¹⁹ <https://en.wikipedia.org/wiki/Z-buffering>

²⁰ <https://open.gl/depthstencils>

- Frame Buffer

OpenGL offers: the color, depth and stencil buffers. This combination of buffers is known as the default framebuffer and as you've seen, a framebuffer is an area in memory that can be rendered to²¹.

- Stencil Buffer

In the simplest case, the stencil buffer is used to limit the area of rendering (stenciling)²²Page 106, 20.

Table 5.8: Compute Buffers

Buffer Type	Access	Location	API/Usage	Function	Description
Compute Buffer	Read/Writ	GPU or Shared	OpenCL, Vulkan, CUDA	Parallel compute data	Buffers used in compute kernels or shaders for general processing.
Atomic Buffer	Read/Writ (Atomic)	GPU	OpenGL, Vulkan	Shared counters/data	Used with atomic ops for synchronization or accumulation.
Acceleration Structure Buffer	Read	GPU	Vulkan RT, DXR	Ray tracing acceleration	Holds spatial hierarchy (BVH) for ray traversal efficiency.
Indirect Draw Buffer	Read	GPU	Vulkan, DirectX	GPU-issued draw	Stores draw/dispatch args to issue commands without CPU.

- DXR: DirectX Raytracing —a D3D12 extension for real-time ray tracing using GPU acceleration.
- Indirect Draw Buffer: A GPU-side buffer holding draw parameters so that GPU (not CPU) can issue rendering work dynamically.

Table 5.9: System-Level and Utility Buffers

Buffer Type	Access	Location	API/Usage	Function	Description
Command Buffer	Write (CPU) / Read (GPU)	Host → GPU	Vulkan, DirectX12	Submit work	Encapsulates commands like draw, dispatch, and memory ops.
Parking / Staging Buffer	Read/Writ	Host-visible	Vulkan, CUDA	Temporary transfer	Temporary CPU-visible buffer for uploading/downloading GPU data.

5.5 OpenGL Shader Compiler

The OpenGL standard is defined in³⁸. OpenGL is primarily designed for desktop computers and servers, whereas OpenGL ES is a subset tailored for embedded systems³⁹.

Although shaders represent only a small part of the entire OpenGL software/hardware stack, implementing a compiler for them is still a significant undertaking. This is because a large number of APIs need to be supported. For instance,

²¹ <https://open.gl/framebuffers>

²² https://en.wikipedia.org/wiki/Stencil_buffer

³⁸ <https://www.khronos.org/registry/OpenGL-Refpages/>

³⁹ https://en.wikipedia.org/wiki/OpenGL_ES

there are over 80 texture-related APIs alone⁴⁰.

A practical approach to implementing such a compiler involves generating LLVM extended intrinsic functions from the shader frontend (parser and AST generator). These intrinsics can then be lowered into GPU-specific instructions in the LLVM backend. The overall workflow is illustrated as follows:

Fragment shader

```
#version 320 es
uniform sampler2D x;
out vec4 FragColor;

void main()
{
    FragColor = texture(x, uv_2d, bias);
}
```

llvm-ir

```
...
!1 = !{"sampler_2d"}
!2 = !{i32 SAMPLER_2D} ; SAMPLER_2D is integer value for sampler2D, for example:
↳ 0x0f02
; A named metadata.
!x_meta = !{!1, !2}

define void @main() #0 {
    ...
    %1 = @llvm.gpu0.texture(metadata !x_meta, %1, %2, %3); ; %1: %sampler_2d, %2: %uv_
↳ 2d, %3: %bias
    ...
}
```

asm of gpu

```
...
// gpu machine code
load $1, tex_a;
sample2d_inst $1, $2, $3 // $1: tex_a, $2: %uv_2d, $3: %bias

.tex_a // The driver sets the GPU descriptor-register index here
```

As shown at the end of the code above, the `.tex_a` memory address contains the Texture Object, which is bound by the driver during online compilation and linking. By binding a Texture Object (software representation) to a Texture Unit (hardware resource) via OpenGL API calls, the GPU can access and utilize Texture Unit hardware efficiently. This binding mechanism ensures that texture sampling and mapping are executed with minimal overhead during rendering.

For more information about LLVM extended intrinsic functions, please refer to⁴¹.

```
gvec4 texture(gsampler2D sampler, vec2 P, [float bias]);
```

⁴⁰ All the api listed in section 8.9 of https://www.khronos.org/registry/OpenGL/specs/es/3.2/GLSL_ES_Specification_3.20.html#texture-functions

⁴¹ <http://jonathan2251.github.io/lbd/funccall.html#add-specific-backend-intrinsic-function>

GPUs provide *Texture Units* to accelerate texture access in fragment shaders. However, *Texture Units* are expensive hardware resources, and only a limited number are available on a GPU. To manage this limitation, the OpenGL driver can associate a *Texture Unit* with a *sampler* variable using OpenGL API calls. This association can be updated or switched between shaders as needed. The following statements demonstrate how to bind and switch *Texture Units* across shaders:




Fig. 5.23: Relationships between the texturing concept⁴².

As shown in Fig. 5.23, the texture object is not bound directly to a shader (where sampling operations occur). Instead, it is bound to a *texture unit*, and the index of this texture unit is passed to the shader. This means the shader accesses the texture object through the assigned texture unit. Most GPUs support multiple texture units, though the exact number depends on the hardware capabilities⁴².

A *texture unit*—also known as a *Texture Mapping Unit (TMU)* or *Texture Processing Unit (TPU)*—is a dedicated hardware component in the GPU that performs texture sampling operations.

The *sampler* argument in the texture sampling function refers to a *sampler2D* (or similar) uniform variable. This variable represents the texture unit index used to access the associated texture object⁴².

Sampler Uniform Variables:

OpenGL provides a set of special uniform variables for texture sampling, named according to the texture target: *sampler1D*, *sampler2D*, *sampler3D*, *samplerCube*, etc.


You can create as many *sampler uniform variables* as needed and assign each one to a specific texture unit index using OpenGL API calls. Whenever a sampling function is invoked with a sampler uniform, the GPU uses the texture unit (and its bound texture object) associated with that sampler⁴².

For Java programmers, JOGL provides same level of API in Java for wrapping to OpenGL C API. As shown in Fig. 5.24, the JOGL *gl.bindTexture()* binds a *Texture Object* to a specific *Texture Unit*. Then, using *gl.getUniformLocation()* and *gl.uniform1i()*, you **associate** the **Texture Unit** with a *sampler uniform variable* in the shader.

For example, *gl.uniform1i(xLoc, 1)* assigns **Texture Unit 1** to the sampler variable at location *xLoc*. Similarly, passing 2 would refer to *Texture Unit 2*, and so on⁴³.

⁴² <http://ogldev.atspace.co.uk/www/tutorial16/tutorial16.html>

⁴³ <http://math.hws.edu/graphicsbook/c6/s4.html>

Fig. 5.24: Binding sampler variables^{Page 109, 43}.

The following Fig. 5.25 illustrates how the OpenGL driver reads metadata from a compiled GLSL object, how the OpenGL API **links sampler uniform variables** to **Texture Units**, and how the GPU **executes** the corresponding **texture instructions**.

Explaining the detailed steps for the figure above:

1. To enable the GPU driver to bind the *texture unit*, the frontend compiler must pass metadata for each *sampler uniform variable* (e.g., *sampler_2d_var* in this example)⁴⁴ to the backend. The backend then allocates and embeds this metadata in the compiled binary file⁴⁵.
2. During the link stage of on-line compilation of the GLSL shader, the GPU driver reads this metadata from the compiled binary file. It constructs an internal table mapping each *sampler uniform variable* to its attributes, such as {*name, type, location*}. This mapping allows the driver to properly populate the *Texture Descriptor* in the GPU's memory, linking the variable to a specific *texture unit*.
3. API:

```

xLoc = gl.getUniformLocation(prog, "x"); // prog: GLSL program, xLoc: location of
↪ sampler variable "x"

```

This API call queries the location of the *sampler uniform variable* named "x" from the internal table that the driver created after parsing the shader metadata.

The returned *xLoc* value corresponds to the location field associated with "x", which will later be used to bind a specific *texture unit* to this sampler variable via *gl.uniform1i(xLoc, unit_index)*.

SAMPLER_2D is the internal representation (usually an integer) that identifies a *sampler2D* type in the shader.

⁴⁴ The type of 'sampler uniform variable' called "sampler variables". <http://math.hws.edu/graphicsbook/c6/s4.html>

⁴⁵ This can be done by llvm metadata. <http://llvm.org/docs/LangRef.html#namedmetadatastructure> <http://llvm.org/docs/LangRef.html#metadata>




Fig. 5.25: Binding Sampler Variables to Texture Instructions

4. API:

```
gl.uniform1i(xLoc, 1);
```

This API call binds the sampler uniform variable x (located at $xLoc$) to **Texture Unit 1**. It works by writing the integer value 1 to the internal GLSL program memory at the location of the sampler variable x , as indicated by $xLoc$.

```
{xLoc, 1} : 1 is 'Texture Unit 1', xLoc is the memory address of 'sampler uniform_
↪variable' x
```

After this call, the OpenGL driver updates the **Texture Descriptor** table in GPU memory with this $\{xLoc, 1\}$ information.

Next, the driver associates the memory address or index of the GPU's texture descriptor with a hardware register or pointer used during fragment shader execution. For example, as shown in the diagram, the driver may write a pointer k to the $.tex_a$ field in memory.

This $.tex_a$ address is used by the GPU to locate the correct **Texture Unit** and access the texture object during shader execution.

5.

```
// gpu machine code
load $1, tex_a;
sample2d_inst $1, $2, $3 // $1: tex_a, $2: %uv_2d, $3: %bias

.tex_a // The driver sets the GPU descriptor-register index here at step 4
```

When the GPU executes texture instructions from a GLSL binary, the corresponding texture unit is selected through a texture descriptor in GPU memory. In this example, $.tex_a$ refers to $\{xLoc, 1\}$. If the architecture provides dedicated texture descriptors, the driver writes the descriptor into the descriptor table⁴⁶.

For example, an NVIDIA texture instruction is as follows:

```
// the content of tex_a bound to texture unit as step 5 above
tex.3d.v4.s32.s32 {r1,r2,r3,r4}, [tex_a, {f1,f2,f3,f4}];

.tex_a
```

The content of tex_a bound to texture unit set by driver as the end of step 4. The pixel of coordinates (x,y,z) is given by $(f1,f2,f3)$ user input. The $f4$ is skipped for 3D texture.

Above $tex.3d$ texture instruction load the calculated color of pixel (x,y,z) from texture image into GPRs $(r1,r2,r3,r4)=(R,G,B,A)$. And fragment shader can re-calculate the color of this pixel with the color of this pixel at texture image⁴⁷.

If it is 1d texture instruction, the $tex.1d$ as follows,

5. GPU Execution of Texture Instruction

```
// GPU machine code
load $1, tex_a;
sample2d_inst $1, $2, $3 // $1: tex_a, $2: %uv_2d, $3: %bias

.tex_a // Set by driver to index of GPU descriptor at step 4
```

⁴⁶ When performing a texture fetch, the addresses to read pixel data from are computed by reading the GPRs that hold the texture descriptor and the GPRs that hold the texture coordinates. It's mostly just general purpose memory fetching. <https://www.gamedev.net/forums/topic/681503-texture-units/>

⁴⁷ page 84: tex instruction, p24: texture memory https://www.nvidia.com/content/CUDA-ptx_isa_1.4.pdf

When the GPU executes the texture sampling instruction (e.g., *sample2d_inst*), it uses the *.tex_a* address, which was assigned by the driver in step 4, to access the appropriate **Texture Descriptor** from GPU memory. This descriptor corresponds to **Texture Unit 1** because of the earlier API call:

```
gl.uniform1i(xLoc, 1);
```

If the GPU hardware provides dedicated **texture descriptor registers** or memory structures, the driver maps *.tex_a* to those structures^{Page 112, 46}.

Example (NVIDIA PTX texture instruction):

```
// The content of tex_a is bound to a texture unit, as in step 4
tex.3d.v4.s32.s32 {r1,r2,r3,r4}, [tex_a, {f1,f2,f3,f4}];

.tex_a
```

Here, the *.tex_a* register holds the texture binding information set by the driver. The vector *{f1, f2, f3}* represents the 3D coordinates (x, y, z) provided by the shader or program logic. The *f4* value is ignored for 3D textures.

This *tex.3d* instruction performs a texture fetch from the bound 3D texture and loads the resulting color values into general-purpose registers:

- *r1*: Red
- *r2*: Green
- *r3*: Blue
- *r4*: Alpha

The **fragment shader** can then use or modify this color value based on further calculations or blending logic^{Page 112, 47}.

If a 1D texture is used instead, the texture instruction would look like:

```
// For compatibility with prior versions of PTX, the square brackets are not
// required and .v4 coordinate vectors are allowed for any geometry, with
// the extra elements being ignored.
tex.1d.v4.s32.f32 {r1,r2,r3,r4}, [tex_a, {f1}];
```

Since the ‘Texture Unit’ is a limited hardware accelerator on the GPU, OpenGL provides APIs that allow user programs to bind ‘Texture Units’ to ‘Sampler Variables’. As a result, user programs can balance the use of ‘Texture Units’ efficiently through OpenGL APIs without recompiling GLSL. Fast texture sampling is one of the key requirements for good GPU performance^{Page 109, 43}.

In addition to the API for binding textures, OpenGL provides the `glTexParameteri` API for texture wrapping⁴⁸. Furthermore, the texture instruction for some GPUs may include *S#* and *T#* values in the operands. Similar to associating ‘Sampler Variables’ to ‘Texture Units’, *S#* and *T#* are memory locations associated with texture wrapping descriptor registers. This allows user programs to change wrapping options without recompiling GLSL.

Even though the GLSL frontend compiler always expands function calls into inline functions, and LLVM intrinsic extensions provide an easy way to generate code through LLVM’s target description (TD) files, the GPU backend compiler is still somewhat more complex than the CPU backend.

(However, considering the effort required for the CPU frontend compiler such as Clang, or toolchains like the linker and GDB/LLDB, the overall difficulty of building a CPU compiler is not necessarily less than that of a GPU compiler.)

Here is the software stack of the 3D graphics system for OpenGL on Linux⁵. The Mesa open source project website is here⁴⁹.

⁴⁸ <https://learnopengl.com/Getting-started/Textures>

⁵ [<https://en.wikipedia.org/wiki/Mesa_\(computer_graphics\)>](https://en.wikipedia.org/wiki/Mesa_(computer_graphics))

⁴⁹ <https://www.mesa3d.org/>

GPU ARCHITECTURE

- *GPU Hardware Units*
- *SM (SIMT)*
 - *SM Hardware*
 - *SM Scheduling*
 - *SIMT and SPMD Pipelines*
- *Processor Units and Memory Hierarchy in NVIDIA GPU*
- *Memory Subsystem*
 - *Address Coalescing and Gather-scatter*
 - *VRAM dGPU*
 - *RegLess-style architectures*
- *Specialized Units*
 - *Geometry Units*
 - *Rasterization Units*
 - *Texture Mapping Units (TMUs)*
 - *Render Output Units (ROPs)*
 - *Framebuffer hierarchy*

6.1 GPU Hardware Units

A GPU (graphics processing unit) is built as a massively generic parallel processor of SIMD/SIMT architecture with several specialized processing units inside shown as Fig. 6.2 from the *section Graphics HW and SW Stack*.

From compiler's view, GPU is shown as Fig. 6.3.

A GPU is not just “many cores”—it's a mix of general-purpose compute clusters, specialized units, and the memory subsystem. It corresponds to the block diagram graph shown in Fig. 6.3.

The stages of the OpenGL rendering pipeline and the GPU hardware units that accelerate them as shown in Fig. 6.4:

Compute Cluster

¹ https://www3.ntu.edu.sg/home/ehchua/programming/opengl/CG_BasicsTheory.html

² https://en.wikipedia.org/wiki/Graphics_processing_unit




Fig. 6.1: Computer Graphics Hardware (figure from book ^{Page 115, 1})




Fig. 6.2: Components of a GPU: GPU has accelerated video decoding and encoding ^{Page 115, 2}




Fig. 6.3: Components of a GPU: SIMD/SIMT + several specialized processing units




Fig. 6.4: The stages of OpenGL pipeline and GPU's acceleration components

- **Role:** Provide large-scale data-parallel execution. Each GPU contains many Streaming Multiprocessors (SMs) or Compute Units (CUs), each capable of executing thousands of threads in parallel.
- **Components:**
 - **Warp Scheduler** —Schedules groups of threads (Warps/Wavefronts), issues instructions in SIMT (Single Instruction, Multiple Threads) fashion.
 - **Registers** —Per-thread private storage, the fastest memory level.
 - **Shared Memory / L1 Cache** —On-chip memory close to the SM. Shared Memory is explicitly managed by the programmer for cooperation across threads, while L1 acts as a transparent cache.
 - **ALUs (FP/INT)** —Execute floating-point and integer arithmetic. They form the bulk of compute resources inside an SM.
 - **SFUs (Special Function Units)** —Execute transcendental functions such as sin, cos, exp, and reciprocal approximations.
 - **Load/Store Units** —Handle global, local, and shared memory access, interact with coalescing and caching logic.
 - **RegLess Staging Operands (RSO)** —Temporary operand buffers used to hide instruction and memory latencies.
- **Usage:**
 - Run programmable shaders (vertex, fragment, geometry, compute).
 - Perform general-purpose compute workloads (GPGPU).
 - Issue texture fetch requests to TMUs.
 - Interact with memory hierarchy via load/store units.
 - Offload certain operations to Tensor or Ray-Tracing units.

Specialized Units

- **Role:** Accelerate fixed-function or specialized stages of the graphics and compute pipeline that are inefficient to run purely in SMs.
- **Components and Usage:**
 - **Geometry Units** —Assemble input vertices into primitives (points, lines, triangles). Perform tessellation (subdivide patches into smaller primitives), clipping (discard geometry outside view), and geometry shading.
Usage: Corresponds to the geometry/tessellation stage in the graphics pipeline.
 - **Rasterization Units** —Convert vector-based primitives into fragments (potential pixels). Interpolate per-vertex attributes (texture coordinates, normals, colors) across the surface of each primitive.
Usage: Bridge between geometry and fragment stages; produces fragments for SM fragment shading.
 - **Texture Mapping Units (TMUs)** —Fetch texture data from memory, apply filtering (bilinear, trilinear, anisotropic), and compute texel addresses (wrap, clamp).
Usage: Invoked during fragment shading inside SMs to provide sampled texture values.
 - **Render Output Units (ROPs)** —Handle late-stage pixel processing. Perform blending operations (alpha, additive), depth and stencil tests, and write final pixel values to the framebuffer in VRAM.
Usage: Final step of the graphics pipeline before display scanout.
 - **Tensor / Matrix Cores** —Perform fused-multiply-add (FMA) on large matrix tiles. Designed for machine learning, AI inference, and linear algebra.

Usage: Accelerate deep learning workloads or matrix-heavy compute kernels.

- **Ray-Tracing Units (RT Cores)** — Traverse bounding volume hierarchies (BVH) and perform ray—primitive intersection tests in hardware.

Usage: Enable real-time ray tracing³ by offloading intersection work from SMs.

- **Video Engines** — Dedicated ASICs for video codec operations such as H.264/H.265/AV1 encode and decode.

Usage: Media playback, streaming, and video encoding without occupying SMs.

- **Display Controller** — Reads final framebuffer images from VRAM and drives display interfaces like HDMI and DisplayPort.

Usage: Outputs rendered frames to monitors or VR headsets.

Memory Subsystem

- **Role:** Deliver high-bandwidth data access to thousands of threads while minimizing latency through caching and access optimization.

- **Components:**

- **L1 / Shared Memory** — Closest to SMs. Shared Memory is explicitly used by programs for intra-block communication, while L1 acts as an automatic cache.

Usage: Boosts performance by keeping frequently accessed data close to execution units.

- **L2 Cache** — Shared across all SMs. Reduces redundant traffic to VRAM and improves latency for reused data.

Usage: Provides intermediate caching layer for both compute and graphics.

- **VRAM (GDDR / HBM)** — External high-bandwidth DRAM. Stores textures, framebuffers, vertex/index buffers, and large compute datasets.

Usage: The main memory backing for all GPU workloads.

- **Interconnect / Memory Controller** — Orchestrates memory requests, manages access to VRAM, and ensures fairness between SMs.

Usage: Handles scheduling and distribution of memory transactions.

- **Memory Coalescing Unit** — Combines multiple per-thread memory requests from a Warp into fewer, wider transactions. Most effective for contiguous access patterns.

Usage: Improves memory bandwidth efficiency and reduces wasted cycles.

- **Gather—Scatter Unit** — Handles irregular or sparse memory accesses where coalescing is not possible. May break requests into multiple smaller transactions.

Usage: Supports workloads such as sparse matrix operations, graph traversal, or irregular data structures.

Data Flow Highlights

- **Graphics pipeline path:** Vertex data → Geometry Units → Rasterization Units → Fragment Shading (SMs) → TMUs (texture fetch) → ROPs (blend/depth/stencil) → VRAM (framebuffer).
- **Compute path:** SMs execute general-purpose kernels → optional offload to Tensor or RT cores → interact with caches → VRAM.
- **Memory behavior:** SMs issue memory requests → Coalescing Unit optimizes if possible → L2 cache → VRAM. For irregular access (e.g., sparse data), Gather—Scatter generates multiple VRAM transactions.
- **Display path:** Final framebuffer stored in VRAM → Display Controller → HDMI / DP scanout.

³ <[https://en.wikipedia.org/wiki/Ray_tracing_\(graphics\)](https://en.wikipedia.org/wiki/Ray_tracing_(graphics))>

All Together

GPU provides the following hardware to accelerate graphics rendering pipeline as follows:

🔄 Simplified Flow (OpenGL → Hardware)

1. Vertex Fetch → VRAM & Memory Controllers.
2. Vertex Shader → SM cores + Geometry Units.
3. Geometry/Tessellation → SM core + Geometry Units.
4. Rasterization → Rasterization units.
5. Fragment Shader → SM cores + TMUs (texture sampling).
6. Depth/Stencil/Blending → ROPs.
7. Framebuffer Write → L2 cache & VRAM → Display Controller.

Variable Rate Shading (VRS) Support

By utilizing certain GPU units as outlined below, Variable Rate Shading (VRS) can be supported⁴.

- Rasterizer (Rasterization Units):
 - Decides how many fragments per pixel (or group of pixels) will actually be shaded.
 - Instead of generating 1 fragment per pixel, it may shade 1 fragment for a 2×2 or 4×4 block and reuse that result.
- Fragment Shader Cores (SMs/CUs):
 - Still run the shading code, but at a reduced frequency (fewer fragment invocations).
- ROPs (and pipeline integration):
 - Apply results to the framebuffer, handling blending/depth as usual.

6.2 SM (SIMT)

Single instruction, multiple threads (SIMT) is an execution model used in parallel computing where a single central “Control Unit” broadcasts an instruction to multiple “Processing Units” for them to all optionally perform simultaneous synchronous and fully-independent parallel execution of that one instruction. **Each PU has its own independent data and address registers, its own independent Memory, but no PU in the array has a Program counter⁵.**

Summary:

- Each Control Unit has a program counter (PC) and controls tens of processing units (PUs).
- Each processing unit (PU) has its own general-purpose register set (GPR) and stack memory.
- A PU is a pipeline execution unit, analogous to a CPU execution unit.

6.2.1 SM Hardware

The leading NVIDIA GPU architecture is illustrated in Fig. 6.5, where the scoreboard is shown without the mask field. This represents a SIMT pipeline with a scoreboard.

⁴ <https://developer.nvidia.com/vrworks/graphics/variable rates shading>

⁵ https://en.wikipedia.org/wiki/Single_instruction,_multiple_threads

⁶ The SIMD Thread Scheduler includes a scoreboard that lets it know which threads of SIMD instructions are ready to run, and then it sends them off to a dispatch unit to be run on the multithreaded SIMD Processor. It is identical to a hardware thread scheduler in a traditional multithreaded processor (see Chapter 3), just that it is scheduling threads of SIMD instructions. Thus, GPU hardware has two levels of hardware schedulers: (1) the Thread Block Scheduler that assigns Thread Blocks (bodies of vectorized loops) to multi-threaded SIMD Processors, which ensures that Thread Blocks are assigned to the processors whose local memories have the corresponding data, and (2) the SIMD Thread Scheduler within a SIMD Processor, which





Figure 4.14 Simplified block diagram of a Multithreaded SIMD Processor. It has 16 SIMD lanes. The SIMD Thread Scheduler has, say, 48 independent threads of SIMD instructions that it schedules with a table of 48 PCs.

Fig. 6.5: Simplified block diagram of a Multithreaded SIMD Processor. (figure from bookPage 120, 6)

Note

A SIMD thread executed by a SIMD processor (also known as an SM) has 16 lanes.

schedules when threads of SIMD instructions should run. Book Figure 4.14 of Computer Architecture: A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design)



- Streaming Multiprocessor SM has two 16-way SIMD units and four special function units. Fermi has 32 SIMD Lanes and Cuda cores. SM has L1 and Read Only Cache (Uniform Cache) GTX480 has 48 SMs.
- In Fermi, ALUs run at twice the clock rate of rest of chip. So each decoded instruction runs on 32 pieces of data on the 16 ALUs over two ALU clocks¹⁰. However after Fermi, the ALUs run at the same clock rate of rest of chip.
- As Fig. 6.6 in Fermi and Volta, it can dual-issue “float + integer” or “integer + load/store” but cannot dual-issue “float + float” or “int + int”.
- Uniform cache: used for storing constant variables in OpenGL (see *uniform of Pipeline Qualifiers*) and in OpenCL/CUDA.

Configurable maximum resident warps and allocated registers per thread as follows:

- Example: Fermi SM (SM 2.x)
 - Hardware limit:
 - * Total registers per SM = $32,768 \times 32\text{-bit}$
 - * Max Warps per SM = 48
 - * Max threads per SM = 1536
 - * Max registers/thread = 63

⁷ Book Figure 4.20 of Computer Architecture: A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design)

⁸ <https://www.tomshardware.com/reviews/geforce-gtx-480,2585-18.html>

⁹ https://www.nvidia.com/content/PDF/fermi_white_papers/NVIDIA_Fermi_Compute_architecture_Whitepaper.pdf?utm_source=chatgpt.com

¹⁰ https://www.cs.cmu.edu/afs/cs/academic/class/15418-s12/www/lectures/02_multicore.pdf

- Configuration: If each thread uses R registers:
 - * Max resident threads = $\text{floor}(32768 / R)$
 - * Max resident Warps = $\text{floor}(\text{Max resident threads} / 32)$
 - * E.g. **R=32: Max resident threads = 32768/32 = 1024, Max resident Warps = 1024/32 = 32.**
- After Fermi, the hardware limit for Maxwell, Pascal, Volta and Ampere are:
 - Hardware limit:
 - * Each SM includes 32 Cuda cores and Lanes $\rightarrow$ 32 active threads.
 - * Total registers per SM = 64K x 32-bit
 - * Max Warps per SM = 64
 - * Max threads per SM = 2048 (64 Warps x 32 threads)
 - * Max registers/thread = 255
 - Notes:
 - * Registers per thread: max number of registers **compiler can allocate to a thread**.
 - * The “registers per thread” limit (255) is a hardware/compiler limit, but the actual number used depends on the kernel. If a kernel uses too many registers per thread, occupancy drops (fewer threads can be resident).
 - * The “max threads per SM = 2048” is a theoretical upper limit; actual resident threads will also depend on shared memory usage, number of thread-blocks per SM, and register usage.

Note

A SIMD thread executed by a Multithreaded SIMD processor, also known as an SM, processes 32 elements.

As configuration above, the 32,768 registers per SM can be configured to each thread allocated 32 registers, Max resident Warps = 32.

Fermi has a mode bit that offers the choice of using 64 KB of SRAM as a 16 KB L1 cache with 48 KB of Local Memory or as a 48 KB L1 cache with 16 KB of Local Memory¹¹.




Fig. 6.8: SM select Thread Blocks to run¹².



Figure 3: Automatic Scalability

Fig. 6.9: Mapping Thread Block to SMs¹³.

6.2.2 SM Scheduling

- A GPU is built around an array of Streaming Multiprocessors (SMs). A multithreaded program is partitioned into blocks of threads that execute independently from each other, so that a GPU with more multiprocessors will automatically execute the program in less time than a GPU with fewer multiprocessors¹³.

- Nvidia's GPUs:

Fermi (2010), Kepler (2012), Maxwell (2014), Pascal (2016), Volta (2017), Turing (2018), Ampere (2020), Ada Lovelace (2022), and Hopper (2022, for data centers).

- Two levels of scheduling:

- Level 1: Thread Block Scheduler

For Fermi/Kepler/Maxwell/Pascal (pre-Volta): Warp-synchronous SIMT (lock-step in Warp):

A Warp includes 32 threads in Fermi GPU. Each Streaming Multiprocessor SM includes 32 Lanes in Fermi GPU, as shown in Fig. 6.8, the Thread Block includes a Warp (32 threads). According Fig. 6.9, more than one block can be assigned and run on a same SM.

When an SM executes a Thread Block, the threads within the block execute together. If any thread in a Warp is not ready because of operand data dependencies, the scheduler switches context between Warps. During a context switch, all the data of the current Warp remains in the register file so it can resume quickly once its operands are ready¹².

¹¹ Page 306 of Computer Architecture: A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design)

¹² <[https://en.wikipedia.org/wiki/Thread_block_\(CUDA_programming\)](https://en.wikipedia.org/wiki/Thread_block_(CUDA_programming))>


¹³ <<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#warps>>

Once a Thread Block is launched on a multiprocessor (SM), all of its Warps are **resident** until their execution finishes. Thus a new block is not launched on an SM until there is sufficient number of free registers for all Warps of the new block, and until there is enough free shared memory for the new block^{Page 124, 12}.

– Level 2: Warp Scheduler

Manages CUDA threads (resident threads) within the same Warp.

A **resident thread** is a thread whose execution context has been allocated on an SM (registers, Warp slot, shared memory). Once resident, the thread is always in exactly one of the following execution states.



Threads retain their registers and per-thread local memory during the stalled state. Therefore, **the context switch incurs almost no overhead compared to CPU threads.**

- * No pipeline flush: illustrate below.
- * No register save/restore
- * No stack frame swapping
- * **No OS involvement**
- * Takes roughly **1 cycle**

No pipeline flush because:

For Fermi/Kepler/Maxwell/Pascal (pre-Volta): Warp-synchronous SIMT (lock-step in Warp):

- * No data is saved/restored when switching to another Warp
- * Switching Warps = selecting a different Warp in the Warp scheduler
- * No pipeline flush

On an NVIDIA GPU, no pipeline flush occurs when a Warp stalls because the Warp's next instruction is **never issued until its operands are ready** as illustrated in [Warp scheduling in Level 1](#). The stalled Warp simply stops issuing instructions, and its pipeline slot is taken by another ready Warp. When the stall condition clears, the Warp re-enters the pipeline by issuing the stalled instruction anew. No state is saved or restored.

For Volta, Turing, Ampere, Hopper: Independent Thread Scheduling:

- * No pipeline flush
 - Stalled threads simply do not issue instructions.
 - Other threads in the same warp continue issuing independently.
 - No pipeline flush needed and No data is saved/restored because instructions are tracked per thread, not per warp.
- Thread Active/Inactive

```
GLSL example for branch divergence
-----

// The value of x is different between threads
if (x > 0.0)
    color = red;
else
    color = blue;

GPUs use conditional instructions like CPUs.

When a shader executes a conditional branch and threads evaluate the
condition differently, the GPU splits execution using a mask register.

predicate = cond          // predicate is the mask register
@predicate instruction

is a form of conditional (predicated) instruction execution on GPUs.

In NVIDIA PTX, it is activemask register.

if EXEC_MASK[thread] == 1
    execute
else
    skip
```

6.2.3 SIMT and SPMD Pipelines

This section illustrates the difference between SIMT and SPMD pipelines using the same pipeline stages: Fetch (F), Decode (D), Execute (E), Memory (M), and Writeback (W).

A GPU contains many SMs. The execution model between SMs is MIMD (Multiple Instructions, Multiple Data) when running different programs, or SPMD (Single Program, Multiple Data). However, within a single SM, the execution model is SIMD/SIMT.”

Low-end GPUs implement SIMD in their pipelines, where all instructions are executed in lockstep. High-end GPUs, however, approximate SPMD in their pipelines, meaning that instructions are interleaved within the pipeline, as shown below.

SPMD Programming Model vs SIMD/SIMT Execution

In the SISD of CPU, a thread is a single pipeline execution unit which can be issued at any specific address.

In a multi-core CPU running SPMD, each core can schedule and execute instructions at any program counter (PC). For example, core-1 may execute I(1–10), while core-2 executes I(31–35). **For GPU, however, within an SM, it is not possible to schedule thread-1 to execute I(1–10) while thread-2 executes I(31–35).**

As a result, **there is no mainstream GPU that is truly hardware-SPMD** (where each thread has its own independent pipeline). All modern GPUs (NVIDIA, AMD, Intel) implement SPMD as a programming model, but under the hood they execute in SIMD lock-step groups (Warps or Wavefronts). GPUs expose an **SPMD programming model** (each thread runs the same kernel on different data). However, the hardware actually executes instructions in **SIMD/SIMT lock-step groups**.

An example to illustrate the difference between Pascal SIMT, Volta SIMT and SPMD.

Divergent Kernel Example:

```
-----
if (tid % 2 == 0) {           // even threads: long loop
    for (...) { loop_body } // many iterations
} else {                     // odd threads: short path
    C[tid] = A[tid] + B[tid];
}
```

Legend: F=Fetch, D=Decode, E=Execute, M=Memory, W=Writeback
S=Stall/masked-off, "..." = loop continues

=====

Pascal (lock-step SIMT with SIMT stack)

```
-----
Cycle →   0   1   2   3   4   5   6   7   8   9  10  11  12 ...
T0 even:  F   D   E   M   W   F   D   E   M   W   F   D   ...
T1 odd :   S   S   S   S   S   S   S   S   S   S   S   S   ...
          (Odd threads wait until even path completes, then:)
          ... F D E M W → done
```

=====

Volta (SIMT with independent thread scheduling)

```
-----
Cycle →   0   1   2   3   4   5   6   7   8   9  10  11 ...
T0 even:  F   D   E   M   W   F   D   E   M   W   F   D   ...
T1 odd :           F   D   E   M   W   done
          (Odd thread issues its short path early,
           interleaved with even loop instructions)
```

=====

True SPMD (CPU-like, fully independent threads)

```
-----
Cycle →   0   1   2   3   4   5   6   7   8   9 ...
T0 even:  F   D   E   M   W   F   D   E   M   W   ...
T1 odd :  F   D   E   M   W   done
          (Threads fetch/execute independently —
           odd thread finishes immediately)
```

Note**SPMD and MIMD**

When run a single program across all cores, SPMD and MIMD pipelines look the same.

The subsection Mapping data in GPU includes more details in Lanes masked.

Scoreboard purpose:

- GPU scoreboard = in-order issue, out-of-order completion
- CPU reorder buffer (ROB) = out-of-order issue + completion, but retire in-order
 - CPUs use a ROB to support out-of-order issue and retirement.

Comparison of Volta and Pascal

In a lock-step GPU without divergence support, the scoreboard entries include only { Warp-ID, PC (Instruction Address), ... }. With divergence support (as in real-world GPUs), the scoreboard entries expand to { Warp-ID, PC, mask, ... }.

Volta (Cuda thread/SIMD Lane with PC, Program Counter and Call Stack)**GPU scoreboard = in-order issue, out-of-order completion**

- SIMT GPU before Volta = scoreboard contains: { Warp ID + PC + Active Mask }
- Volta = scoreboard contains: { Warp ID + PC per thread (+ readiness per thread) }

Example for mutex¹⁴

```
//
__device__ void insert_after(Node *a, Node *b)
{
    Node *c;
    lock(a); lock(a->next);
    ...
    unlock(c); unlock(a);
}
```

Assume that the mutex is contended across SMs but not within the same SM. On average, each thread spends 10 cycles executing the insert_after operation without resource contention, and 20 cycles when accounting for contention. Therefore, the average total execution time for 32 threads in an SM is:

- Volta: 20 cycles
- Pascal: 640 cycles (20 cycles × 32 threads, due to lack of independent progress inside a Warp)

6.3 Processor Units and Memory Hierarchy in NVIDIA GPU¹⁵

Illustrate L1, L2 and Global Memory used by SM and whole chip of GPU as Fig. 6.12.

The Fig. 6.11 illustrates the memory hierarchy in NVIDIA GPU. The Cache flow for 3D Model Information, Animation Parameters, and GLSL Variables is as follows:

3D Model Information:

¹⁴ See the same Figures from <https://images.nvidia.com/content/volta-architecture/pdf/volta-architecture-whitepaper.pdf>

¹⁵ chatgpt: Give me a memory hierarchy for L1, L2, local memory, shared memory for these processing units of hierarchy in reST and separate DOT graph.

¹⁶ Book Figure 4.17 of Computer Architecture: A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design)




Figure 4.18 GPU Memory structures. GPU Memory is shared by all **Grids** (vectorized loops), Local Memory is shared by all threads of SIMD instructions within a thread block (body of a vectorized loop), and Private Memory is private to a single CUDA Thread.

Fig. 6.10: GPU memory (figure from book^{Page 128, 16}).




Fig. 6.11: Processor Units and Memory Hierarchy in NVIDIA GPU **Local Memory is shared by all threads and Cached in L1 and L2.** In addition, the **Shared Memory is provided to use per-SM, not cacheable.**




Figure 1: *The GPU Devotes More Transistors to Data Processing*

Fig. 6.12: **L1 Cache: Per-SM, Coherent across all 16 Lanes in the same SM.** L2 Cache: Coherent across all SMs and GPCs. Global Memory (DRAM: HBM/GDDR). Both HBM and GDDR are DRAM. GDDR (Graphics DDR) — optimized for GPUs (GDDR5, GDDR6, GDDR6X). HBM (High Bandwidth Memory) — 3D-stacked DRAM connected via TSVs (Through-Silicon Vias) for extremely high bandwidth and wide buses ^{Page 124, 13}.

- VBO/IBO → Global → L2 → L1 → Registers
- Material constants → Uniform Cache → Registers

Animation Parameters:

- Bone matrices → Uniform Cache → Registers
- Morph targets → Global → L2 → L1 → Registers
- Shared bone data (compute) → Shared Memory → Registers

GLSL Variables:

- uniform (include uniform variable and UBO) → Uniform Cache → Registers
- in (vertex attributes) → Global → L2 → L1 → Registers
- out (varyings) → Registers → Interpolators
- buffer (SSBO) → Global → L2 → L1 → Registers
- shared → Shared Memory → Registers
- local arrays → Registers or Local Memory

More details of the NVIDIA GPU memory hierarchy are described as follows:

- **Registers**
 - Per-thread, fastest memory, located in CUDA cores, as illustrated also in Fig. 6.6.
 - *Configurable maximum resident warps and allocated registers per thread* following Fig. 6.6.
 - Latency: ~1 cycle.
- **Uniform / Constant cache**
 - Stored constant variables in OpenGL and OpenCL/CUDA, as illustrated in Fig. 6.6.
- **Local Memory**
 - Per-thread, stored in global DRAM.
 - Cached in L1 and L2.
 - Latency: high, depends on cache hit/miss.
- **Shared Memory**
 - Per-SM, shared across threads in a Thread Block as shown in Fig. 6.6.
 - On-chip, programmer-controlled.
 - Latency: ~20 cycles.
- **L1 Cache**
 - Per-SM, unified with shared memory.
 - Hardware-managed.
 - Latency: ~20 cycles.
- **L2 Cache**
 - Shared across the entire GPU chip.
 - Coherent across all SMs and GPCs as shown in Fig. 6.12.
- **Global Memory (DRAM: HBM/GDDR)**

- Visible to all SMs across all GPCs.
- Highest latency (~400–800 cycles).

GPU Hierarchy Context

- **GigaThread Engine (chip-wide scheduler)**
 - Contains multiple GPCs.
 - * Fermi (2010): up to 4 GPCs per chip.
 - * Pascal GP100 (Tesla P100): 6 GPCs.
 - * Volta GV100 (Tesla V100): 6 GPCs.
 - Distributes Thread Blocks to all GPCs.
- **GPC (Graphics Processing Cluster)**
 - Contains multiple TPCs.
- **TPC (Texture Processing Cluster)**
 - Groups 1–2 SMs.
- **SM (Streaming Multiprocessor)**
 - Contains CUDA cores, registers, shared memory, L1 cache.
- **CUDA Cores**
 - Execute threads with registers and access the memory hierarchy.
- A Warp of 32 threads is mapped across 16 Lanes. If each Lane has 2 Chimes, it may support dual-issue or time-sliced execution as Fig. 6.14.

In the following matrix multiplication code, the 8096 elements of matrix $A = B \times C$ are mapped to Thread Blocks, SIMD Threads, Lanes, and Chimes as illustrated in the Fig. 6.15. In this example, it run on **time-sliced execution**.

Listing 6.1: MATMUL CUDA Example

```
// Invoke MATMUL with 256 threads per Thread Block
__host__
int nblocks = (n + 255) / 512;
matmul<<<nblocks, 255>>>(n, A, B, C);
// MATMUL in CUDA
__device__
void matmul(int n, double A, double *B, double *C) {
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < n) A[i] = B[i] + C[i];
}
```

Explain the mapping and execution in Fig. 6.15 for *MATMUL CUDA Example* using the terminology from Fig. 6.13 and the previous sections of this book, presented in the table below.

¹⁷ Book Figure 4.12 of Computer Architecture: A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design)

¹⁸ Book Figure 4.13 of Computer Architecture: A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design)

Type	More descriptive name	Closest old term outside of GPUs	Official CUDA/NVIDIA GPU term	Book definition
Program abstractions	Vectorizable Loop	Vectorizable Loop	Grid	A vectorizable loop, executed on the GPU, made up of one or more Thread Blocks (bodies of vectorized loop) that can execute in parallel.
	Body of Vectorized Loop	Body of a (Strip-Mined) Vectorized Loop	Thread Block	A vectorized loop executed on a multithreaded SIMD Processor, made up of one or more threads of SIMD instructions. They can communicate via Local Memory.
	Sequence of SIMD Lane Operations	One iteration of a Scalar Loop	CUDA Thread	A vertical cut of a thread of SIMD instructions corresponding to one element executed by one SIMD Lane. Result is stored depending on mask and predicate register.
Machine object	A Thread of SIMD Instructions	Thread of Vector Instructions	Warp	A traditional thread, but it contains just SIMD instructions that are executed on a multithreaded SIMD Processor. Results stored depending on a per-element mask.
	SIMD Instruction	Vector Instruction	PTX Instruction	A single SIMD instruction executed across SIMD Lanes.
Processing hardware	Multithreaded SIMD Processor	(Multithreaded) Vector Processor	Streaming Multiprocessor	A multithreaded SIMD Processor executes threads of SIMD instructions, independent of other SIMD Processors.
	Thread Block Scheduler	Scalar Processor	Giga Thread Engine	Assigns multiple Thread Blocks (bodies of vectorized loop) to multithreaded SIMD Processors.
	SIMD Thread Scheduler	Thread scheduler in a Multithreaded CPU	Warp Scheduler	Hardware unit that schedules and issues threads of SIMD instructions when they are ready to execute; includes a scoreboard to track SIMD Thread execution.
	SIMD Lane	Vector Lane	Thread Processor	A SIMD Lane executes the operations in a thread of SIMD instructions on a single element. Results stored depending on mask.
Memory hardware	GPU Memory	Main Memory	Global Memory	DRAM memory accessible by all multithreaded SIMD Processors in a GPU.
	Private Memory	Stack or Thread Local Storage (OS)	Local Memory	Portion of DRAM memory private to each SIMD Lane.
	Local Memory	Local Memory	Shared Memory	Fast local SRAM for one multithreaded SIMD Processor, unavailable to other SIMD Processors.
	SIMD Lane Registers	Vector Lane Registers	Thread Processor Registers	Registers in a single SIMD Lane allocated across a full thread block (body of vectorized loop).

Figure 4.12 Quick guide to GPU terms used in this chapter. We use the first column for hardware terms. Four groups cluster these 11 terms. From top to bottom: Program Abstractions, Machine Objects, Processing Hardware, and Memory Hardware. Figure 4.21 on page 309 associates vector terms with the closest terms here, and Figure 4.24 on page 313 and Figure 4.25 on page 314 reveal the official CUDA/NVIDIA and AMD terms and definitions along with the terms used by OpenCL.

Fig. 6.13: Terms in Nvidia's gpu (figure from book^{Page 132, 17})




Fig. 6.14: In **dual-issue mode**, Chime A carries floating-point data while Chime B carries integer data—both issued by the same CUDA thread. In contrast, under **time-sliced execution**, Chime A and Chime B carry either floating-point or integer data independently, and are assigned to separate CUDA threads.




Figure 4.13 The mapping of a Grid (vectorizable loop), Thread Blocks (SIMD basic blocks), and threads of SIMD instructions to a vector-vector multiply, with each vector being 8192 elements long. Each thread of SIMD instructions calculates 32 elements per instruction, and in this example each Thread Block contains 16 threads of SIMD instructions and the Grid contains 16 Thread Blocks. The hardware Thread Block Scheduler assigns Thread Blocks to multithreaded SIMD Processors and the hardware Thread Scheduler picks which thread of SIMD instructions to run each clock cycle within a SIMD Processor. Only SIMD Threads in the same Thread Block can communicate via Local Memory. (The maximum number of SIMD Threads that can execute simultaneously per Thread Block is 16 for Tesla-generation GPUs and 32 for the later Fermi-generation GPUs.)

Fig. 6.15: Mapping 8192 elements of matrix multiplication for Nvidia's GPU (figure from ^{Page 132, 18}). SIMT: 16 SIMD threads in one Thread Block.

Table 6.1: Summary terms for GPU.

Terms	Structure	Description
Grid, Giga Thread Engine	Each loop (Grid) consists of multiple Thread Blocks.	Grid is Vectorizable Loop as Fig. 6.13. The hardware scheduler Giga Thread Engine schedules the Thread Blocks to SMs.
Thread Block	In this example, each Grid has 16 Giga Thread ¹⁹ .	Each Thread Block is assigned 512 elements of the vectors to work on. As Fig. 6.15, it assigns 16 Thread Block to 16 SMs. Giga Thread is the name of the scheduler that distributes Thread Blocks to Multiprocessors, each of which has its own SIMD Thread Scheduler ¹⁹ . More than one Block can be mapped to a same SM as the explanation in “Level 1: Thread Block Scheduler” for Fig. 6.9.
Streaming Multiprocessor, SM, GPU Core (Warp) ²⁰	Each SIMD Processor has 16 SIMD Threads.	Each SIMD processor includes local memory, as in Fig. 6.10. Local memory is shared among SIMD Lanes within a SIMD processor but not across different SIMD processors. A Warp has its own PC and may correspond to a whole function or part of a function. Compiler and runtime may assign functions to the same or different Warps ²¹ .
Cuda core	Fermi has 32 Cuda cores in a SM as Fig. 6.7.	A CUDA core is the scalar execution unit inside an SM. It is capable of executing one integer or floating-point instruction from one Lane of a Warp. The CUDA core is analogous to an ALU pipeline stage in a CPU.
Cuda Thread	Each SM can configure to have different number of resident threads.	Fermi can configure Max resident threads = $32768/32 = 1024$ for 32 registers/per thread in a SM as mentioned earlier. A CUDA thread is the basic unit of execution defined in CUDA’s programming model. Each thread executes the kernel code independently with its own registers, program counter (PC), and per-thread local memory. Each Thread has its TLR (Thread Level Registers) allocated from Register file (32768×32 -bit) by SIMD Processor (SM) as Fig. 6.6.
SIMD Lane	Each SIMD Thread has 32 Lanes.	A vertical cut of a thread of SIMD instructions corresponding to one element executed by one SIMD Lane. It is a vector instruction with processing 32-elements. A Warp of 32 threads is mapped across 32 Lanes. Lane = per-thread execution slot inside a Warp. If each Lane has 2 Chimes, it may support dual-issue or time-sliced execution as Fig. 6.14.
Chime	Each Lane has 2 Chimes.	A Chime represents one “attempt” or opportunity for issuing instructions from Warps. In Fermi (SM2.x): Each SM has 2 Warp schedulers. Each Warp scheduler has 2 dispatch units (dual-issue, but with constraints, it can issue “float + load/store” for “fadd and load C[i]” in this example).

References

- [NVIDIA GPU Architecture Overview](#)
- [Understanding Warps and Threads](#)

6.4 Memory Subsystem

6.4.1 Address Coalescing and Gather-scatter

Brief description is shown in Fig. 6.16.

¹⁹ Figure 4.15 Floor plan of the Fermi GTX 480 GPU of A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design). **Giga Thread** is the name of the scheduler that distributes Thread Blocks to Multiprocessors, each of which has its own SIMD Thread Scheduler.

²⁰ Copilot: Is GPU core meaning SM in NVidia?

²¹ Book Figure 4.14 and 4.24 of Computer Architecture: A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design)




Fig. 6.16: Coalescing and Gather-scatter

The Load/Store Units (LD/ST) is important because memory latency is huge compared to ALU ops. Some GPUs provide Address Coalescing and gather-scatter to accelerate memory access.

- Address Coalescing: **Memory coalescing is the process of merging memory requests from threads in a Warp (NVIDIA: 32 threads, AMD: 64 threads) into as few memory transactions as possible.**
 - Cache miss (global memory/DRAM): Coalescing = big performance improvement.
 - Cache hit (L1/L2): Coalescing = smaller benefit, since cache line fetch already amortizes cost.
 - Note that unlike vector architectures, GPUs don't have separate instructions for sequential data transfers, strided data transfers, and gather-scatter data transfers. All data transfers are gather-scatter! To regain the efficiency of sequential (unit-stride) data transfers, GPUs include special Address Coalescing hardware to recognize when the SIMD Lanes within a thread of SIMD instructions are collectively issuing sequential addresses. That runtime hardware then notifies the Memory Interface Unit to request a block transfer of 32 sequential words. To get this important performance improvement, the GPU programmer must ensure that adjacent CUDA Threads access nearby addresses at the same time that can be coalesced into one or a few memory or cache blocks, which our example does²².
- Gather-scatter data transfer: **HW support sparse vector access is called gather-scatter.** The VMIPS instructions are LVI (load vector indexed or gather) and SVI (store vector indexed or scatter)²³.

1. Address Coalescing in GPU Memory Transactions

Definition: Memory coalescing is the process of merging memory requests from threads in a Warp (NVIDIA: 32 threads, AMD: 64 threads) into as few memory transactions as possible.

How It Works:

- If threads access **contiguous and aligned addresses**, the hardware combines them into a single memory transaction.
- If threads access **strided or random addresses**, the GPU must issue multiple transactions, wasting bandwidth.

Examples:

- *Coalesced (efficient):*

```
// Each thread accesses consecutive elements
value = A[threadId];
```

→ One transaction for 32 threads.

- *Non-coalesced (inefficient):*

```
// Each thread accesses strided elements
value = A[threadId * 100];
```

→ Many transactions required due to striding.

2. Gather—Scatter in Sparse Matrix Access

Definition: Gather—scatter refers to memory operations where each GPU thread in a Warp loads from or stores to irregular memory addresses. This is common in sparse matrix operations, where non-zero elements are stored in compressed formats.

Sparse Matrix Example (CSR format):

- *CSR (Compressed Sparse Row)* stores three arrays:
 - `values[]`: non-zero entries of the matrix

²² Page 300 of Computer Architecture: A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design)

²³ Page 280 of Computer Architecture: A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design)

- colIndex[]: column indices for each non-zero
- rowPtr[]: index into values[] for each row

- Sparse matrix-vector multiplication (SpMV):

```
for row in matrix:
    for idx = rowPtr[row] to rowPtr[row+1]:
        col = colIndex[idx];      // gather index
        val = values[idx];        // gather nonzero
        y[row] += val * x[col];    // scatter result
```

Characteristics:

- **Gather:** Each thread loads from potentially scattered locations (values[idx] or x[col]).
- **Scatter:** Results may be written back to irregular output locations (y[row]).
- **Challenge:** These accesses often break memory coalescing, leading to multiple memory transactions. An example is shown as follows:

Summary:

- Gather—scatter is fundamental for sparse matrix access but typically results in non-coalesced memory patterns.
- Address coalescing is critical for high GPU throughput; restructuring data to improve coalescing often provides significant performance gains.

6.4.2 VRAM dGPU




Fig. 6.17: iGPU versus dGPU

DMA = Direct Memory Access, meaning a hardware unit can read/write memory without CPU involvement.




Fig. 6.18: CPU send Uniform Updates Data to dGPU for rendering




Fig. 6.19: CPU send Uniform Updates Data to iGPU for rendering

Reason:

1. Even though the **uniform updates are “zero-copy” in Integrated GPUs**, the GPU must fetch **all rendering data** (textures, vertex buffers, render targets) from system RAM, which is:

- 4–10× lower bandwidth than VRAM
- shared with the CPU
- higher latency
- less optimized for parallel GPU access

2. Since CPU and GPU have different requirements, a shared memory design cannot match the performance of dedicated GPU memory.

3. In systems with shared memory (like integrated GPUs), both the CPU and GPU access the same physical memory (DRAM). This leads to several forms of contention:

- a. Cache Coherency Overhead
- b. DMA Contention
- c. Bus & Memory Controller Bottleneck

Advantages:

A discrete GPU has its own dedicated memory (VRAM) while an integrated GPU (iGPU) shares memory with the CPU as shown in Fig. 6.17.

Dedicated GPU memory (VRAM) outperforms shared CPU-GPU memory due to higher bandwidth, lower latency, parallel access optimization, and no contention with CPU resources.

Feature	Shared Memory (CPU + iGPU)	Dedicated GPU Memory (dGPU)
Bandwidth	Lower (DDR/LPDDR)	Higher (GDDR/HBM)
Latency	Higher	Lower
Parallel Access	Limited	Optimized for many threads
Cache Coherency	Required (with CPU)	Not required
DMA Bandwidth	Shared with CPU	GPU has exclusive DMA access
Memory Contention	Yes	No
Performance	Lower: Bandwidth bottlenecks, CPU-GPU interference and Cache/DMA conflicts	Higher: Wide memory bandwidth, Parallel thread access and Low latency memory access

Dedicated memory allows the GPU to run high-throughput workloads without interference from the CPU. It provides

(1). **wide bandwidth**, (2). **optimized parallel access**, and (3). **low-latency paths**, avoiding cache and DMA conflicts for superior performance.**

(1). **Wide bandwidth:** Dedicated GPU memory (VRAM) is often based on GDDR6, GDDR6X, or HBM2/3, which are much faster than standard system RAM (DDR4/DDR5).

Typical bandwidths:

- GDDR6: ~448—768 GB/s
- HBM2: up to 1 TB/s+
- DDR5 (shared memory): ~50—80 GB/s

Impact: Faster access to textures, vertex buffers, and framebuffers—critical for rendering and compute tasks.

(2). **Optimized parallel access:**

- VRAM is optimized for the massively parallel architecture of GPUs.
- It allows thousands of threads to access memory simultaneously without stalling.

Shared system memory is optimized for CPU access patterns, not thousands of GPU threads.

(3). **Low-latency paths:**

- Dedicated memory is physically closer to the GPU die.
- No need to traverse the PCIe bus like discrete GPUs accessing system RAM.

In shared memory systems (like integrated GPUs), memory access may have to go through a memory controller shared with the CPU, adding delay.

6.4.3 RegLess-style architectures²⁴

Note

RegLess remains a research concept, not (as far as public evidence shows) a commercial design in shipping GPUs.

Difference: Add **Staging Buffer** between Register Files and Execution Unit.

This section outlines the transition from traditional GPU operand coherence using a monolithic register file and L1 data cache, to a RegLess-style architecture that employs operand staging and register file-local coherence.

✓ Operand Delivery in Traditional GPU: Fig. 6.20:




Fig. 6.20: **Operand Delivery in Traditional GPU (Traditional Model: Register File + L1 Cache)**

Architecture:

²⁴ <https://cccp.eecs.umich.edu/papers/jklooste-micro17.pdf>

- Large monolithic register file per SM (e.g., 256KB, Maxwell, Pascal, Volta and Ampere have 64K x 32-bit register file per SM, see *Configurable maximum resident warps and allocated registers per thread*)
- Coherent with L1 data cache via write-through or write-back policies

Challenges:

- High energy cost due to cache coherence traffic
- Complex invalidation and synchronization logic
- Register pressure limits Warp occupancy (limit the number of active Warps)
- Redundant operand tracking across register file and cache

Example:

```

v1 = normalize(N)
v2 = normalize(L)
v3 = dot(v1, v2)
v4 = max(v3, 0.0)
v5 = mul(v4, color)

# All operands reside in register file and may be cached in L1

```

✓ Operand Delivery in RegLess GPU (with L1 Cache in LD Path): Fig. 6.21:




Fig. 6.21: **Operand Delivery in RegLess GPU (with L1 Cache in LD Path)**

Description

- **Global Memory:** Source of all operands and data.
- **L1 Cache:** Participates in memory hierarchy; may serve LD requests.
- **Register File:** Receives operands via LD; stages them into Staging Buffer for Transient Operands.
- **Staging Buffer:** Holds transient operands for immediate execution.
- **Execution Unit:** Consumes operands from Staging Buffer for Transient Operands and Register File for Persistent Operands.

Notes

- **L1 Cache is not part of staging**—it only serves LDs.
- **Dashed arrows:** Compiler-controlled operand movement.
- **Solid arrows:** Operand delivery to execution.
- **Green self-loop:** Internal coherence within Register File.

RegLess Model: Staging-Aware Register File

Architecture:

- Smaller register file (e.g., 64–128KB per SM)
- For Transient Operands, no L1 cache coherence required
- Operands staged dynamically based on lifetime

Key Concepts:

- Region slicing: compiler divides computation into operand regions
- Operand tagging: transient, intermediate, persistent
- Metadata compression: region-level hints, not per-instruction lifetimes

Benefits:

- ~75% reduction in register file size
- ~11% energy savings
- Simplified coherence model
- Improved Warp occupancy

Example with Operand Staging:

```
v1 = normalize(N)           # transient
v2 = normalize(L)           # transient
v3 = dot(v1, v2)            # intermediate
v4 = max(v3, 0.0)           # intermediate
v5 = mul(v4, color)         # persistent

# v1 and v2 staged briefly, v3–v4 may be staged or registered, v5 fully
# registered
```

Compiler-Hardware Interface

Compiler Responsibilities:

- Emit structured IR with operand usage hints
- Slice computation graph into regions
- Avoid explicit staging register allocation

Hardware Responsibilities:

- Interpret operand lifetime metadata
- Dynamically stage operands or allocate registers
- For Transient Operands, eliminate L1 cache coherence logic

Metadata Compression Techniques:

- Region-level tagging
- Operand class encoding
- Profile-guided optimization
- Off-chip metadata tables (e.g., DEER)

Conclusion

The move to RegLess-style coherence simplifies GPU operand management, reduces energy, and enables more efficient shader execution. Compiler-guided operand staging and region slicing allow hardware to dynamically optimize operand placement without burdening the instruction stream with excessive metadata.

6.5.1 Geometry Units

Function:

Raw Vertices & Primitives → Transformed Vertices & Primitives

The GS does not output arbitrary primitive types. GS only output:

- point_strip
- line_strip
- triangle_strip

Suppose the GLSL geometry shader looks like this:

An example of GLSL geometry shader

```
#version 330 core

layout(triangles) in;
layout(triangle_strip, max_vertices = 3) out;

void main()
{
    // Forward the input triangle unchanged.
    for (int i = 0; i < 3; ++i) {
        gl_Position = gl_in[i].gl_Position;
        EmitVertex();
    }

    EndPrimitive();
}
```

The corresponding PTX instructions and pipeline flow as Fig. 6.24.

The **Geometry Unit** in a GPU is a collection of fixed-function and programmable stages responsible for transforming assembled primitives (points, lines, triangles, patches) into screen-space primitives ready for rasterization. The emit and cut are compiler intrinsics that map to control messages to the Geometry Unit. When we say emit and cut in NVIDIA PTX (or HLSL/GLSL geometry shaders), they're not ALU instructions that run in the SM like add or mul. Instead, they act like special control instructions that tell the GPU's fixed-function Geometry Unit what to do with the vertex data currently in the SM's output registers illustrated in Fig. 6.25.

Unlike GLSL textures, which are converted into a specific hardware ISA, the Geometry Shader in Fig. 6.22 maps directly to the Geometry Units instead of the SMs.

Geometry Unit bridges the **vertex shading** stage and the **rasterization** stage as shown in Fig. 6.26.

Role

- Organize and process geometry data after vertex shading.
- Perform primitive-level operations such as assembly, tessellation, clipping, viewport transform, and primitive setup.
- Provide hardware acceleration for geometry amplification or reduction before rasterization.

Components

- **Primitive Assembly (Input Assembler)**
 - Groups vertices into primitives (triangles, lines, patches).




Fig. 6.24: Fetch a sequence of Geometry instructions and pass to Geometry Unit




Fig. 6.25: Micro-level flow: SM → Geometry Unit via Emit/Cut




Fig. 6.26: Geometry Unit with its sub-functions (assembly, tessellation, clipping, viewport transform, etc.)

- Fetches indices and vertex attributes from memory.
- Prepares data structures for downstream geometry stages.
- **Tessellation Engine (optional, OpenGL 4.0+ / DirectX 11+)**
 - Subdivides patches into finer primitives.
 - Contains Tessellation Control Shader, Primitive Generator, and Tessellation Evaluation Shader.
 - Used in terrain rendering, displacement mapping, and adaptive LOD.
- **Geometry Shader (optional, programmable stage)**
 - Can generate new primitives or discard existing ones.
 - Enables shadow volume extrusion, point sprite expansion, or procedural geometry.
 - High flexibility but often limited in performance due to amplification.
- **Culling & Clipping**
 - Removes back-facing or out-of-view primitives.
 - Clips primitives against the view frustum or user-defined clipping planes.
 - Optimizes rendering by reducing fragment processing workload.
- **Viewport Transform**
 - Maps Normalized Device Coordinates (NDC) to screen-space pixel coordinates.
 - Applies viewport scaling, offset, and depth range mapping.
- **Primitive Setup**
 - Converts screen-space primitives into edge equations and interpolation rules.
 - Prepares slopes and barycentric coefficients for attribute interpolation in rasterization.
 - Ensures that per-fragment attributes (e.g., texture coordinates, normals) are interpolated correctly.

Usage

- Reduces workload on the fragment stage by culling invisible primitives.
- Provides tessellation and geometry shaders for advanced rendering effects.
- Ensures efficient and accurate rasterization setup.
- Works closely with specialized GPU fixed-function blocks such as **PolyMorph Engines** (NVIDIA) or **Geometry Processors** (AMD).

References

- Wikipedia —[Graphics pipeline](#)
- NVIDIA —[DirectX 11 GPU Architecture \(Geometry and PolyMorph Engine\)](#)
- Intel —[3D Pipeline Overview \(including Geometry Stage\)](#)
- LearnOpenGL —[Geometry Shader](#)
- Microsoft Docs —[Tessellation and Geometry Pipeline](#)

6.5.2 Rasterization Units²⁵

Function:

Transformed Vertices & Primitives → Fragments

Overview

The rasterization unit is a critical component of the graphics pipeline in modern GPUs. It converts geometric primitives (typically triangles) into fragments that correspond to pixels on the screen. This process is essential for rendering 3D scenes into 2D images.

The pipeline flow for Rasterization Units is shown as [Fig. 6.27](#).




Fig. 6.27: Rasterization pipeline

Key Functions

- **Triangle Setup:** Computes edge equations and bounding boxes for each triangle.
- **Scan Conversion:** Determines which pixels are covered by the triangle.
- **Attribute Interpolation:** Calculates interpolated values (e.g., texture coordinates, depth) for each fragment.
- **Fragment Generation:** Produces fragment data for downstream shading and blending stages.

²⁵ copilot: Please provide detailed information about the Rasterization Unit and its pipeline, including a separated dot graph and relevant website references in reStructuredText (reST) format.

Hardware Architecture

Modern GPUs implement rasterization in highly parallel hardware blocks to maximize throughput. A simplified block diagram includes:

- **Primitive Assembly Unit:** Groups vertices into triangles.
- **Triangle Setup Engine:** Prepares edge equations and bounding boxes.
- **Rasterizer Core:** Performs scan conversion and fragment generation.
- **Early-Z Unit:** Performs early depth testing to discard hidden fragments.

Clipping and culling (in the geometry/primitive stage) only guarantee:

- triangles outside the frustum are removed
- triangles behind the camera are removed
- triangles partly outside the frustum are clipped
- degenerate triangles are removed

But they do not determine which triangle is in front of another at a given pixel. For example when two triangles overlap on screen, both pass clipping and culling, and get rasterized and produce fragment.

- **Fragment Queue:** Buffers fragments for shading.

Optimization Techniques

- **Tile-Based Rasterization:** Divides the screen into tiles to reduce memory bandwidth.
- **Early-Z Culling:** Discards fragments before shading if they fail depth tests.
- **Compression:** Reduces data transfer costs between pipeline stages.

Use Cases

- Real-time rendering in games and simulations.
- 3D Gaussian Splatting acceleration for AI-based rendering.
- Mobile GPUs with power-efficient rasterization pipelines.

References

- [GauRast: Enhancing GPU Triangle Rasterizers](#)
- [NVIDIA Ada GPU Architecture PDF](#)
- [Stanford CS248A Lecture on Rasterization](#)

6.5.3 Texture Mapping Units (TMUs)²⁶

Function:

Fragments → Processed Fragments

Overview

A Texture Mapping Unit (TMU) is a fixed-function hardware block inside a GPU responsible for *fetching, filtering, and preparing texture data* that shaders (sampled in fragment or compute stages) use during rendering.

As explained in previous [section OpenGL Shader Compiler](#), the texture instruction using TMU to accelerate calculation as the following explanation with [Fig. 6.28](#).

²⁶ <http://math.hws.edu/graphicsbook/c6/s4.html>

TMUs sit between the shader cores (SMs/CUs) and the memory subsystem. They provide high-performance, specialized texture access operations that would be too slow or costly to emulate in general-purpose ALUs is shown as Fig. 6.28.




Fig. 6.28: The flow of issuing texture instruction from SM to TMU.

Pipeline Role

- In the **OpenGL / Direct3D graphics pipeline**, TMUs are mainly used in the *fragment shading stage*, where textured surfaces are shaded with data from 2D/3D textures.
- In **compute shaders**, TMUs are also used for image load/store operations and texture sampling.

Key Responsibilities

1. Texture Addressing

- Compute the correct texture coordinate for a given fragment or pixel.
- Handle the following wrapping modes are shown as Fig. 6.29 and as Fig. 6.30:

Texture coordinates usually range from (0,0) to (1,1) but what happens if we specify coordinates outside this range? OpenGL provides the following wrapping modes for outside this range.

- Clamp-to-border (GL_CLAMP_TO_BORDER)
 - * When a texture coordinate falls outside the [0,1] range, the GPU does not sample the nearest texel.
 - * Instead, it returns a user-defined border color for that texture.
 - * This is useful for effects like shadow maps, where sampling outside the valid area should produce a consistent value.
- Repeat (GL_REPEAT): Wraps coordinates around (tiles the texture).
- Clamp-to-edge (GL_CLAMP_TO_EDGE): Uses the edge texel when coordinates are out of range.
- Mirrored repeat (GL_MIRRORED_REPEAT): Mirrors the texture each repetition.
 - * For the middle row (t(V) in the range 0.0 to 1.0), the mirroring operation applies only a left-right swap. For the top and bottom rows, the mirroring includes both left-right and up-down swaps.

²⁷ <https://learnopengl.com/Getting-started/Textures>




Fig. 6.29: Texture Warpping




Fig. 6.30: Texture Warpping^{Page 149, 27}

- Convert normalized texture coordinates into actual memory addresses.

2. Texture Fetching

- Retrieve texels (texture elements) from texture memory (L1 texture cache, then L2/VRAM on miss).
- Handle different texture layouts:
 - 1D, 2D, 3D textures
 - Cubemaps
 - Texture arrays
- Support compressed texture formats (e.g., DXT, ASTC, ETC2).

3. Texture Filtering

Give a Texture coordinates, OpenGL has to figure out which **texture pixel (also known as a texel)** to map the texture coordinate to.

- Perform *interpolation* between texels to produce smooth visual results.
- Filtering requires multiple texel reads + weighted average calculations.
- Common filtering modes as the following are shown as Fig. 6.33:
 - Nearest-neighbor (point sampling) (GL_NEAREST)
 - * When set to GL_NEAREST, OpenGL selects the color of the texel that center is closest to the texture coordinate shown as the example in Fig. 6.31. ‘+’ is the coordinates of texel. ‘Returns’ is the color of result.



Fig. 6.31: GL_NEAREST^{Page 149, 27}

– Bilinear (GL_LINEAR)

- * The return color is the mix of four neighboring pixels. The smaller the distance from the texture coordinate to a texel's center, the more that texel's color contributes to the sampled color shown as the example in Fig. 6.32.



Fig. 6.32: GL_LINEAR^{Page 149, 27}

- Trilinear (with mipmaps)
- Anisotropic filtering (for angled surfaces)
- Let's see how these methods work when using a texture with a low resolution on a large object (texture is therefore scaled upwards and individual texels are noticeable). The GL_NEAREST and GL_LINEAR as the following Fig. 6.33. As a result, GL_LINEAR produces a blurrier image with smoother edges.



Fig. 6.33: Texture Filter: GL_NEAREST has sharp color and jagged edge^{Page 149, 27}

4. Mipmap Level of Detail (LOD) Selection

- Choose the correct mipmap level based on screen-space derivatives of texture coordinates.
- Prevent aliasing and improve cache efficiency.
- Optionally blend between mip levels for trilinear filtering.

5. Texture Caching

- TMUs have a **dedicated texture cache** optimized for 2D/3D spatial locality.
- Neighboring threads in a Warp often fetch adjacent texels, improving cache hits.
- Caches reduce memory latency and improve bandwidth utilization.

6. Specialized Operations

- Texture gather: fetch 4 neighboring texels around a coordinate.
- Shadow mapping: compare fetched depth texel against reference value.
- Multisample textures: fetch per-sample data for MSAA.
- Border color application for out-of-bounds accesses.

Microarchitecture Aspects

- Each **Streaming Multiprocessor (SM)** or **Compute Unit (CU)** is paired with several TMUs.
- The number of TMUs is a key spec in GPU datasheets (e.g., “64 TMUs”).
- TMU throughput is often measured in **texels per clock cycle**.
- Modern GPUs balance **TMUs per ALU** to ensure shading and texture workloads are not bottlenecked.

Performance Considerations

- **Bandwidth-limited:** TMUs rely heavily on memory bandwidth. Mipmapping and caches reduce this pressure.
- **Latency hiding:** texture fetches may take hundreds of cycles, so GPUs rely on massive multithreading to hide stalls.
- **Workload dependent:** texture-heavy games or rendering pipelines are often limited by TMU throughput.

Summary

TMUs are highly specialized GPU units that:

- Translate texture coordinates into addresses.
- Fetch texels efficiently with dedicated caches.
- Perform filtering and LOD computations in hardware.
- Deliver high throughput for texture operations that are essential in realistic rendering.

Without TMUs, all these operations would fall on general-purpose ALUs, resulting in drastically lower performance and efficiency.

6.5.4 Render Output Units (ROPs)²⁸

Function:

Processed Fragments → Pixels

Overview

Render Output Units (ROPs), also known as Raster Operations Pipelines, are the final stage in the GPU graphics pipeline before pixel data is written to the framebuffer. ROPs handle pixel-level operations such as blending, depth and stencil testing, multisample resolve, and writing to memory. They are crucial for assembling the final image that appears on screen.

Pipeline Responsibilities

- **Fragment Reception:** Accepts shaded fragments from the pixel shader.
- **Depth and Stencil Testing:** Compares fragment depth/stencil values against buffers.
- **Blending:** Combines fragment color with existing framebuffer data.

²⁸ copilot: Please provide detailed information about the Render Output Units (ROPs) and its pipeline, including a dot graph and relevant website references in reStructuredText (reST) format.

- **Multisample Resolve:** Merges multiple samples into a final pixel (for MSAA).
- **Framebuffer Write:** Commits final pixel data to memory for display.

The pipeline flow is shown as Fig. 6.34.




Fig. 6.34: The pipeline for Render Output Units (ROPs)

Performance Considerations

- **ROP Count:** More ROPs can increase pixel throughput, especially at high resolutions.
- **Memory Bandwidth:** ROPs are tightly coupled with memory controllers; bandwidth limits can bottleneck performance.
- **Antialiasing Support:** Hardware MSAA and resolve operations are often implemented in ROPs.
- **Compression:** Some GPUs use framebuffer compression to reduce bandwidth usage.

Vendor-Specific Notes

- **NVIDIA:** Refers to these units as ROPs; tightly integrated with memory partitions.
- **AMD:** Calls them Render Backends (RBs); RDNA architecture decouples ROPs from shader engines.
- **Intel & ARM:** Implement simplified ROPs for power-efficient mobile rendering.

References

- [Render Output Unit - Wikipedia](#)
- [What is a ROP on a GPU? - CORSAIR](#)
- [TechPowerUp Forums: ROPs and TMUs](#)

6.5.5 Framebuffer hierarchy

The framebuffer is a collection of per-pixel attachments. The color buffer is only one part of it. A complete framebuffer has a clear hierarchy: Framebuffer → Attachments → Per-pixel storage (color/depth/stencil).

```

Framebuffer (FBO)
├── Color Attachment 0 → Color Buffer (RGBA)
├── Color Attachment 1 → Color Buffer (RGBA)
├── Depth Attachment → Depth Buffer (float)
├── Stencil Attachment → Stencil Buffer (8-bit)
└── Depth-Stencil Attachment → Combined Buffer
  
```

The LCD/LED screen does not directly read the GPU color buffer. Instead, the GPU's display engine scans out the final framebuffer and sends it to the panel's TCON, which then drives the physical pixels.

GPU Rendering Pipeline





Fig. 6.35: TCON

. _software-struct:

SOFTWARE STRUCTURE

- *Vector Processor*
- *General purpose GPU*
 - *Mapping data in GPU*
 - *Work between CPU and GPU in Cuda*
- *OpenCL, Vulkan and Spir-v*
 - *Summary*
- *Unified IR Conversion Flows*
 - *Graphics Compilation Flow (Microsoft DirectX & OpenGL)*
 - *ML and GPU Compilation*
 - * *NVIDIA IR Conversion Flow*
 - * *AMD IR Conversion Flow*
 - * *ARM IR Conversion Flow*
 - * *Imagination Technologies IR Conversion Flow*
 - * *Comparison Summary*
 - * *References*
- *Accelerate ML/DL on OpenCL/SYCL*

As the previous section illustrated, GPU is a SIMT (SIMD) for data parallel application. This section introduces the GPU evolved from Graphics GPU to the General purpose GPU (GPGPU) and the software architecture of GPUs and explores AI software frameworks designed for GPUs, NPU, and CPUs.

7.1 Vector Processor

As described in the Computer Architecture: A Quantitative Approach book, the vector processor VMIPS introduces the **Vector Length Register (VLR) and Vector Mask (VM)** to support SIMD execution. **The Vector Mask functions similarly to conditional instructions in CPUs.** In vector processors, VM acts as a form of conditional execution mechanism.

✓ Vector-Length Registers: Handling Loops Not Equal to 64

```
for (i=0; i <n; i=i+1)
    Y[i] = a * X[i] + Y[i];
```

In the preceding code, the value of n is not known at compile time.

Solution:

The compiler converts the loop into multiple iterations, each of which processes up to the maximum vector length (MVL), as shown below. For VMIPS, the MVL is 64.

```
low = 0;
VL = (n % MVL); /*find odd-size piece using modulo op % */
for (j = 0; j <= (n/MVL); j=j+1) { /*outer loop*/
    for (i = low; i < (low+VL); i=i+1) /*runs for length VL*/
        Y[i] = a * X[i] + Y[i] ; /*main operation*/
    low = low + VL; /*start of next vector*/
    VL = MVL; /*reset the length to maximum vector length*/
}
```

The inner loop of the preceding code is vectorizable with length VL, which is equal to either $(n \% \text{MVL})$ or MVL. The VLR register must be set twice in the code, once at each place where the variable VL in the code is assigned.

✓ Vector Mask Registers: Handling IF Statements in Vector Loops

```
for (i = 0; i < 64; i=i+1)
    if (X[i] != 0)
        X[i] = X[i] - Y[i];
```

For the VMIPS vector processor, the above code can be implemented using the Vector Length Register (VLR) as shown below.

Assembly code of Vector Processor (from page 276 of Quantitative)

```
LV V1,Rx          ;load vector X into V1
LV V2,Ry          ;load vector Y
L.D F0,#0         ;load FP zero into F0
SNEVS.D V1,F0     ;sets VM(i) to 1 if V1(i)!=F0
SUBVV.D V1,V1,V2  ;subtract under vector mask
SV V1,Rx          ;store the result in X
```

- Code reference here³.

7.2 General purpose GPU

Because GLSL shaders provide a C-like programming model, a software framework can use GPU shaders for data-parallel computation instead of only graphics rendering. This can accelerate work and allow CPU and GPU execution to overlap. Furthermore, any language that lets CPU code retrieve results from a GPU shader can support a GPGPU framework¹.

³ subsection Vector Mask Registers: Handling IF Statements in Vector Loops of Computer Architecture: A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design)

¹ https://en.wikipedia.org/wiki/General-purpose_computing_on_graphics_processing_units

7.2.1 Mapping data in GPU

As described in the previous section on GPUs, the subset of the array calculation $y[] = a * x[] + y[]$.

```
// Invoke DAXPY with 256 threads per Thread Block
__host__
int nblocks = (n+255) / 256;
daxpy<<<nblocks, 256>>>(2.0, x, y);
// DAXPY in CUDA
__device__
void daxpy(double a, double *x, double *y) {
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    a*x[i] + y[i];
}
```

- name<<<dimGrid, dimBlock>>>(...parameter list ...):
 - dimGrid: Number of Blocks in Grid
 - dimBlock: 256 Threads in Block

Assembly code of PTX (from page 300 of Quantitative book)

```
// The following sequence of PTX instructions is for one iteration of the
// DAXPY loop above.
shl.u32 R8, blockIdx, 9          ; Thread Block ID * Block size (512)
add.u32 R8, R8, threadIdx        ; R8 = i = my CUDA Thread ID
shl.u32 R8, R8, 3                 ; byte offset
ld.global.f64 RD0, [X+R8]        ; RD0 = X[i]
ld.global.f64 RD2, [Y+R8]        ; RD2 = Y[i]
mul.f64 RD0, RD0, RD4             ; Product in RD0 = RD0 * RD4 (scalar a)
add.f64 RD0, RD0, RD2            ; SuminRD0 = RD0 + RD2 (Y[i])
st.global.f64 [Y+R8], RD0        ; Y[i] = sum (X[i]*a + Y[i])
```

✓ Conditional Branching in GPUs⁴:

Assembly code of PTX (from referring page 302 of Quantitative book)

```
__device__
void lane-mask-ex( double *X, double *Y, double *Z) {
    if (X[i] != 0)
        X[i] = X[i] - Y[i];
    else X[i] = Z[i];
}
```

- Code from here⁵.

The following two instructions illustrate **conditional (predicated) instruction execution on GPUs**.

```
predicate = cond          // predicate is the mask register
@predicate instruction
```

⁴ subsection Conditional Branching in GPUs, page 300 - 303 of Computer Architecture: A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design)

⁵ Code written by referring page 302 of Computer Architecture: A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design)

This IF statement could compile to the following PTX instructions (assuming that R8 already has the scaled thread ID), with *Push, *Comp, *Pop indicating the branch synchronization markers inserted by the PTX assembler that push the old mask, complement the current mask, and pop to restore the old mask:

Assembly code of PTX (from referring page 302 of Quantitative book)

```
ld.global.f64 RD0, [X+R9]      ; RD0 = X[i]
setp.neq.s32 P1, RD0, #0       ; P1 is predicate register 1
@!P1, bra ELSE1, *Push         ; Push old mask, set new mask bits
                                ; if P1 false, go to ELSE1
ld.global.f64 RD2, [Y+R8]      ; RD2 = Y[i]
sub.f64 RD0, RD0, RD2          ; Difference in RD0
st.global.f64 [X+R8], RD0      ; X[i]=RD0
ELSE1:
ld.global.f64 RD0, [Z+R8]      ; RD0 = Z[i]
st.global.f64 [X+R8], RD0      ; X[i] = RD0
ENDIF1:
ret, *Pop                      ; pop to restore old mask
```

The PTX:

```
setp.neq.s32 P1, RD0, #0
```

On actual NVIDIA hardware (SASS), the instruction typically becomes:

```
ISETP.NE.AND P1, PT, R0, RZ, PT ; PT is always-true predicate
```

This instruction compares the 32-bit signed integer value in register RD0 with the constant 0. The result of the comparison is written to the predicate register P1.

Semantics:

```
P1 = (RD0 != 0)
```

Each thread (lane) in the warp evaluates this comparison independently.

Example (8-thread warp):

```
RD0 values:  [5, 3, 0, 7, 1, 0, 4, 2]
P1 result:   [1, 1, 0, 1, 1, 0, 1, 1]
```

This instruction **does not modify the active thread mask**. It only produces a predicate value that will be used by later predicated instructions or branches.

```
@!P1 bra ELSE1, *Push
```

This is a predicated branch instruction.

The prefix @!P1 means the instruction executes only for threads where the predicate P1 is false.

Semantics:

```
if (!P1)
    branch to ELSE1
```

If all threads agree on the predicate value, the warp simply branches or falls through. However, if some threads have P1 = 1 and others have P1 = 0, **control flow divergence occurs**.

Mask Stack Operation

The `*Push` modifier indicates that the hardware must update the SIMT control-flow stack.

When divergence occurs:

1. The current active mask is pushed onto the stack.
2. The warp execution mask is split into two masks:
 - `mask_then = active_mask & P1`
 - `mask_else = active_mask & !P1`
3. Execution proceeds with one mask while the other path is saved for later execution.

Conceptual behavior:

```
push(active_mask)

mask_then = active_mask & P1
mask_else = active_mask & !P1

execute THEN block using mask_then
later execute ELSE block using mask_else
```

4. Threads in the THEN mask execute the fall-through path, while threads in the ELSE mask branch to `ELSE1`.
5. Keep in mind, however, that the only choice for a SIMD Lane in a clock cycle is to perform the operation specified in the PTX instruction or be idle; two SIMD Lanes cannot simultaneously execute different instructions^{Page 157, 4}.

The following table explains how the elements of `saxpy()` are mapped to the Lanes of a SIMD Thread (Warp), which belongs to a Thread Block (Core) within a Grid.

Table 7.1: Mapping saxpy code to Fig. 6.15.

saxpy()	Instance in Fig. 6.15	Description
block-Dim.x	The index of Thread Block	blockDim: in this example configured as Fig. 6.15 is 16(Thread Blocks) * 16(SIMD Threads) = 256
block-Idx.x	The index of SIMD Thread	blockIdx: the index of Thread Block within the Grid
thread-Idx.x	The index of elements	threadIdx: the index of the SIMD Thread within its Thread Block

- With Fermi, each 32-wide thread of SIMD instructions is mapped to 16 physical SIMD Lanes, so each SIMD instruction in a thread of SIMD instructions takes two clock cycles to complete.
- You could say that it has 16 Lanes, the vector length would be 32, and the Chime is 2 clock cycles.
- The map of `y[0..31] = a * x[0..31] * y[0..31]` to `<Core, Warp, Cuda Thread>` of GPU as the following table. `x[0..31]` map to 32 Cuda Threads; **two Cuda Threads map to one SIMD Lane** as Fig. 6.14..

Table 7.2: Map `<Core, Warp>` to saxpy

	Warp-0	Warp-1	...	Warp-15
Core-0	<code>y[0..31] = a * x[0..31] * y[0..31]</code>	<code>y[32..63] = a * x[32..63] + y[32..63]</code>	...	<code>y[480..511] = a * x[480..511] + y[480..511]</code>
...	...	...	...	...
Core-15	<code>y[7680..7711] = a * ...</code>	...	...	<code>y[8160..8191] = a * x[8160..8191] + y[8160..8191]</code>

- Each Cuda Thread runs the GPU function code *saxpy*. Fermi has a register file of size 32768 x 32-bit. As shown in Fig. 6.6, the number of registers in a Thread Block is: 16 (SM) * 32 (Cuda Threads) * 64 (TLR, Thread Level Register) = 32768 x 32-bit (Register file).
- When mapping to fragments/pixels in a graphics GPU, $x[0..15]$ corresponds to a two-dimensional tile of fragments/pixels at $pixel[0..3][0..3]$, since images use tile-based grouping to cluster similar colors together.

7.2.2 Work between CPU and GPU in Cuda

The previous *daxpy()* GPU code did not include the host (CPU) side code that triggers the GPU function.

The following example shows the host (CPU) side of a CUDA program that calls *saxpy* on the GPU²:

```
#include <stdio.h>

__global__
void saxpy(int n, float a, float * x, float * y)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < n) y[i] = a*x[i] + y[i];
}

int main(void)
{
    ...
    cudaMemcpy(d_x, x, N*sizeof(float), cudaMemcpyHostToDevice);
    cudaMemcpy(d_y, y, N*sizeof(float), cudaMemcpyHostToDevice);
    ...
    saxpy<<<(N+255)/256, 256>>>(N, 2.0, d_x, d_y);
    ...
    cudaMemcpy(y, d_y, N*sizeof(float), cudaMemcpyDeviceToHost);
    ...
}
```

The *main()* function runs on the CPU, while *saxpy()* runs on the GPU. The CPU copies data from *x* and *y* to the corresponding device arrays *d_x* and *d_y* using *cudaMemcpy*.

The *saxpy* kernel is launched with the following statement:

```
saxpy<<<(N+255)/256, 256>>>(N, 2.0, d_x, d_y);
```

This launches the kernel with Thread Blocks containing 256 threads, and uses integer arithmetic to determine the number of Thread Blocks needed to process all *N* elements in the arrays. The expression $(N+255)/256$ ensures full coverage of the input data.

Using *cudaMemcpyHostToDevice* and *cudaMemcpyDeviceToHost*, the CPU can pass data in *x* and *y* to the GPU, and retrieve the results back to *y*.

Since both memory transfers are handled by DMA and do not require CPU operation, the performance can be improved by running CPU and GPU independently, each accessing their own cache.

After the DMA copy from CPU memory to GPU memory, the GPU performs the full matrix operation loop for $y[i] = a * x[i] + y[i]$; using a single Grid of threads.

DMA *memcpy* maps the data in CPU memory to each L1 cache of a core on GPU memory.

² <https://devblogs.nvidia.com/easy-introduction-cuda-c-and-c/>

Many GPUs support scatter and gather operations to access DRAM efficiently for stream processing tasks⁶Page 156, 17.

When the GPU function is dense computation in array such as MPEG4 encoder or deep learning for tuning weights, it may get much speed up⁸. However when GPU function is matrix addition and CPU will idle for waiting GPU's result. It may slow down than doing matrix addition by CPU only. Arithmetic intensity is defined as the number of operations performed per word of memory transferred. It is important for GPGPU applications to have high arithmetic intensity else the memory access latency will limit computational speedup^{Page 156, 1}.

Wiki here⁹ includes GPU-accelerated applications for speedup as follows:

General Purpose Computing on GPU, has found its way into fields as diverse as machine learning, oil exploration, scientific image processing, linear algebra, statistics, 3D reconstruction and even stock options pricing determination. In addition, section "GPU accelerated video decoding and encoding" for video compressing⁹ gives the more applications for GPU acceleration.

Table 7.3: The differences for speedup in architecture of CPU and GPU

Item	CPU	GPU
Application	Non-data parallel	Data parallel
Architecture	SISD, small vector (eg.4*32bits)	Large SIMD (eg.16*32bits)
Cache	Smaller and faster	Larger and slower (ref. The following Note)
ILP	Pipeline	Pipeline
	Superscalar, SMT	SIMT
	Super-pipeline	
Core	Smaller threads for SMT (2 or 4)	Larger threads (16 or 32)
Branch	Conditional-instructions	Mask & conditional-instructions

Note

GPU-Cache

In theory, for data-parallel applications using GPU's SMT, the GPU can schedule more threads and aims for throughput rather than speedup of a single thread, as seen in SISD on CPUs.

However, in practice, GPUs provide only a small L1 cache, similar to CPUs, and handle cache misses by scheduling another thread.

As a result, GPUs often lack L2 and L3 caches, which are common in CPUs with deeper cache hierarchies.

7.3 OpenCL, Vulkan and Spir-v

Table 7.4: OpenCL and OpenGL SW system

Name of SW	GPU language	Level of GPU language
OpenCL	OpenCL	C99 dialect (with C pointer, ...)
OpenGL	GLSL	C-like (no C pointer, ...)
Vulkan	SPIR-V	IR

⁶ Reference "Gather-Scatter: Handling Sparse Matrices in Vector Architectures": section 4.2 Vector Architecture of A Quantitative Approach 5th edition (The Morgan Kaufmann Series in Computer Architecture and Design)

⁷ The whole chip shares a single L2 cache, but the different units will have individual L1 caches. <https://computergraphics.stackexchange.com/questions/355/how-does-texture-cache-work-considering-multiple-shader-units>

⁸ <https://www.manchestervideo.com/2016/06/11/speed-h-264-encoding-budget-gpu/>

⁹ https://en.wikipedia.org/wiki/Graphics_processing_unit

GLSL	OpenCL
SPIR-V	
GPU Machine Code	

Fig. 7.1: OpenCL and GLSL(OpenGL)




Fig. 7.2: Offline Compilation of OpenCL Kernels into SPIR-V Using Open Source Tooling

- clang: Compiles OpenCL to SPIR-V for the runtime and driver. Alternatively, it compiles OpenCL to LLVM, after which the SPIR-V—LLVM translator converts LLVM to SPIR-V.
- clspv: Compiles OpenCL to SPIR-V directly.




Fig. 7.3: GPU Compiler Components and Flow

The flow and relationships among GLSL, OpenCL, SPIR-V (Vulkan/OpenCL), LLVM IR, and the GPU compiler are shown in the Fig. 7.1, Fig. 7.2 and Fig. 7.3. As shown in Fig. 7.3, OpenCL-C to SPIR-V (OpenCL) can be compiled using **clang + llvm-spirv** tools or a proprietary converter.

As shown in Fig. 7.3, both GLSL and OpenCL use frontend tools to generate SPIR-V. The driver can invoke either the GLSL or OpenCL compiler based on metadata fields in the SPIR-V, as illustrated in Fig. 7.4 and the following figures, which describe offline compilation from GLSL/OpenCL to SPIR-V and online execution using the generated SPIR-V files.




Fig. 7.4: Compiling and Deploying GPU Code from GLSL, Vulkan, and OpenCL

Based on the flows above, the public standards OpenGL and OpenCL provide tools for transferring these data format, as illustrated in Fig. 7.5. The corresponding LLVM IR and SPIR-V formats are listed below.




Fig. 7.5: Conversion between GLSL, OpenCL C, SPIR-V and LLVM-IR

¹¹ <https://www.khronos.org/blog/offline-compilation-of-opencl-kernels-into-spir-v-using-open-source-tooling>

References/add-matrix.ll

```
; ModuleID = 'add-matrix.ll'
target datalayout = "e-p:64:64:64-i1:8:8-i8:8:8-i16:16:16-i32:32:32-i64:64:64-
→f32:32:32-f64:64:64-v16:16:16-v24:32:32-v32:32:32-v48:64:64-v64:64:64-v96:128:128-
→v128:128:128-v192:256:256-v256:256:256-v512:512:512-v1024:1024:1024-G1"
target triple = "spir64-unknown-unknown"

; Function Attrs: nounwind
define spir_func <4 x i32> @add_mat(<4 x i32> %a, <4 x i32> %b) #0 {
entry:
    %sum = add <4 x i32> %a, %b
    ret <4 x i32> %sum
}

attributes #0 = { nounwind }

!spirv.MemoryModel = !{!0}
!opencl.enable.FP_CONTRACT = !{}
!spirv.Source = !{!1}
!opencl.spir.version = !{!2}
!opencl.used.extensions = !{!3}
!opencl.used.optional.core.features = !{!3}
!spirv.Generator = !{!4}

!0 = !{i32 2, i32 2}
!1 = !{i32 0, i32 0}
!2 = !{i32 1, i32 2}
!3 = !{}
!4 = !{i16 6, i16 14}
```

References/add-matrix.spvasm

```
; SPIR-V
; Version: 1.0
; Generator: Khronos LLVM/SPIR-V Translator; 14
; Bound: 10
; Schema: 0

    OpCapability Addresses
    OpCapability Linkage
    OpCapability Kernel
    %1 = OpExtInstImport "OpenCL.std"
    OpMemoryModel Physical64 OpenCL
    OpSource Unknown 0
    OpName %add_mat "add_mat"
    OpName %a "a"
    OpName %b "b"
    OpName %entry "entry"
    OpName %sum "sum"
    OpDecorate %add_mat LinkageAttributes "add_mat" Export
    %uint = OpTypeInt 32 0
    %v4uint = OpTypeVector %uint 4
    %4 = OpTypeFunction %v4uint %v4uint %v4uint
```

(continues on next page)

(continued from previous page)

```

%add_mat = OpFunction %v4uint None %4
    %a = OpFunctionParameter %v4uint
    %b = OpFunctionParameter %v4uint
    %entry = OpLabel
    %sum = OpIAdd %v4uint %a %b
    OpReturnValue %sum
    OpFunctionEnd

```

Convert between spirv and llvm-ir

```

% pwd
$HOME/git/lbd/References
% llvm-as -o add-matrix.bc add-matrix.ll
% llvm-spirv -o add-matrix.spv add-matrix.bc
% spirv-dis -o add-matrix.spvasm add-matrix.spv
// Convert spirv to llvm-ir again and check the converted llvm-ir is same
// with origin.
% llvm-spirv -r add-matrix.spv -o add-matrix.spv.bc
% llvm-dis add-matrix.spv.bc -o add-matrix.spv.bc.ll
% diff add-matrix.ll add-matrix.spv.bc.ll
1c1
< ; ModuleID = 'add-matrix.ll'
---
> ; ModuleID = 'add-matrix.spv.bc'

```

Install llvm-spriv and llvm with Brew-install

```

% brew install spirv-llvm-translator
% brew install llvm

```

The following explains how the driver identifies whether the SPIR-V source is from GLSL or OpenCL.

SPIR-V binaries contain metadata that can help identify whether they were generated from OpenCL, GLSL, or another language.

- Execution Model

Defined by the *OpEntryPoint* instruction. It is a strong indicator of the source language.

ExecutionModel	Typical Source	Notes
Kernel	OpenCL	Used only by OpenCL C
GLCompute	GLSL or HLSL	Used in compute shaders
Fragment	GLSL or HLSL	For pixel shaders
Vertex	GLSL or HLSL	For vertex shaders

- Capabilities

Declared using *OpCapability*. They provide clues about the SPIR-V's execution model and source.

Capability	Likely Source
Kernel	OpenCL
Addresses	OpenCL
Linkage	OpenCL
Shader	GLSL or HLSL

- Extensions

Declared using *OpExtension*. Some are tied to specific compilers or languages.

Extension	Likely Source
SPV_KHR_no_integer_wrap_decoration	OpenCL
SPV_INTEL_unified_shared_memory	OpenCL (Intel)
SPV_AMD_shader_ballot	GLSL (graphics)

- Memory Model

Defined by *OpMemoryModel*.

- *OpenCL* → OpenCL source
- *GLSL450* → GLSL or HLSL source

- How to Inspect

Use the *spirv-dis* tool to disassemble SPIR-V to human-readable form:

```
spirv-dis kernel.spv -o kernel.spvasm
```

Look for these at the top of the file:

Example (GLSL):

```
OpCapability Shader
OpMemoryModel Logical GLSL450
OpEntryPoint GLCompute %main "main"
```

Example (OpenCL):

```
OpCapability Kernel
OpCapability Addresses
OpMemoryModel Logical OpenCL
OpEntryPoint Kernel %foo "foo"
```

7.3.1 Summary

Feature	Indicates
OpEntryPoint Kernel	OpenCL
OpCapability Shader	GLSL or HLSL
OpMemoryModel OpenCL	OpenCL
OpMemoryModel GLSL450	GLSL or HLSL

- Comparison of OpenCL and OpenGL compute shaders.

- Same:

Both support general-purpose GPU computing.

- Difference:

OpenCL supports GPUs and other accelerator devices/processors. OpenCL C runs on the device, while host code uses the OpenCL runtime. A compute shader is written in GLSL and runs in the OpenGL graphics environment, where it can readily access OpenGL API data¹⁰.

- OpenGL/GLSL vs Vulkan/spir-v.

- High level of API and shader: OpenGL, GLSL.

- Low level of API and shader: Vulkan, spir-v.

Though OpenGL api existed in higher level with many advantages from sections above, sometimes it cannot compete in efficiency with direct3D providing lower levels api for operating memory by user program¹³. Vulkan api is lower level⁷ s C/C++ api to fill the gap allowing user program to do these things in OpenGL to compete against Microsoft direct3D. Here is an example¹⁴. Meanwhile glsl is C-like language. The vulkan infrastructure provides tool, glslangValidator¹⁵, to compile glsl into an Intermediate Representation (IR) called spir-v off-line. As a result, it saves part of compilation time from glsl to gpu instructions on-line since spir-v is an IR of level closing to llvm IR¹⁶. In addition, vulkan api reduces gpu drivers efforts in optimization and code generation¹³. These standards provide user programmer option to using vulkan/spir-v instead of OpenGL/glsl, and allow them pre-compiling glsl into spir-v off-line to saving part of on-line compilation time.

With vulkan and spir-v standard, the gpu can be used in OpenCL for Parallel Programming of Heterogeneous Systems¹⁷¹⁸. Similar with Cuda, a OpenCL example for fast Fourier transform (FFT) is here¹⁹. Once OpenCL grows into a popular standard when more computer languages or framework supporting OpenCL language, GPU will take more jobs from CPU²⁰.

Most GPUs have 16 or 32 Lanes in a SIMD processor (Warp), vulkan provides Subgroup operations to data parallel programming on Lanes of SIMD processor²¹.

Subgroup operations provide a fast way for moving data between Lanes intra Warp. Assuming each Warp has four Lanes. The following table lists result of reduce, inclusive and exclusive operations.

Table 7.5: Lists each Lane's value after **Reduce**, **Inclusive** and **Exclusive** operations respectively

Lane	0	1	2	3
Initial value	a	b	c	d
Reduce	OP(abcd)	OP(abcd)	OP(abcd)	OP(abcd)
Inclusive	OP(a)	OP(ab)	OP(abc)	OP(abcd)
Exclusive	not define	OP(a)	OP(ab)	OP(abc)

¹⁰ <https://stackoverflow.com/questions/15868498/what-is-the-difference-between-opengl-and-opengl-compute-shader>

¹³ Vulkan offers lower overhead, more direct control over the GPU, and lower CPU usage...By allowing shader pre-compilation, application initialization speed is improved...A Vulkan driver only needs to do GPU specific optimization and code generation, resulting in easier driver maintenance ...¹² https://en.wikipedia.org/wiki/Vulkan#OpenGL_vs._Vulkan

¹² <https://en.wikipedia.org/wiki/Vulkan>

¹⁴ <https://github.com/SaschaWillems/Vulkan/blob/master/examples/triangle/triangle.cpp>

¹⁵ glslangValidator is the tool used to compile GLSL shaders into SPIR-V, Vulkan's shader format. https://vulkan.lunarg.com/doc/sdk/latest/windows/spirv_toolchain.html

¹⁶ SPIR 2.0: LLVM IR version 3.4. SPIR-V 1.X: 100% Khronos defined Round-trip lossless conversion to LLVM. https://en.wikipedia.org/wiki/Standard_Portable_Intermediate_Representation

¹⁷ <https://www.khronos.org/opencl/>

¹⁸ https://en.wikipedia.org/wiki/Compute_kernel

¹⁹ <https://en.wikipedia.org/wiki/OpenCL>

²⁰ The OpenCL standard defines host APIs for C and C++; third-party APIs exist for other programming languages and platforms such as Python,[15] Java, Perl[15] and .NET.[11]:15 <https://en.wikipedia.org/wiki/OpenCL>

²¹ <https://www.khronos.org/blog/vulkan-subgroup-tutorial>

- Reduce: e.g. `subgroupAdd`. Inclusive: e.g. `subgroupInclusiveAdd`. Exclusive: e.g. `subgroupExclusiveAdd`.
- For examples:
 - ADD operation: $OP(abcd) = a+b+c+d$.
 - MAX operation: $OP(abc) = MAX(a,b,c)$.
- When Lane i is inactive, its value is none.
 - For instance if Lane 0 is inactive, then MUL operation: $OP(abcd) = b*c*d$.

The following is a code example.

An Example of Subgroup Operations in GLSL for Vulkan

```
vec4 sum = vec4(0, 0, 0, 0);
if (gl_SubgroupInvocationID < 16u) {
    sum += subgroupAdd(in[gl_SubgroupInvocationID]);
}
else {
    sum += subgroupInclusiveMul(in[gl_SubgroupInvocationID]);
}
subgroupMemoryBarrier();
```

- NVIDIA GPUs provide `__syncwarp()` for synchronization within a warp; the compiler can also synchronize lanes in the same warp.

To help lanes in the same SIMD processor work efficiently, data-uniformity analysis provides optimization opportunities in register allocation, transformation, and code generation²².

LLVM IR expansion from CPU to GPU is becoming increasingly influential. In fact, LLVM IR has been expanding steadily from version 3.1 until now, as I have observed.


7.4 Unified IR Conversion Flows

This section outlines the intermediate representation (IR) flows for graphics (Microsoft DirectX, OpenGL) and OpenCL compilation across major GPU vendors: NVIDIA, AMD, ARM, Imagination Technologies, and Apple.

7.4.1 Graphics Compilation Flow (Microsoft DirectX & OpenGL)

Graphics shaders are compiled from high-level languages (HLSL, GLSL) into vendor-specific GPU binaries via intermediate representations like DXIL and SPIR-V.

 Each node in the graph is color-coded to indicate its category or role within the structure.



- **Vendor Driver** will call **Vendor Compiler** for on-line compilation.




Fig. 7.6: Graphics and OpenCL Compiler IR Conversion Flow

IR	SSA Form	Virtual Registers	Notes
PTX IR	Yes	Yes	Standalone SSA IR; unlimited virtual registers; lowered to SASS.
GCN IR (LLVM AMDGPU)	Yes	Yes	LLVM IR with AMD intrinsics; backend lowers to GCN ISA.
Burst IR (PowerVR / IMG)	Yes (SSA-like)	Yes	Driver-side IR; physical register allocation done by driver.
Apple Meta-IR	No	No	Not a shader IR; used for tiling, dependencies, pipeline state.

- OpenCL C is the device side code in C language while Host side code is C/C++.
- OpenCL C is compiled to SPIR-V in later versions of OpenCL, while earlier versions used SPIR. SPIR-V has now largely replaced SPIR as the standard intermediate representation.

Table 7.6: Comparison of PTX, GCN IR, Burst IR, and Metal IR

IR Layer	Abstraction Level	Register Model
PTX (NVIDIA)	Virtual ISA; portable across GPU generations; hides hardware scheduling	Virtual registers (%r, %f); mapped to physical registers during SASS lowering
GCN IR (AMD)	Machine IR; tightly coupled to GCN/RDNA architecture; exposes Wavefront semantics	Explicit vector (vN) and scalar (sN) registers; register pressure affects occupancy. AMD's compiler backend can lower vector operations to scalar instructions on low-end GPUs, while preserving vector operations on high-end architectures.
Burst IR (Imagination)	Power-aware IR; optimized for burst-core scheduling and latency hiding	Operand staging model; abstracted register usage; mapped late to USC ISA
Metal IR (Apple)	LLVM-inspired IR; abstracted from developers; tuned for tile shading and threadgroup fusion	Region-based register allocation; dynamic renaming; not exposed as physical register model

²² <https://llvm.org/docs/ConvergenceAndUniformity.html>

❗ NVIDIA, AMD, ARM and Imagination all have exposed LLVM IR and convert

SPIR-V IR to LLVM IR.

- SPIR:
 - For OpenCL development, the IR started from SPIR (LLVM-based IR).
 - SPIRV's Limitation: tightly coupled to specific LLVM versions, making it brittle across.
- SPIR-V:
 - A complete redesign: binary format, not tied to LLVM.
 - Designed for Vulkan, but also supports OpenCL and OpenGL.
 - Enables cross-vendor portability, **shader reflection**, and custom extensions.
 - Used in graphics and compute pipelines, including ML workloads via Vulkan compute.
 - A Vulkan shader written in GLSL is compiled to SPIR-V, then passed to the GPU driver.
 - An OpenCL kernel written in C can be compiled to SPIR-V, then lowered to LLVM IR internally by vendors like AMD or NVIDIA.
 - **Shader reflection:**

Means reading information out of the compiled shader binary (SPIR-V) so the engine can understand what the shader needs —without looking at the original GLSL/HLSL source code.

The key idea: SPIR-V preserves metadata about resources, types, and interfaces, and reflection tools can extract that metadata programmatically.

SPIRV-Cross tool can convert SPIR-V file to GLSL source code.

⚠ Apple

- Uses LLVM IR Partially. Apple supports SPIR-V in Metal and OpenCL, but LLVM IR is not always exposed.
- Metal shaders are compiled via MetalIR, which is LLVM-inspired but not standard LLVM IR. Metal IR is not standard LLVM IR and is not exposed to developers.
- Apple's ML compiler stack may use LLVM IR internally, but it's abstracted from developers.
- Apple is not a vendor of GPU IP, so it does not expose LLVM IR in its ML or graphics APIs for the reasons below:
 - Security: opaque compilation prevents tampering
 - Performance tuning: Apple controls the entire stack for optimal hardware use
 - Developer simplicity: high-level APIs reduce friction

Notes:

- **HLSL → DXIL → DirectX** is Microsoft's graphics pipeline, used on Windows and Xbox.
- **GLSL → SPIR-V → OpenGL/Vulkan** is cross-platform and supported by all vendors.
- Final GPU ISA varies by vendor:
 - NVIDIA: PTX → SASS
 - AMD: LLVM IR → GCN/RDNA
 - ARM: Mali ISA
 - Imagination: USC ISA
 - Apple: Metal GPU ISA

Notes:

- **OpenCL C → SPIR → Vendor Driver → GPU ISA** is the standard compilation path.
- Some vendors (e.g., AMD, NVIDIA) may bypass SPIR and compile directly to LLVM IR or PTX.
- Apple deprecated OpenCL in favor of Metal, but legacy support remains.
- **All major GPU vendors use microcode**, but only for complex, multi-cycle, or evolving instructions such as matrix engines, texture units, and transcendental units. Microcode allows new hardware designs to be implemented without modifying the instruction set.


✓ References

- OpenCL Specification: <https://www.khronos.org/opencl/>
- SPIR-V Specification: <https://www.khronos.org/spir>
- DirectX Shader Compiler: <https://github.com/microsoft/DirectXShaderCompiler>
- Imagination E-Series GPU: <https://www.imaginationtech.com/>
- Apple Metal API: <https://developer.apple.com/metal/>

7.4.2 ML and GPU Compilation

This section outlines the intermediate representation (IR) flows used by NVIDIA, AMD, and ARM in machine learning and GPU compilation pipelines. It includes both inference engines and compiler toolchains.

❓ Each node in the graph is color-coded to indicate its category or role within the structure. In AI, the term *runtime* is usually used instead of *driver*.



NVIDIA IR Conversion Flow

NVIDIA supports both TensorRT-based inference and MLIR-based compilation targeting CUDA GPUs is shown as Fig. 7.7.

- SASS stands for Streaming ASSEMBLER, and it represents the native instruction set architecture (ISA) for NVIDIA GPUs.
- TensorRT is a C++ library and runtime developed by NVIDIA for deep learning inference—the phase where trained models make predictions.
 - It works with models trained in frameworks like TensorFlow, PyTorch, and ONNX, converting them into highly optimized engines for execution on NVIDIA hardware^{23,24}.
- **CUDA Kernel IR** is a bridge between LLVM IR and PTX/SASS, or a direct output from TensorRT.
- **LLVM IR** is foundational in many flows, but **TensorRT may skip it** and directly emit CUDA kernels.

²³ <https://resources.nvidia.com/en-us-inference-resources/nvidia-tensorrt>

²⁴ <https://www.geeksforgeeks.org/deep-learning/what-is-tensorrt/>




Fig. 7.7: NVIDIA IR Conversion Flow

- Although MLIR dialects may be publicly available, they are typically hardware-dependent and tailored to specific vendors' GPU architectures. As a result, their applicability is limited to the corresponding hardware platforms.
- MLIR GPU Dialects is public but it is for Nvidia's GPU.




Fig. 7.8: Tensor-RT

Tensor-RT is a SDK bundles tools for model import, graph optimization, precision calibration, engine building, runtime execution, and specialized LLM acceleration.

TensorRT-SDK includes:²⁵

- TensorRT compiler —converts neural networks into optimized inference engines
- TensorRT runtime —executes the compiled engine on NVIDIA GPUs
- TensorRT Model Optimizer —performs graph-level transformations, quantization, fusion
- TensorRT-LLM —specialized compiler/runtime for large language models
- TensorRT for RTX —consumer GPU-optimized inference stack

²⁵ <https://developer.nvidia.com/tensorrt>

- TensorRT Cloud —cloud-based inference services
- Tools & utilities —trtexec, calibration tools, profiling tools
- Kernel libraries —highly optimized CUDA kernels for inference
- Precision calibration —FP16, BF16, FP8, INT8, INT4, FP4 support
- Integration APIs —C++/Python APIs for engine building and execution

AMD IR Conversion Flow

AMD uses ROCm and MIOpen for ML workloads, with LLVM-based compilation targeting GCN or RDNA architectures is shown as Fig. 7.9.




Fig. 7.9: AMD IR Conversion Flow

- **ROCm device libraies (.bc):** ROCm device libraries (.bc files) are LLVM bitcode libraries containing AMD’s **built-in GPU functions**. Before lowering to GCN-IR, the compiler **links** your generated LLVM IR with these **.bc libraries**, so all intrinsic calls (math, atomics, wave ops, etc.) are resolved at the LLVM level. This ensures the final GCN machine code has concrete implementations for all built-ins.
 - They are LLVM bitcode archives containing AMD’s GPU-side implementations of:
 - * OCML —optimized math library (sin, cos, exp, log, pow, etc.)
 - * OCKL —kernel library (barriers, wave ops, LDS operations)
 - * OCLC —language runtime (OpenCL builtins, address space helpers)
 - * Atomic / synchronization intrinsics
 - * Vectorized math routines
 - * FP mode handling, denorm behavior, rounding rules
- **ROCm** is not just a compiler or driver —it includes a full runtime stack that enables AMD GPUs to execute compute workloads across HIP (Heterogeneous-compute Interface for Portability), OpenCL, and ML frameworks. It’s **analogous to NVIDIA’s CUDA runtime** but built around open standards like HSA (Heterogeneous System Architecture)²⁶ and LLVM.
- **MIOpen Graph IR** includes different form and structure. **(Pre-MLIR)** and **(Post-GCN)** are different.
 - Developers interact with MIOpen via high-level APIs (e.g., miopenConvolutionForward) —not via direct IR manipulation.

²⁶ HSA is an open standard developed to simplify programming across heterogeneous systems —especially those combining CPUs and GPUs. It defines:

- Agents: CPUs, GPUs, and other compute units treated uniformly
- Queues: Asynchronous command queues for dispatching kernels
- Memory model: Shared virtual memory across agents
- Signals: Lightweight synchronization primitives

- While MIOpen itself is open source (GitHub repo), its graph IR format and transformation passes are internal.

ARM IR Conversion Flow

ARM supports both CPU/NPU deployment (e.g., Ethos-U/N) and GPU execution (e.g., Mali via Vulkan). The IR flow diverges depending on the target is shown as Fig. 7.10.




Fig. 7.10: ARM IR Conversion Flow

- Node “**Mali GPU (Vulkan)**” is the SPIR-V compilation flow that illustrated in the previous section.
- **Ethos-N** is ARM’s NPU. **Cortex-M** is ARM’s CPU.

ARM Codegen generally emits instructions for CPU/NPU execution, but for certain NN operations (especially those requiring vendor-specific acceleration), it may generate function calls into the Ethos-N driver, which then orchestrates execution on the NPU.

✓ Common Case: Direct NPU/CPU Instruction Generation

- For operations that are well-supported by the NPU or CPU, the codegen backend emits hardware-specific instructions or IR directly.
- These are scheduled for execution on the CPU or passed to the Ethos-N via its driver stack.

⚙️ Special Case: Function Calls to Ethos-N Driver

- For complex or fused neural network operations (e.g., custom activation functions, quantized convolutions, or optimized memory layouts), the codegen may emit calls (**LLVM-IR `call`**) to precompiled driver functions.
- These function calls act as entry points into the Ethos-N runtime, which handles:
 - Memory management
 - Scheduling
 - Firmware-level execution
 - Hardware-specific optimizations

Imagination Technologies IR Conversion Flow

Notes:

- E-Series GPUs support up to **200 TOPS INT8/FP8** for edge AI workloads [B](https://www.techpowerup.com/336545/imagination-announces-e-series-gpu-ip-with-burst-processors-and-up-to-200-tops?copilot_analytics_metadata=eyJldmVudEluZm9fY29udmVyc2F0aW9uSWQjOiJMSINQWjRaWjY5Y2ZuN0VnWnJEVzEiLCJldmVudEluZm9fb3D%3D&citationMarker=9F742443-6C92-4C44-BF58-8F5A7C53B6F1).
- The architecture is **programmable**, supporting **graphics and AI** workloads simultaneously.

- ## Comparison Summary

References

- ## 7.4. Unified IR Conversion Flows

7.5 Accelerate ML/DL on OpenCL/SYCL




Fig. 7.12: ML graph scheduling in both the compiler and runtime

As shown in Fig. 7.12, the Device, such as a GPU or a CPU+NPU, is capable of running the entire ML graph. However, if the Device has only an NPU, then operations like Avg-Pool, which require CPU support, must run on the Host side. This introduces communication overhead between the Host and the Device.

Similar to OpenGL shaders, the “kernel” function may be compiled either on-line or off-line and then sent to the GPU as a programmable function.

To run machine-learning (ML) workloads efficiently, GPU/NPU platforms implement scheduling software in both the graph compiler and runtime. If OpenCL were extended to support ML graphs, graph compilers such as TVM and open-source runtimes could reuse scheduling work²⁷. CUDA Graphs follow a similar idea^{28,29}.

- SYCL uses C++ templates to optimize and generate code for OpenCL and CUDA. It provides a consistent language, APIs, and ecosystem for writing and tuning code for different accelerator architectures, CPUs, GPUs, and FPGAs³⁰.
 - SYCL uses generic programming with templates and generic lambda functions to enable higher-level application software to be cleanly coded with optimized acceleration of kernel code across an extensive range of acceleration backend APIs, such as OpenCL and CUDA³¹.
- DPC++ (OneDPC) compiler: Based on SYCL, DPC++ can compile the DPC++ language for both CPU host and GPU device. DPC++ (Data Parallel C++) is a language developed by Intel and may be adopted into standard C++. The GPU-side (kernel code) is written in C++ but does not support exception handling^{32,33}.
 - Features of Kernel Code:
 - * Not supported:

Dynamic polymorphism, dynamic memory allocations (therefore no object management using new or delete operators), static variables, function pointers, runtime type information (RTTI), and **exception handling**. No virtual member functions, and no variadic functions, are allowed to be called from kernel code. Recursion is not allowed within kernel code.
 - * Supported:

²⁷ <https://easychair.org/publications/preprint/GjhX>

²⁸ <https://developer.nvidia.com/blog/cuda-graphs/>

²⁹ <https://pytorch.org/blog/accelerating-pytorch-with-cuda-graphs/>

³⁰ <https://www.khronos.org/sycl/>

³¹ <https://github.com/codeplaysoftware/sycl-for-cuda/blob/cuda/sycl/doc/GetStartedWithSYCLCompiler.md>

³² <https://www.intel.com/content/www/us/en/developer/tools/oneapi/dpc-compiler.html#gs.cxolyx>

³³ <https://link.springer.com/book/10.1007/978-1-4842-5574-2>




Fig. 7.13: SYCL = C++ template and compiler for Data Parallel Applications on AI on CPUs, GPUs and HPGAs.

Lambdas, operator overloading, templates, classes, and static polymorphism³⁴.

³⁴ Page 14 of DPC++ book.

OPEN SOURCES

- <https://registry.khronos.org/OpenGL-Refpages/>
- <https://www.mesa3d.org>
- <https://www.opengl.org/sdk/>, <https://www.opengl.org/sdk/libs/>

RESOURCES

9.1 Build steps

<https://github.com/Jonathan2251/gpu/blob/master/README.md>

9.2 Alternate formats

The book is also available in the following formats:

9.3 Search this website

- search